

1. Relation Algebra

Q1: $\pi_{theater} \sigma_{title=Zootopia} Schedule$

Q2: $\pi_{(Location.Theater, Location.Address)} \sigma_{Movies.Director=Steven Spielberg} (Movies \times Location)$

Q3: $\pi_{(Location.Address, Location.PhoneNumber)} \sigma_{Location.Theater=Le Champo}$

Q4: $\pi_{(M1.Actor, M2.Actor)} \sigma_{(M1.Title=M2.Title \ \& \ M1.Actor \neq M2.Actor)} (\rho_{M1} Movies \times \rho_{M2} Movies)$

2. Relation $\{S: 20'000t, 10t \text{ per block} \mid 2'000b\}$

Relation $\{R: 100'000t, 10t \text{ per block} \mid 10'000b\}$

Assumption: neither is indexed.

Constraint: 52 blocks of memory.

a. Nested-Join: using two nested for loops (technically 4 but more on that later).

i. For each block in S, $b_i < b_S$: (1)

1. For each block in R, $b_j < b_R$: (2)

a. For each tuple in S, $t_m < t_S$: (3)

i. For each tuple in R, $t_n < t_R$: (4)

1. Test if (t_m, t_n) share condition $\theta_{S.t_S=R.t_R}$

ii. End for

b. End for

2. End for

ii. End for

The cost of the two innermost loops is (10×10) tuples per $(1b_i + 1b_j) = 2b$ per each call on the second for-loop. Since memory can hold at most 52 blocks, we can abstract this cost away and consider only the two outermost loops on b_i and b_j .

The cost, then, is:

- Worst case: $(b_S + b_S \cdot b_R)$ for transfers, $(b_S + b_S \cdot 1)$ for seeks.
- Best case: $(b_S + b_R)$ for transfers, $(1t + 1t)$ for seeks.

b. Sort-Merge: using divide and conquer.

Create sorted chunks that can fit in memory (52 blocks).

- Recursively call the following till the end of the relation:
 - a. Read M blocks of relation into memory
 - b. Sort the in-memory blocks as sorted chunk
 - c. Write sorted chunk to disk
- Given that the input is much larger than the memory, we need several merge passes. Therefore, we read b_b blocks per chunk so that we can merge a group of $(\frac{M}{b_b} - 1)$ chunks per pass.
- A pass reduces the number of chunks by a factor of $(\frac{M}{b_b} - 1)$, and creates chunks longer by the same factor.
- Repeated passes are performed till all runs have been merged into one.

- Total number of merge passes required: $\text{ceil}(\log_{\text{floor}(\frac{M}{b_b}-1)}(\frac{b_s}{M}))$.
- $b_b < 52$ determines the trade-off between number of passes, and disk I/O operation time per pass.
- The cost is then:
 - Number of transfers: $2b_s \cdot (\log_{(\frac{M}{b_b}-1)}(\frac{b_s}{M}))$
 - Number of seeks: $\frac{2b_s}{b_b} \cdot (\log_{(\frac{M}{b_b}-1)}(\frac{b_s}{M}))$
- c. Hash-Join: applicable for equijoins and natural joins.
 - Perfect hashing breaks down S, R into 40 partitions $S: \{40p, \frac{50b}{p}\}$ and $R: \{40p, \frac{250b}{p}\}$
 - Total cost, where $b_b = 10$:
 - Block transfers: $10 \cdot (2'000 + 10'000) = 120'000$
 - Seeks: $2 \cdot (\frac{2'000}{10} + \frac{10'000}{10}) = 2'400$
- 3. See separate XML files.
 - For DTD, keys are declared as attributes. Primary key is defined as *#REQUIRED*. Foreign key is defined as *#IMPLIED*.
 - For XML Schema, primary keys are identified by the *name* attribute. The *id* attributes can serve as foreign keys.
- 4. See separate JSON files.