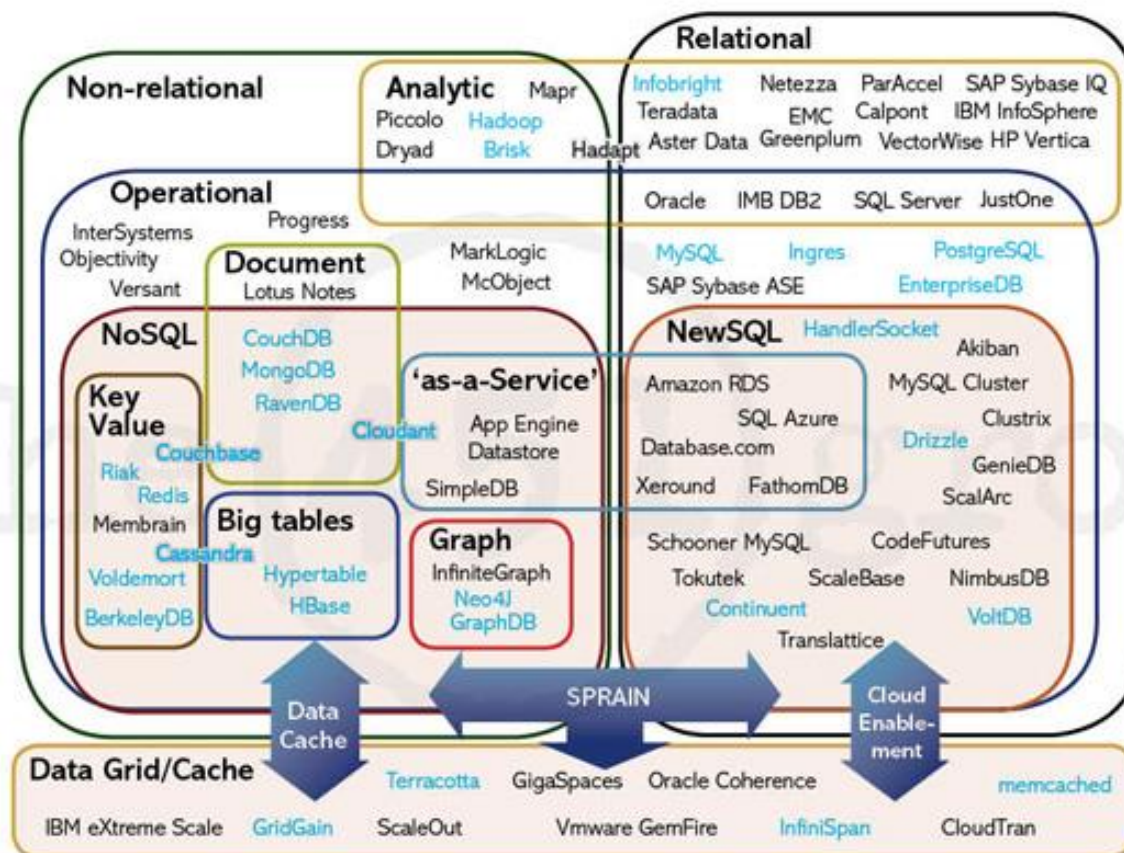




noSQL Databases

Content

- noSQL: Concept and Theory
 - CAP Theory
 - ACID vs EASE
 - noSQL vs RDBMS
- noSQL databases
 - Key-value stores
 - Column Family
- Graph databases



noSQL: Concept

- NoSQL is a non-relational database management system, different from traditional RDBMS
- Carlo Strozzi used the term NoSQL in 1998 to name his lightweight, open-source relational database that did not expose the standard SQL interface
- In 2009, Eric Evans reused the term to refer databases which are non-relational, distributed, and does not conform to ACID
- The NoSQL term should be used as in the Not-Only-SQL and not as No to SQL or Never SQL

HOW TO WRITE A CV



Leverage the NoSQL boom

Motives Behind NoSQL

- Big data.
- Scalability.
- Data format.
- Manageability.

Scalability

- Scale up, Vertical scalability.
 - Increasing server capacity.
 - Adding more CPU, RAM.
 - Managing is hard.
 - Possible down times
- Scale out, Horizontal scalability.
 - Adding servers to existing system with little effort, aka Elastically scalable.
 - Shared nothing.
 - Use of commodity/cheap hardware.
 - Heterogeneous systems.
 - Controlled Concurrency (avoid locks).
 - Service Oriented Architecture.
 - Decentralized to reduce bottlenecks.
 - Avoid Single point of failures.
 - Asynchrony.

CAP Theorem


- Also known as Brewer's Theorem by Prof. Eric Brewer, published in 2000 at UC Berkeley.
- <http://www.cs.berkeley.edu/~brewer/cs262b-2004/PODC-keynote.pdf>



Eric Brewer 2001

CAP Theory

- **Consistency**
 - All replicas contain the same version of data
 - Client always has the same view of the data (no matter what node)
- **Availability**
 - System remains operational on failing nodes
 - All clients can always read and write
- **Partition tolerance**
 - multiple entry points
 - System remains operational on system split (communication malfunction)
 - System works well across physical network partitions



CAP Theorem: satisfying all three at the same time is impossible


CAP Theorem

“Of three properties of a shared data system: data consistency, system availability and tolerance to network partitions, only two can be achieved **at any given moment.**”

Proven by Nancy Lynch et al. MIT labs.


- What the CAP theorem really says:
 - If you cannot limit the number of faults and requests can be directed to any server and you insist on serving every request you receive then you cannot possibly be consistent
- How it is interpreted:
 - You must always give something up: consistency, availability or tolerance to failure and reconfiguration

Proof: A Trivial Two-node System




A Simple Proof

Available and partitioned
Not consistent, we get back old data.




A Simple Proof

Consistent and partitioned
Not available, waiting...

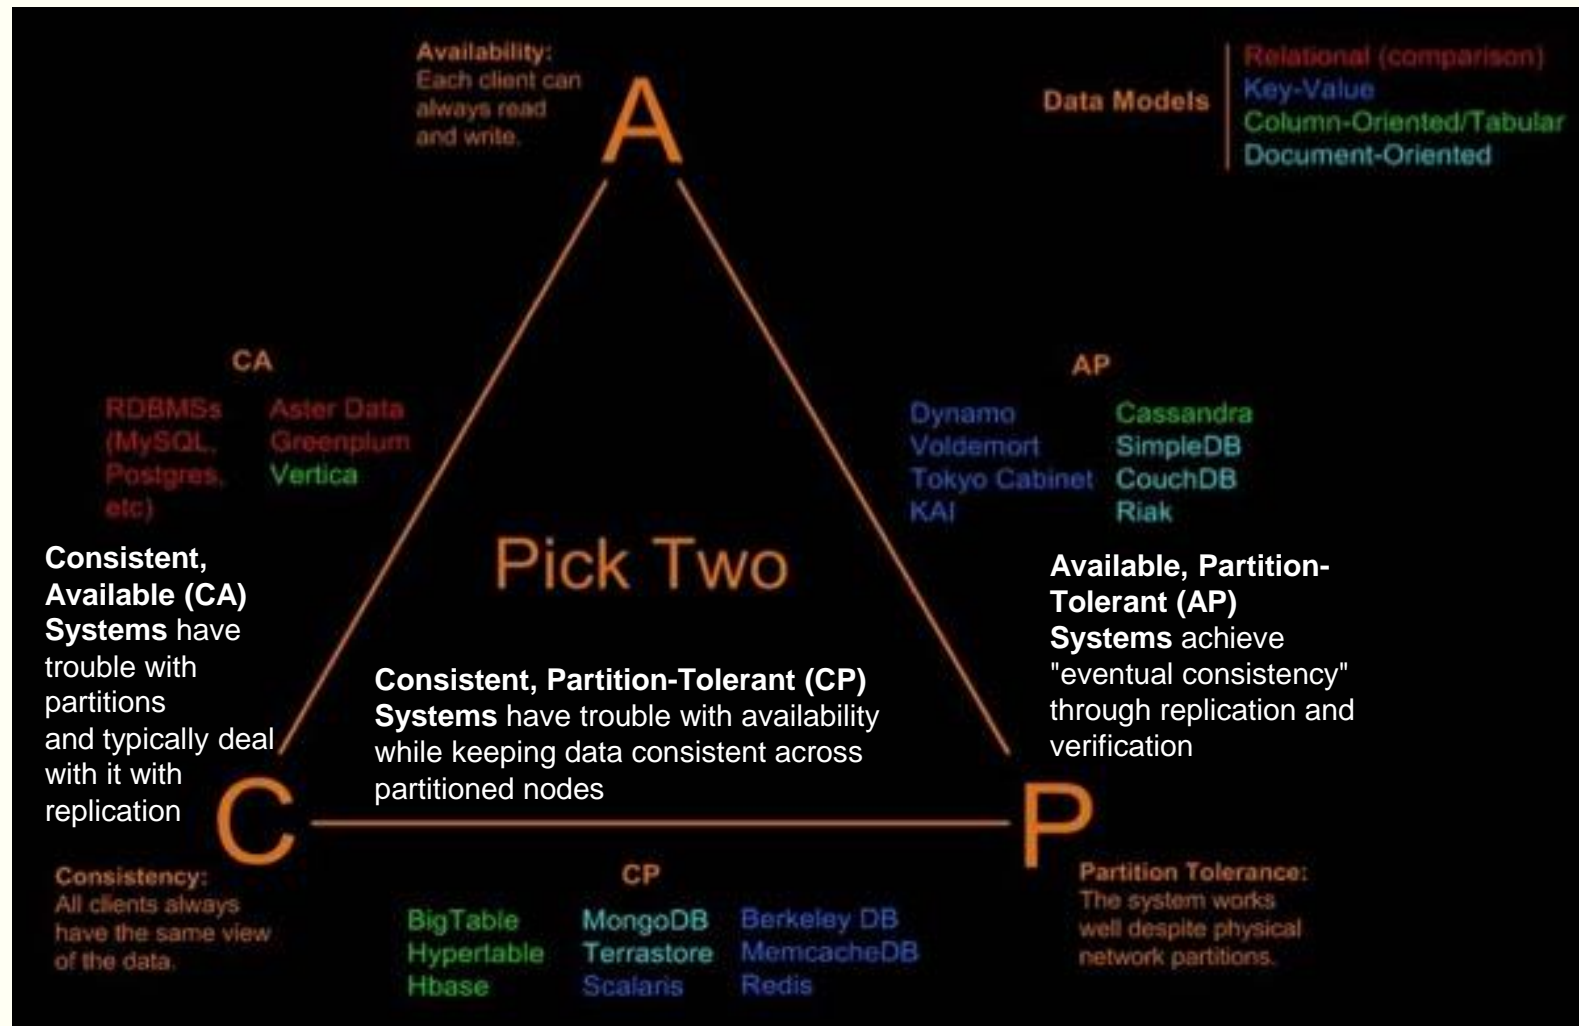


A Simple Proof

Consistent and Available
No partition.



Visual Guide to NoSQL Systems



Consistency: ACID vs BASE

ACID

- Databases require 4 properties:
 - **Atomicity**: When an update happens, it is “all or nothing”
 - **Consistency**: The state of various tables must be consistent (relations, constraints) at all times.
 - **Isolation**: Concurrent execution of transactions produces the same result as if they occurred sequentially.
 - **Durability**: Once committed, the results of a transaction persist against various problems like power failure etc.
- These properties ensure that data is protected even with complex updates and system failures.
- Any data store can achieve Atomicity, Isolation and Durability but do you always need consistency? No.
- By giving up ACID properties, one can achieve higher performance and scalability.

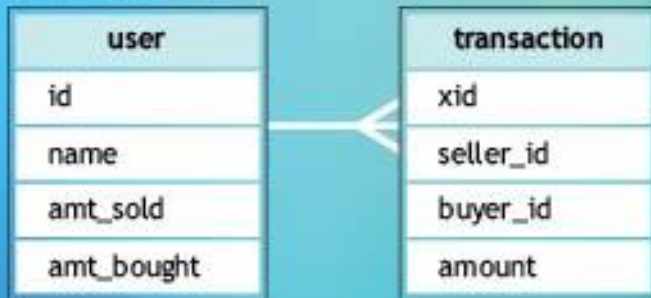
BASE

- Acronym contrived to be the opposite of ACID
 - Basically Available,
 - Soft state,
 - Eventually Consistent
- Characteristics
 - Weak consistency – stale data OK
 - Availability first
 - Best effort
 - Approximate answers OK
 - Simpler and faster

A Toy Example

Pritchett, D.:
BASE: An Acid Alternative
(queue.acm.org/detail.cfm?id=1394128)

Sample Schema



```
Begin transaction
  Insert into transaction(id, seller_id, buyer_id, amount);
  Queue message "update user("seller", seller_id, amount)";
  Queue message "update user("buyer", buyer_id, amount)";
End transaction
For each message in queue
  Begin transaction
    Dequeue message
    If message.balance == "seller"
      Update user set amt_sold=amt_sold + message.amount
        where id=message.id;
    Else
      Update user set amt_bought=amt_bought + message.amount
        where id=message.id;
    End if
  End transaction
End for
```

RDB ACID to NoSQL BASE


Pritchett, D.: BASE: An Acid Alternative (queue.acm.org/detail.cfm?id=1394128)

Atomicity

Consistency

Isolation

Durability



Basically

Available

Soft-state


(State of system may change over time)

Eventually consistent

(Asynchronous propagation)

Data constraints
Smaller,
Schema-driven,
Normalized,
Relational,
Pre-social network

Unstructured data
Big data
Non-relational,
Schema-less,
Distributed,
open-linked data





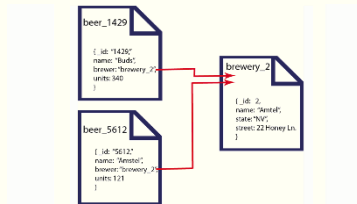
A Clash of cultures

- ACID:
 - Strong consistency.
 - Less availability.
 - Pessimistic concurrency.
 - Complex.
- BASE:
 - Availability is the most important thing. Willing to sacrifice for this (CAP).
 - Weaker consistency (Eventual).
 - Best effort.
 - Optimistic.
 - Simple and fast.

noSQL Data Models

- Key/Value Pairs
- Row/tabular
- Columns
- Documents
- Graphs
- and correspondingly...

| key | value |
|-----------|-------|
| firstName | Bugs |
| lastName | Bunny |
| location | Earth |



Categories of NoSQL storages

- Key-Value

- memcached
- Redis
- Dynamo



- Tabular

- BigTable, HBase



- Column Family

- Cassandra




- Document-oriented

- MongoDB



- Graph (beyond noSQL)


- Neo4j
- TITAN



Key-Value Stores

- “Dynamo: Amazon’s Highly Available Key-Value Store” (2007)
- Data model:
 - Global key-value mapping
 - Highly fault tolerant (typically)
- Examples:
 - Riak, Redis, Voldemort

Key-Value



Column Family (BigTable)

- Google's "Bigtable: A Distributed Storage System for Structured Data" (2006)
- Data model:
 - A big table, with column families
 - Map-reduce for querying/processing
- Examples:
 - HBase, HyperTable, Cassandra


BigTable

| | | | | | |
|---|---|---|---|---|--|
| | | | | 1 | |
| 1 | | | | 1 | |
| | 1 | | 1 | | |
| | | | | | |
| | 1 | 1 | | | |
| | | | | 1 | |
| | 1 | | | 1 | |
| | 1 | | | 1 | |
| | | 1 | | 1 | |
| | | | | 1 | |

Document Databases

- Data model
 - Collections of documents
 - A document is a key-value collection
 - Index-centric, lots of map-reduce
- Examples
 - CouchDB, MongoDB


Document




Graph Databases

- Data model:
 - Nodes with properties
 - Named relationships with properties
 - Hypergraph, sometimes
- Examples:
 - Neo4j, Sones GraphDB, OrientDB, InfiniteGraph, AllegroGraph

Graph DB



Complexity



Key Value Stores and Relational Table

- KV-stores seem very simple. They can be viewed as two-column (key, value) tables with a single key column.
- But they can be used to implement more complicated relational tables:

| State | ID | Population | Area | Senator_1 |
|------------|-----|------------|---------|-----------|
| Alabama | 1 | 4,822,023 | 52,419 | Sessions |
| Alaska | 2 | 731,449 | 663,267 | Begich |
| Arizona | 3 | 6,553,255 | 113,998 | Boozman |
| Arkansas | 4 | 2,949,131 | 53,178 | Flake |
| California | 5 | 38,041,430 | 163,695 | Boxer |
| Colorado | 6 | 5,187,582 | 104,094 | Bennet |
| ... | ... | | | |



Index

KV-stores and Relational Tables

- The KV-version of the previous table includes one table indexed by the actual key, and others by an ID.

| State | ID | ID | Population | ID | Area | ID | Senator_1 |
|------------|-----|-----|------------|-----|---------|-----|-----------|
| Alabama | 1 | 1 | 4,822,023 | 1 | 52,419 | 1 | Sessions |
| Alaska | 2 | 2 | 731,449 | 2 | 663,267 | 2 | Begich |
| Arizona | 3 | 3 | 6,553,255 | 3 | 113,998 | 3 | Boozman |
| Arkansas | 4 | 4 | 2,949,131 | 4 | 53,178 | 4 | Flake |
| California | 5 | 5 | 38,041,430 | 5 | 163,695 | 5 | Boxer |
| Colorado | 6 | 6 | 5,187,582 | 6 | 104,094 | 6 | Bennet |
| ... | ... | ... | ... | ... | ... | ... | ... |

KV-stores and Relational Tables

- We can add indices with new KV-tables:
- Thus KV-tables are used for column-based storage, as opposed to row-based storage typical in older DBMS.

| State | ID |
|------------|-----|
| Alabama | 1 |
| Alaska | 2 |
| Arizona | 3 |
| Arkansas | 4 |
| California | 5 |
| Colorado | 6 |
| ... | ... |


Index

| ID | Population |
|-----|------------|
| 1 | 4,822,023 |
| 2 | 731,449 |
| 3 | 6,553,255 |
| 4 | 2,949,131 |
| 5 | 38,041,430 |
| 6 | 5,187,582 |
| ... | ... |

| Senator_1 | ID |
|-----------|-----|
| Sessions | 1 |
| Begich | 2 |
| Boozman | 3 |
| Flake | 4 |
| Boxer | 5 |
| Bennet | 6 |
| ... | ... |


Index_2

OR: the value field can contain complex data

Key-Values: Examples

- Amazon:
 - Key: customerID
 - Value: customer profile (e.g., buying history, credit card, ..)
- Facebook, Twitter:
 - Key: UserID
 - Value: user profile (e.g., posting history, photos, friends, ...)
- iCloud/iTunes:
 - Key: Movie/song name
 - Value: Movie, Song
- Distributed file systems
 - Key: Block ID
 - Value: Block

System Examples

- **Google File System, Hadoop Dist. File Systems (HDFS)**
- **Amazon**
 - Dynamo: internal key value store used to power Amazon.com (shopping cart)
 - Simple Storage System (S3)




- **BigTable/HBase/Hypertable:** distributed, scalable data storage
- **Cassandra:** “distributed data management system” (Facebook)
- **Memcached:** in-memory key-value store for small chunks of arbitrary data (strings, objects)
- **eDonkey/eMule:** peer-to-peer sharing system




Key-Value Store

- Also called a Distributed Hash Table (DHT)
- Main idea: partition set of key-values across many machines



Challenges



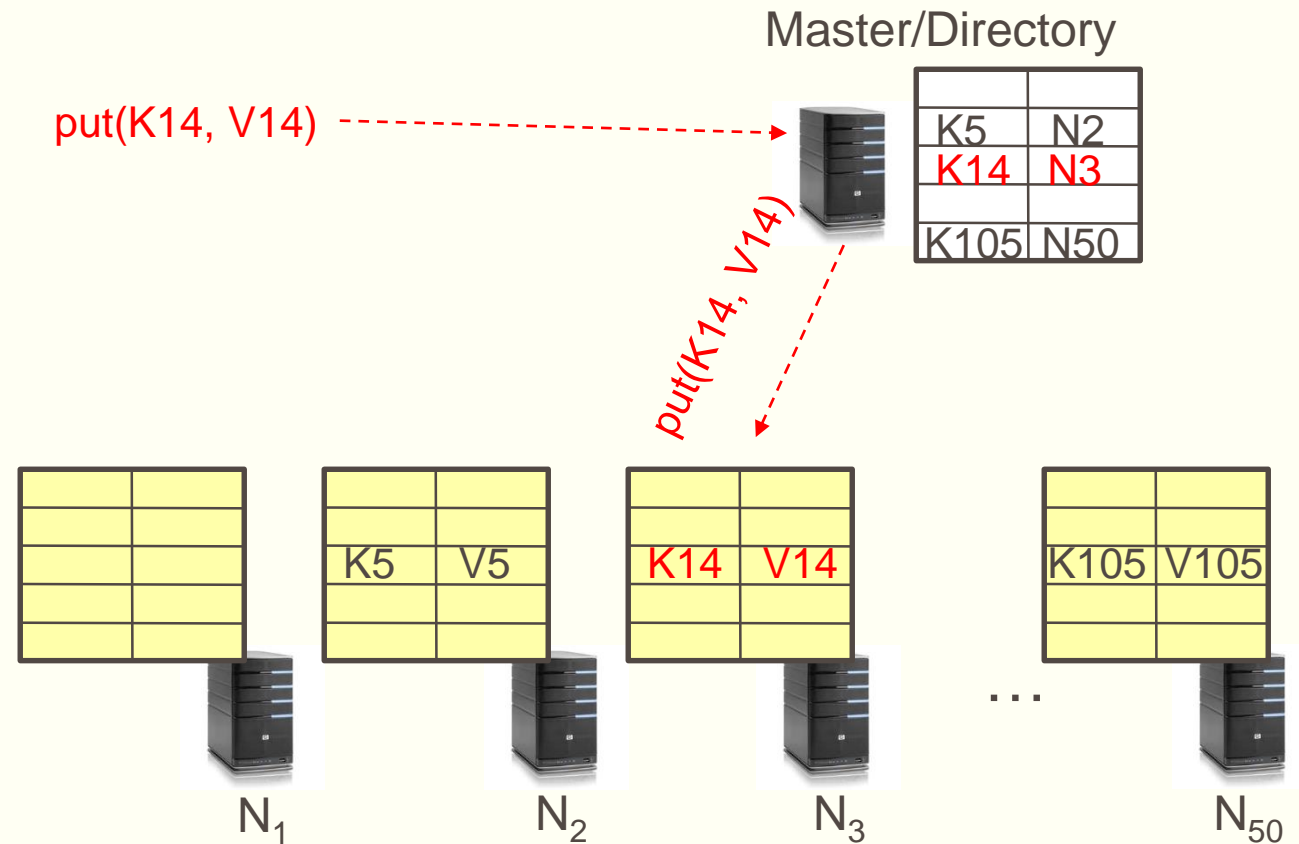
- **Fault Tolerance:** handle machine failures without losing data and without degradation in performance
- **Scalability:**
 - Need to scale to thousands of machines
 - Need to allow easy addition of new machines
- **Consistency:** maintain data consistency in face of node failures and message losses
- **Heterogeneity** (if deployed as peer-to-peer systems):
 - Latency: 1ms to 1000ms
 - Bandwidth: 32Kb/s to several GB/s

Key Operators

- `put(key, value)`: where do you store a new (key, value) tuple?
- `get(key)`: where is the value associated with a given “key” stored?
- And, do the above while providing
 - Fault Tolerance
 - Scalability
 - Consistency

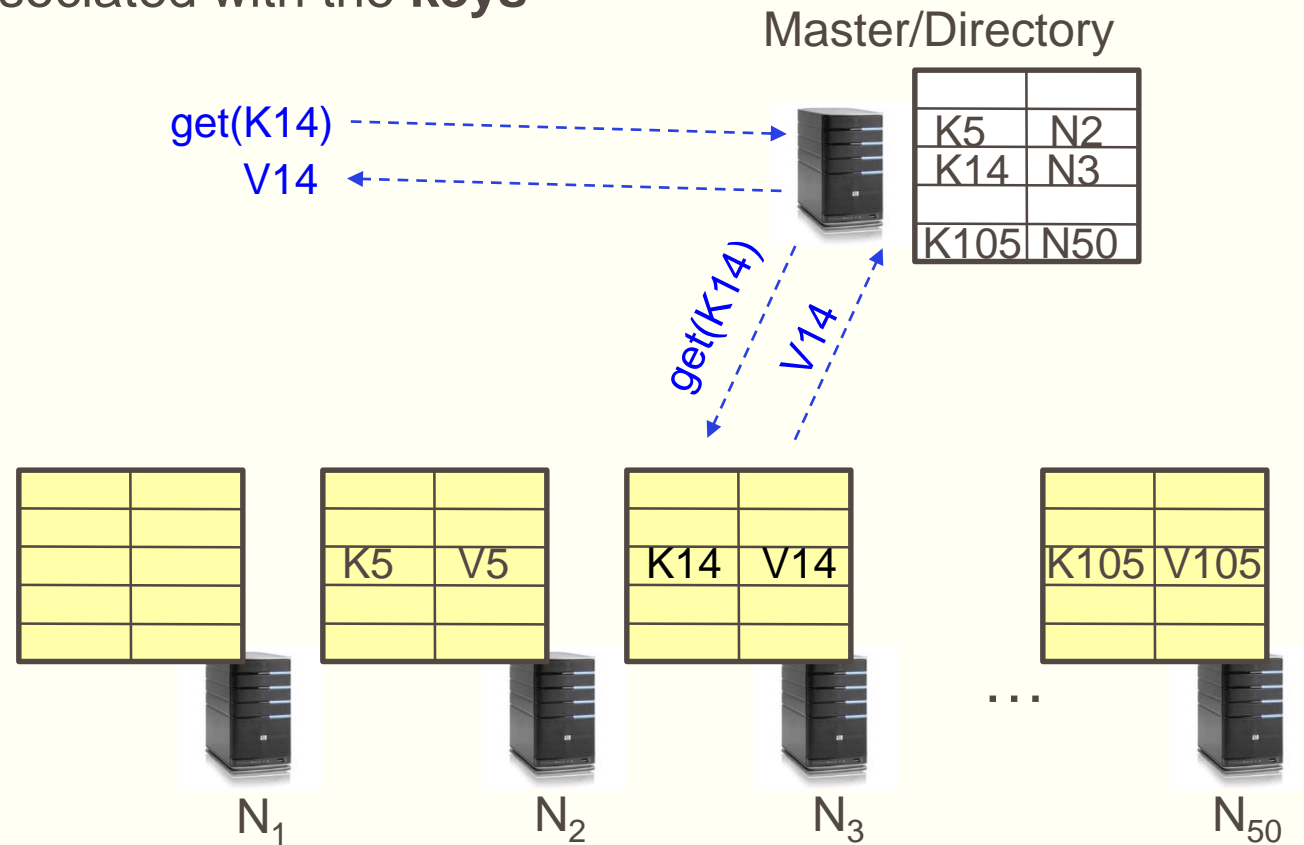
Directory-Based Architecture

- Have a node maintain the mapping between **keys** and the **machines (nodes)** that store the **values** associated with the **keys**



Directory-Based Architecture

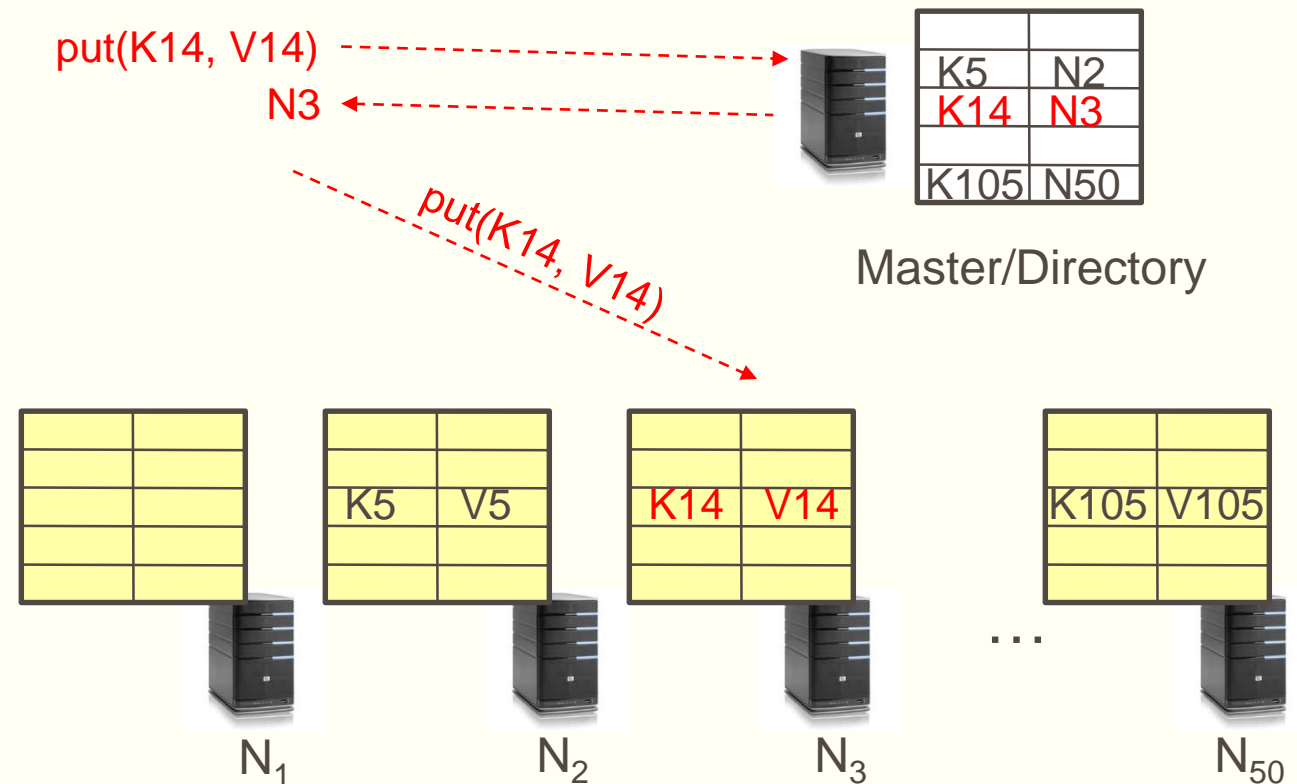
- Have a node maintain the mapping between **keys** and the **machines (nodes)** that store the **values** associated with the **keys**



Having the master relay the requests : recursive query

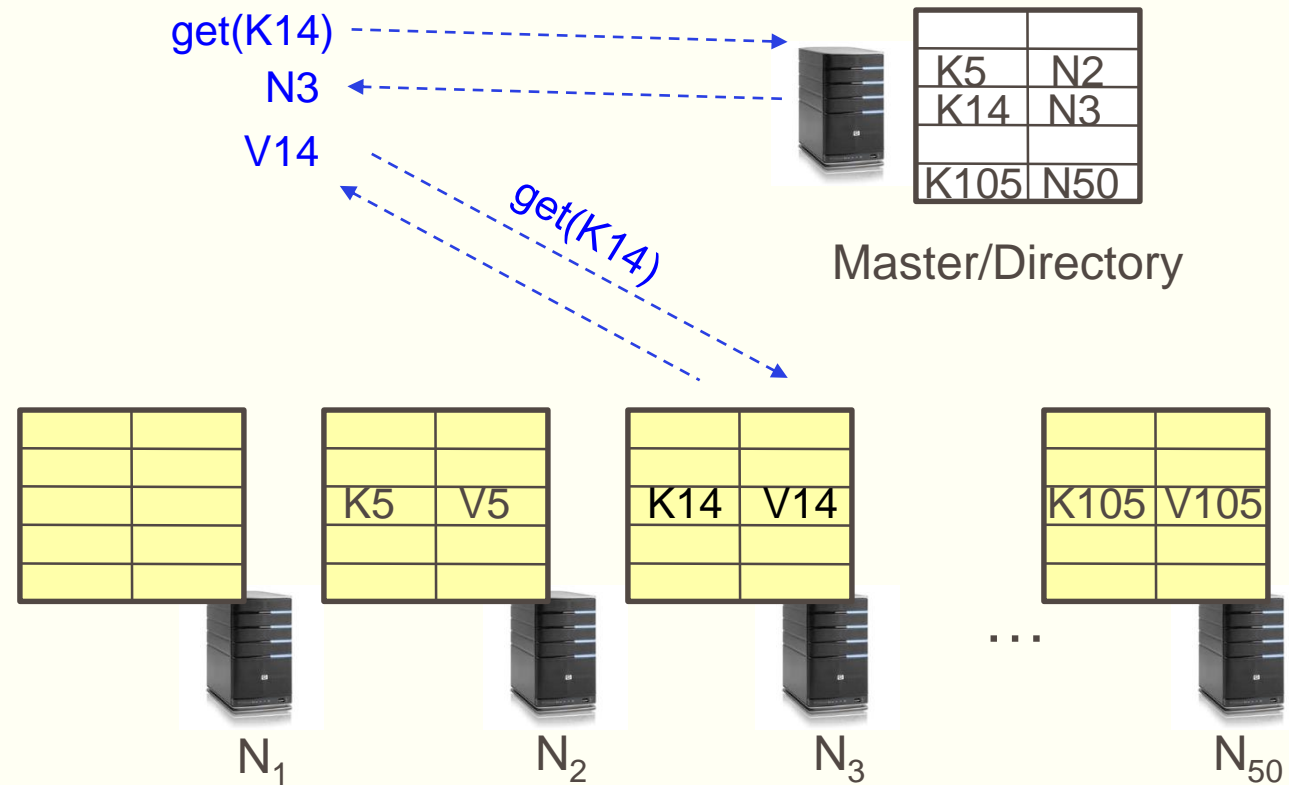
Directory-Based Architecture

- Another method: **iterative query** (this slide)
 - Return node to requester and let requester contact node

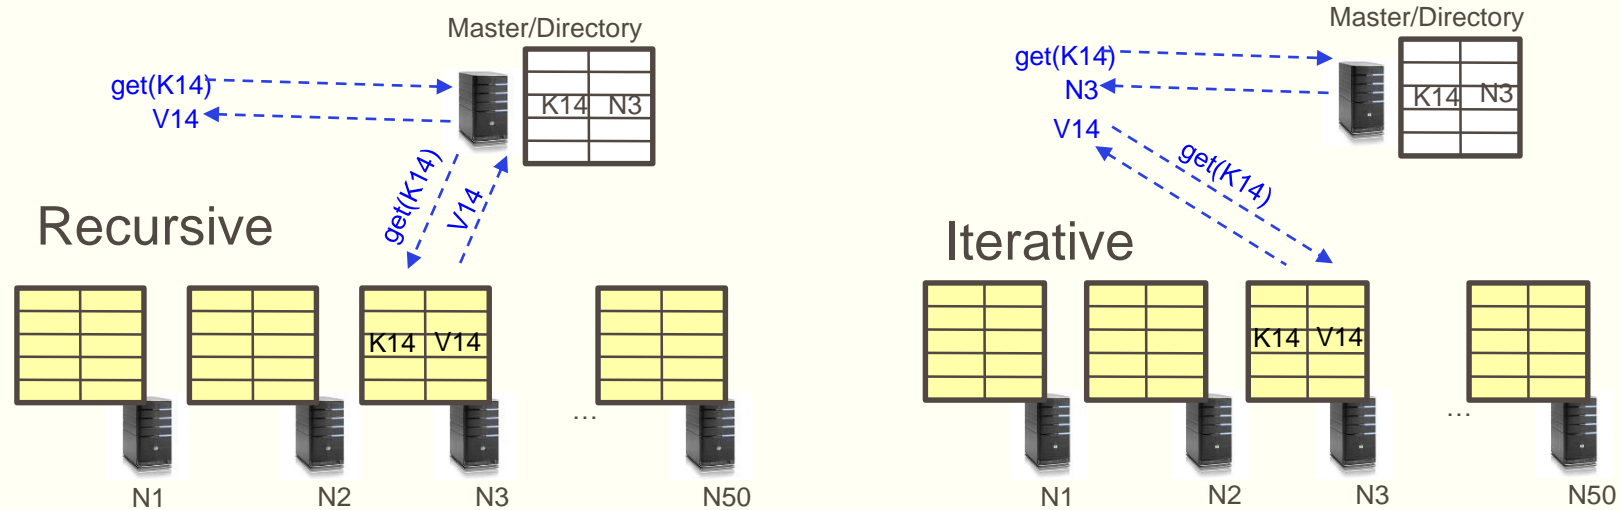


Directory-Based Architecture

- Another method: **iterative query**
 - Return node to requester and let requester contact node



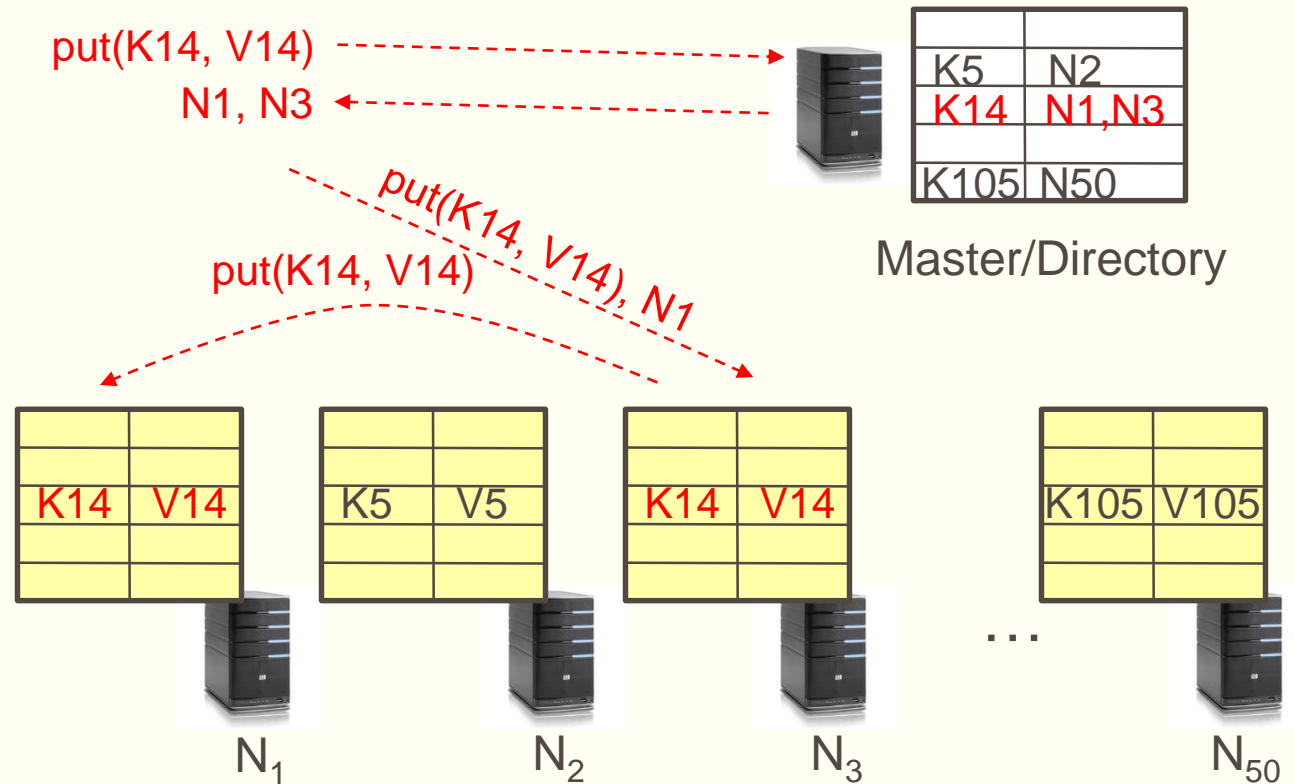
Recursive Vs. Iterative Query



- Recursive Query:
 - Advantages:
 - Faster (latency), as typically master/directory closer to nodes
 - Easier to maintain consistency, as master/directory can serialize puts()/gets()
 - Disadvantages: scalability bottleneck, as all “Values” go through master/directory
- Iterative Query
 - Advantages: more scalable
 - Disadvantages: slower (latency), harder to enforce data consistency

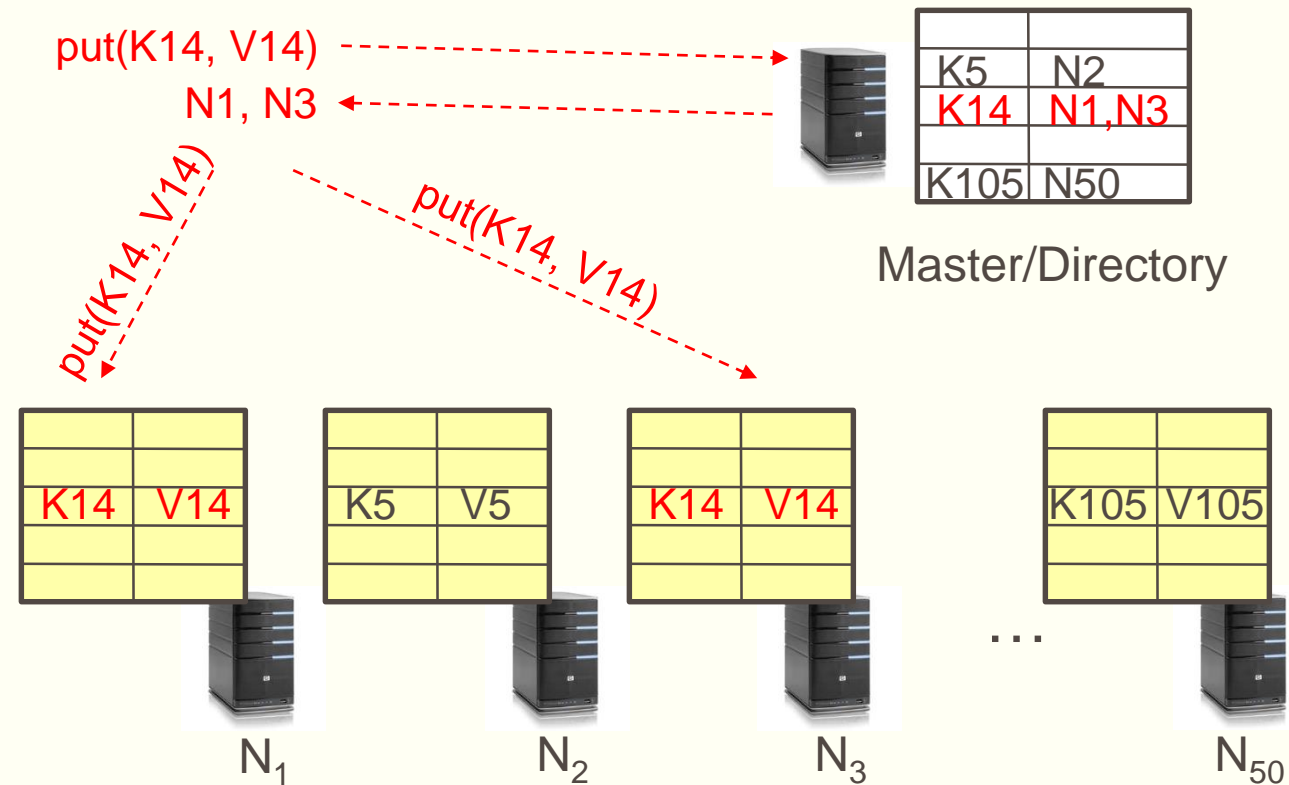
Fault Tolerance

- Replicate value on several nodes
- Usually, place replicas on different racks in a datacenter to guard against rack failures



Fault Tolerance

- Again, we can have
 - **Recursive** replication (previous slide)
 - **Iterative** replication (this slide)

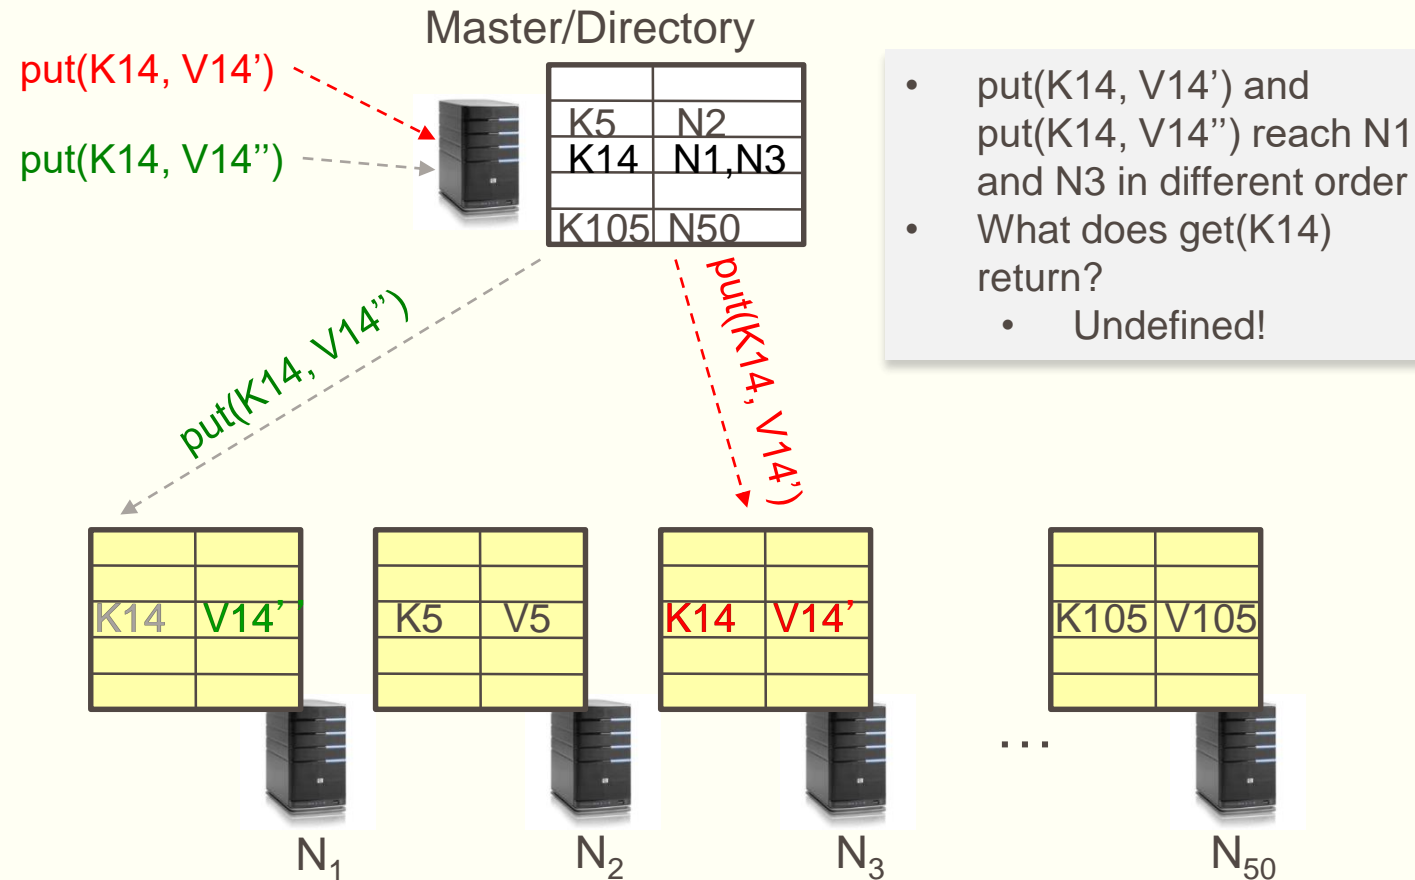


Consistency

- How close does a distributed system emulate a single machine in terms of read and write semantics?
- **Q:** Assume **put(K14, V14')** and **put(K14, V14'')** are concurrent, what value ends up being stored?
- **Q:** Assume a client calls **put(K14, V14)** and then **get(K14)**, what is the result returned by **get()**?
- Above semantics, not trivial to achieve in distributed systems

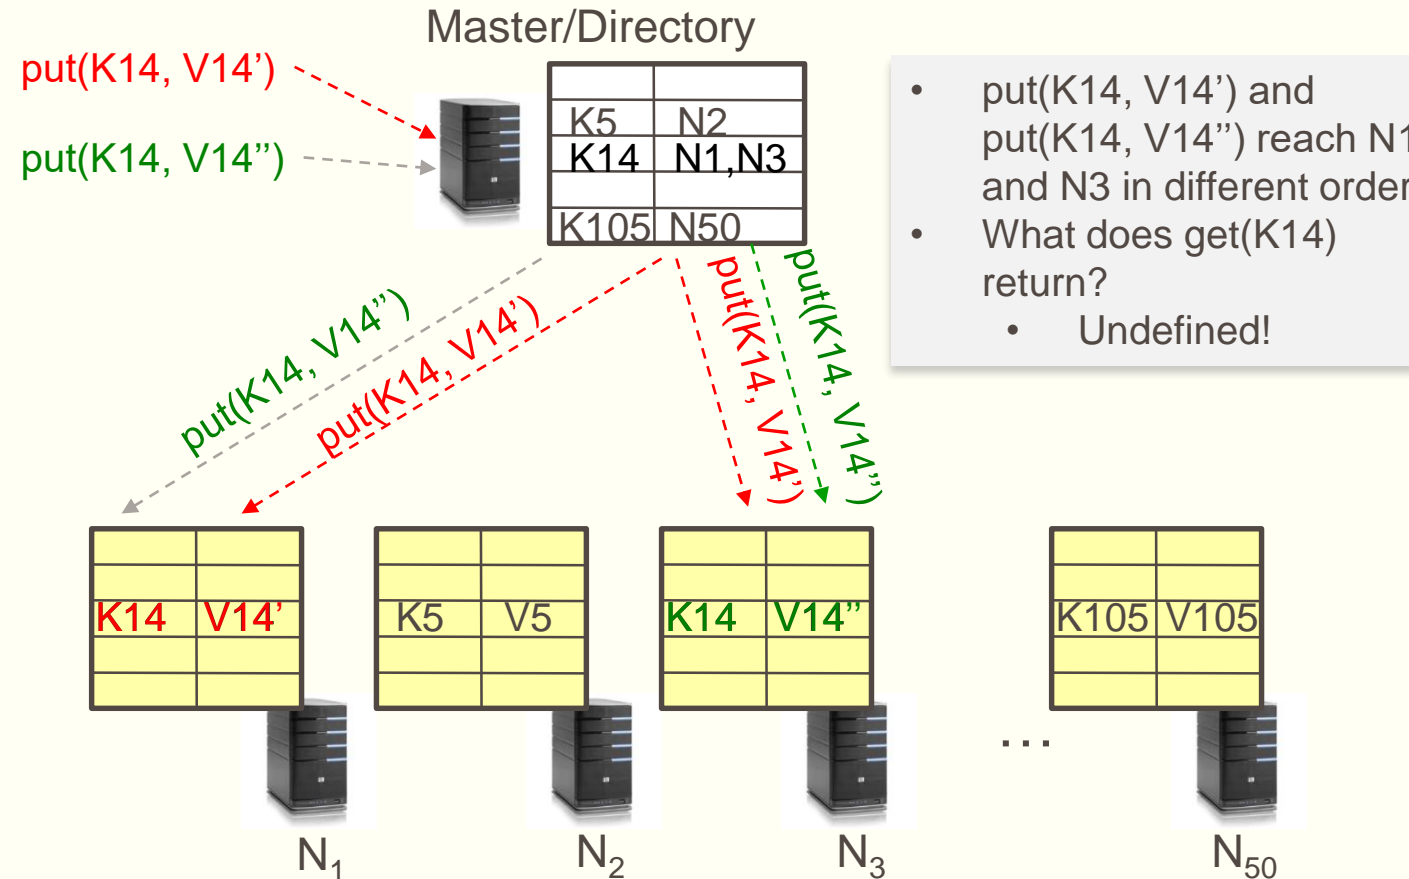
Concurrent Writes (Updates)

- If concurrent updates (i.e., puts to same key) may need to make sure that updates happen in the same order



Concurrent Writes (Updates)

- If concurrent updates (i.e., puts to same key) may need to make sure that updates happen in the same order




Quorum Consensus

- Improve **put()** and **get()** operation performance
- Define a replica set of size N
- **put()** waits for acks from at least W replicas
- **get()** waits for responses from at least R replicas
- $W+R > N$
- Why does it work?
 - There is at least one node that contains the update


Quorum Consensus Example

- $N=3$, $W=2$, $R=2$
- Replica set for K14: {N1, N3, N4}
- Assume put() on N3 fails



Quorum Consensus Example

- Now, issuing `get()` to any two nodes out of three will return the answer



Scalability

- Storage: use more nodes
- Request Throughput:
 - Can serve requests from all nodes on which a value is stored in parallel
 - Large “values” can be broken into blocks (HDFS files are broken up this way)
 - Master can replicate a popular value on more nodes
- Master/directory scalability:
 - Replicate it
 - Partition it, so different keys are served by different masters/directories

Scalability: Load Balancing

- Directory keeps track of the storage availability at each node
 - Preferentially insert new values on nodes with more storage available
- What happens when a new node is added?
 - Move values from the heavy loaded nodes to the new node
- What happens when a node fails?
 - Need to replicate values from failed node to other nodes

Replication Challenges

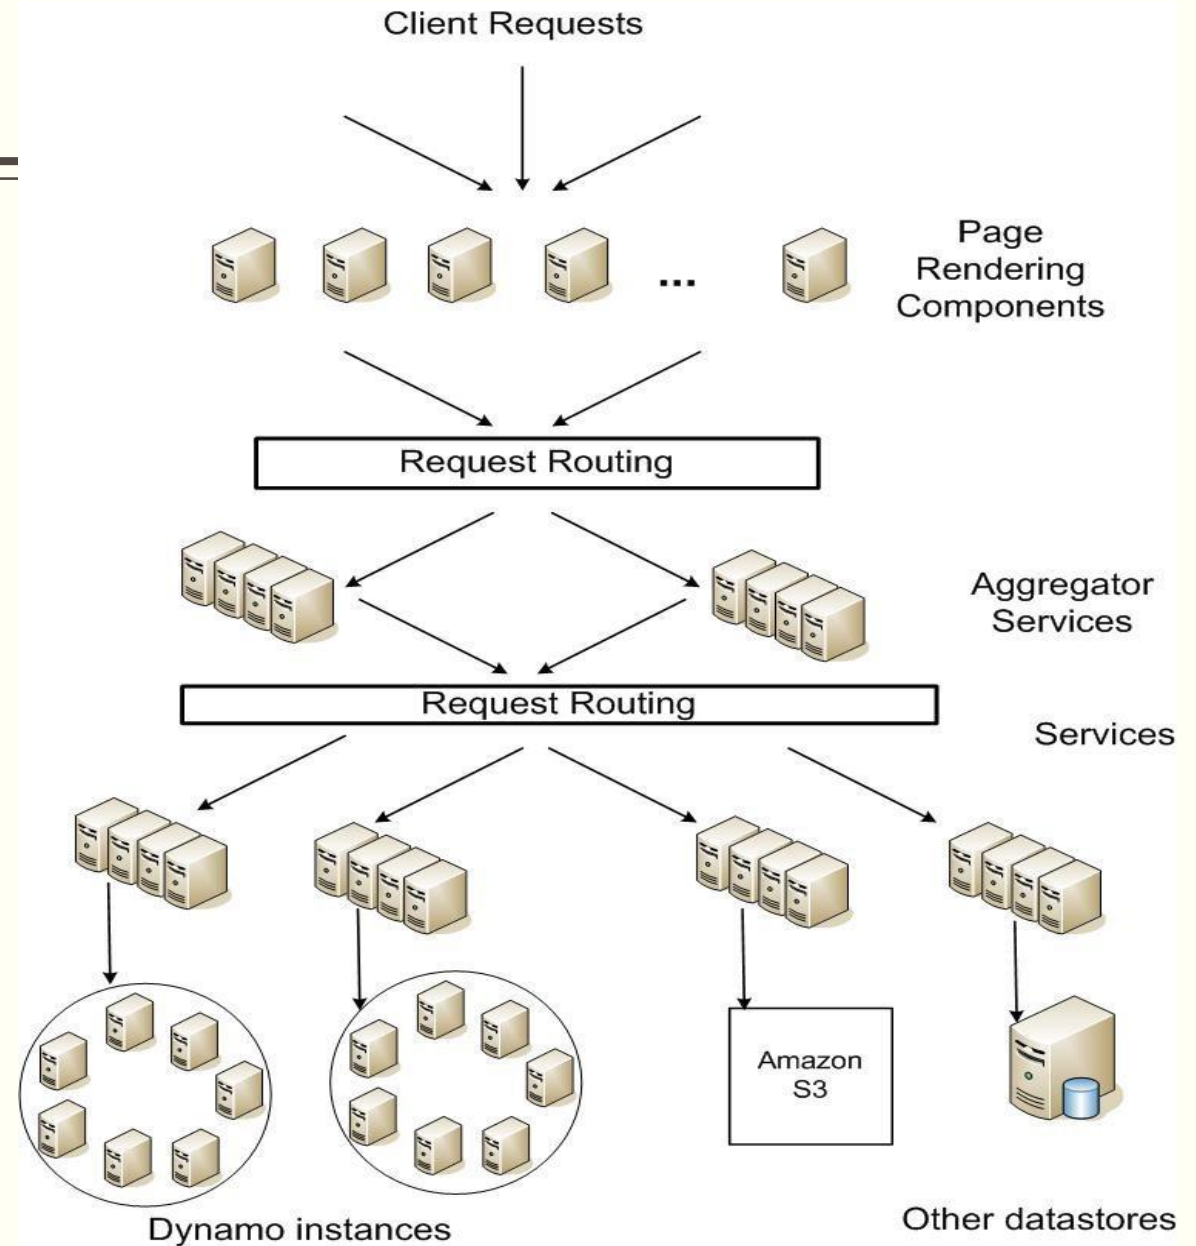
- Need to make sure that a value is replicated correctly
- How do you know a value has been replicated on every node?
 - Wait for acknowledgements from every node
- What happens if a node fails during replication?
 - Pick another node and try again
- What happens if a node is slow?
 - Slow down the entire put()? Pick another node
- In general, with multiple replicas
 - Slow puts and fast gets

Summary: Key-Value Store

- Very large-scale storage systems
- Two operations
 - `put(key, value)`
 - `value = get(key)`
- Challenges
 - Fault Tolerance → replication
 - Scalability → serve `get()`'s in parallel; replicate/cache hot tuples
 - Consistency → quorum consensus to improve put/get performance
- System case study: Dynamo

Amazon Platform Architecture

- simple read and write operations to a data item that is uniquely identified by a key.
- Most of Amazon's services can work with this simple query model and do not need any relational schema.
- targeted applications - store objects that are relatively small (usually less than 1 MB)
- Dynamo targets applications that operate with **weaker consistency** (the "C" in ACID) if this results in high availability.

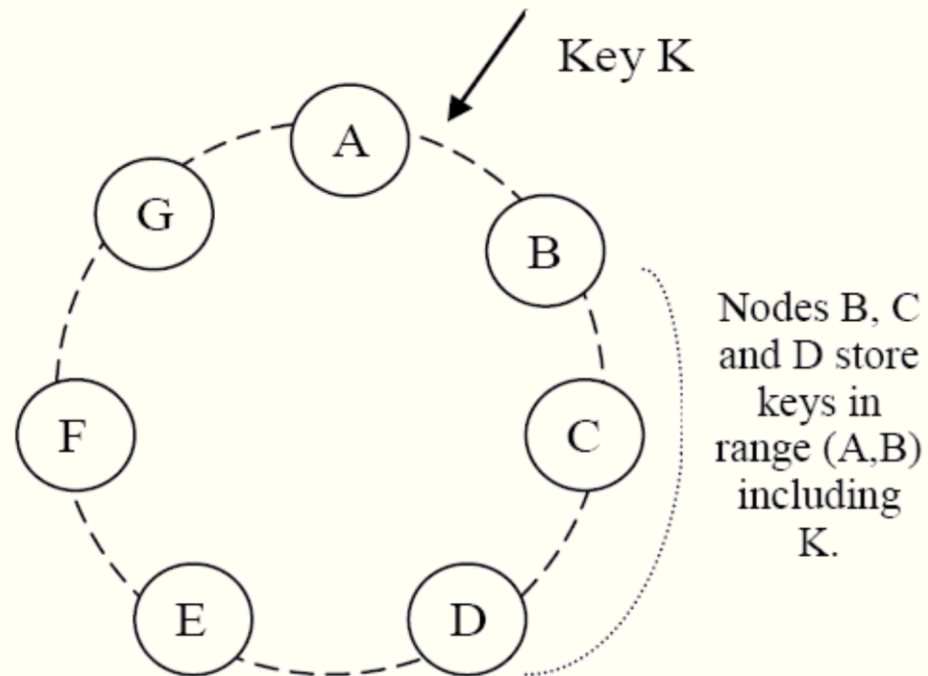


System Architecture

- Partitioning
- High Availability for writes
- Handling temporary failures
- Recovering from permanent failures
- Membership and failure detection

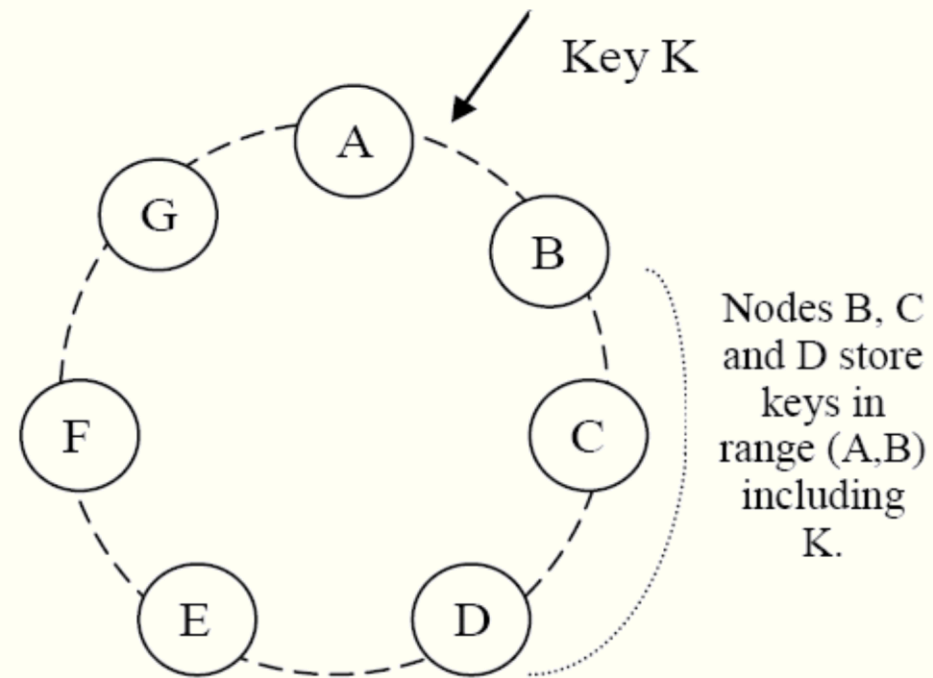
Partition Algorithm

- *Consistent hashing*: the output range of a hash function is treated as a fixed circular space or “ring”.
- *Virtual Nodes*: Each node can be responsible for more than one virtual node.



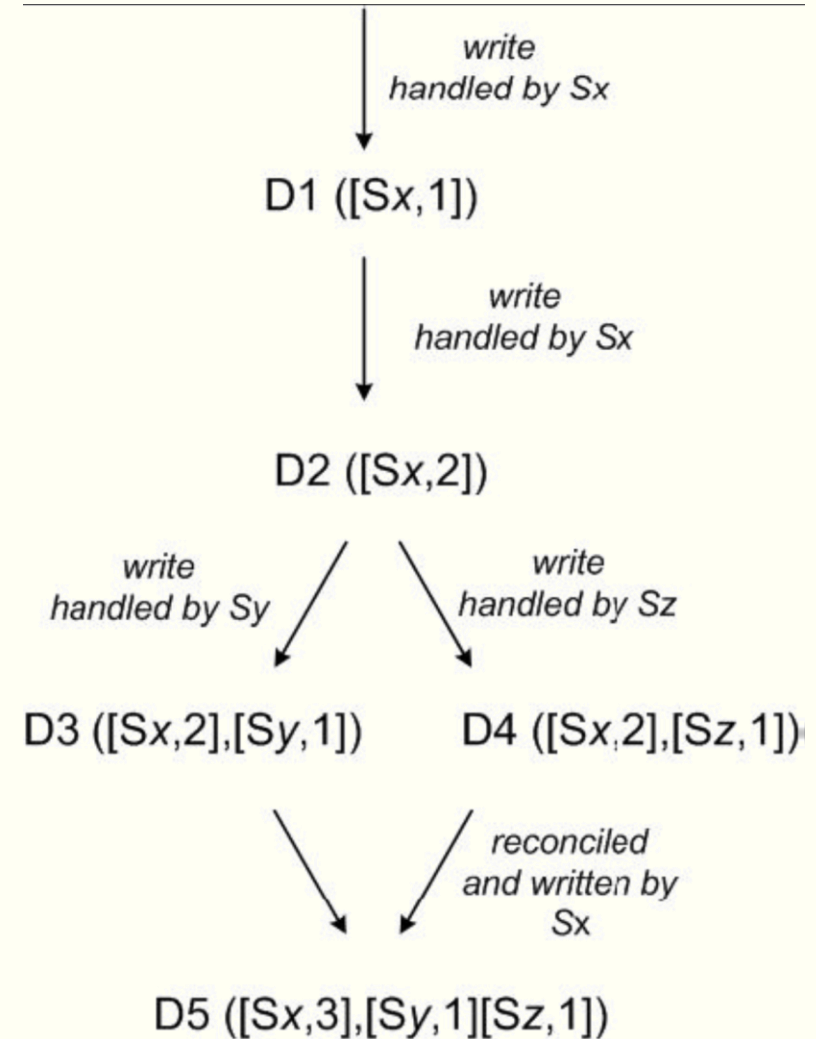
Replication

- Each data item is replicated at N hosts.
- “preference list”: The list of nodes that is responsible for storing a particular key.



Vector Clock

- A vector clock is a list of (node, counter) pairs.
- Every version of every object is associated with one vector clock.
- *If the counters on the first object's clock are less-than-or-equal to all of the nodes in the second clock, then the first is an ancestor of the second and can be forgotten.*



Summary of Techniques Used in Dynamo and Their Advantages

| Problem | Technique | Advantage |
|------------------------------------|---|---|
| Partitioning | Consistent Hashing | Incremental Scalability |
| High Availability for writes | Vector clocks with reconciliation during reads | Version size is decoupled from update rates. |
| Handling temporary failures | Sloppy Quorum and hinted handoff | Provides high availability and durability guarantee when some of the replicas are not available. |
| Recovering from permanent failures | Anti-entropy using Merkle trees | Synchronizes divergent replicas in the background. |
| Membership and failure detection | Gossip-based membership protocol and failure detection. | Preserves symmetry and avoids having a centralized registry for storing membership and node liveness information. |