



MapReduce I

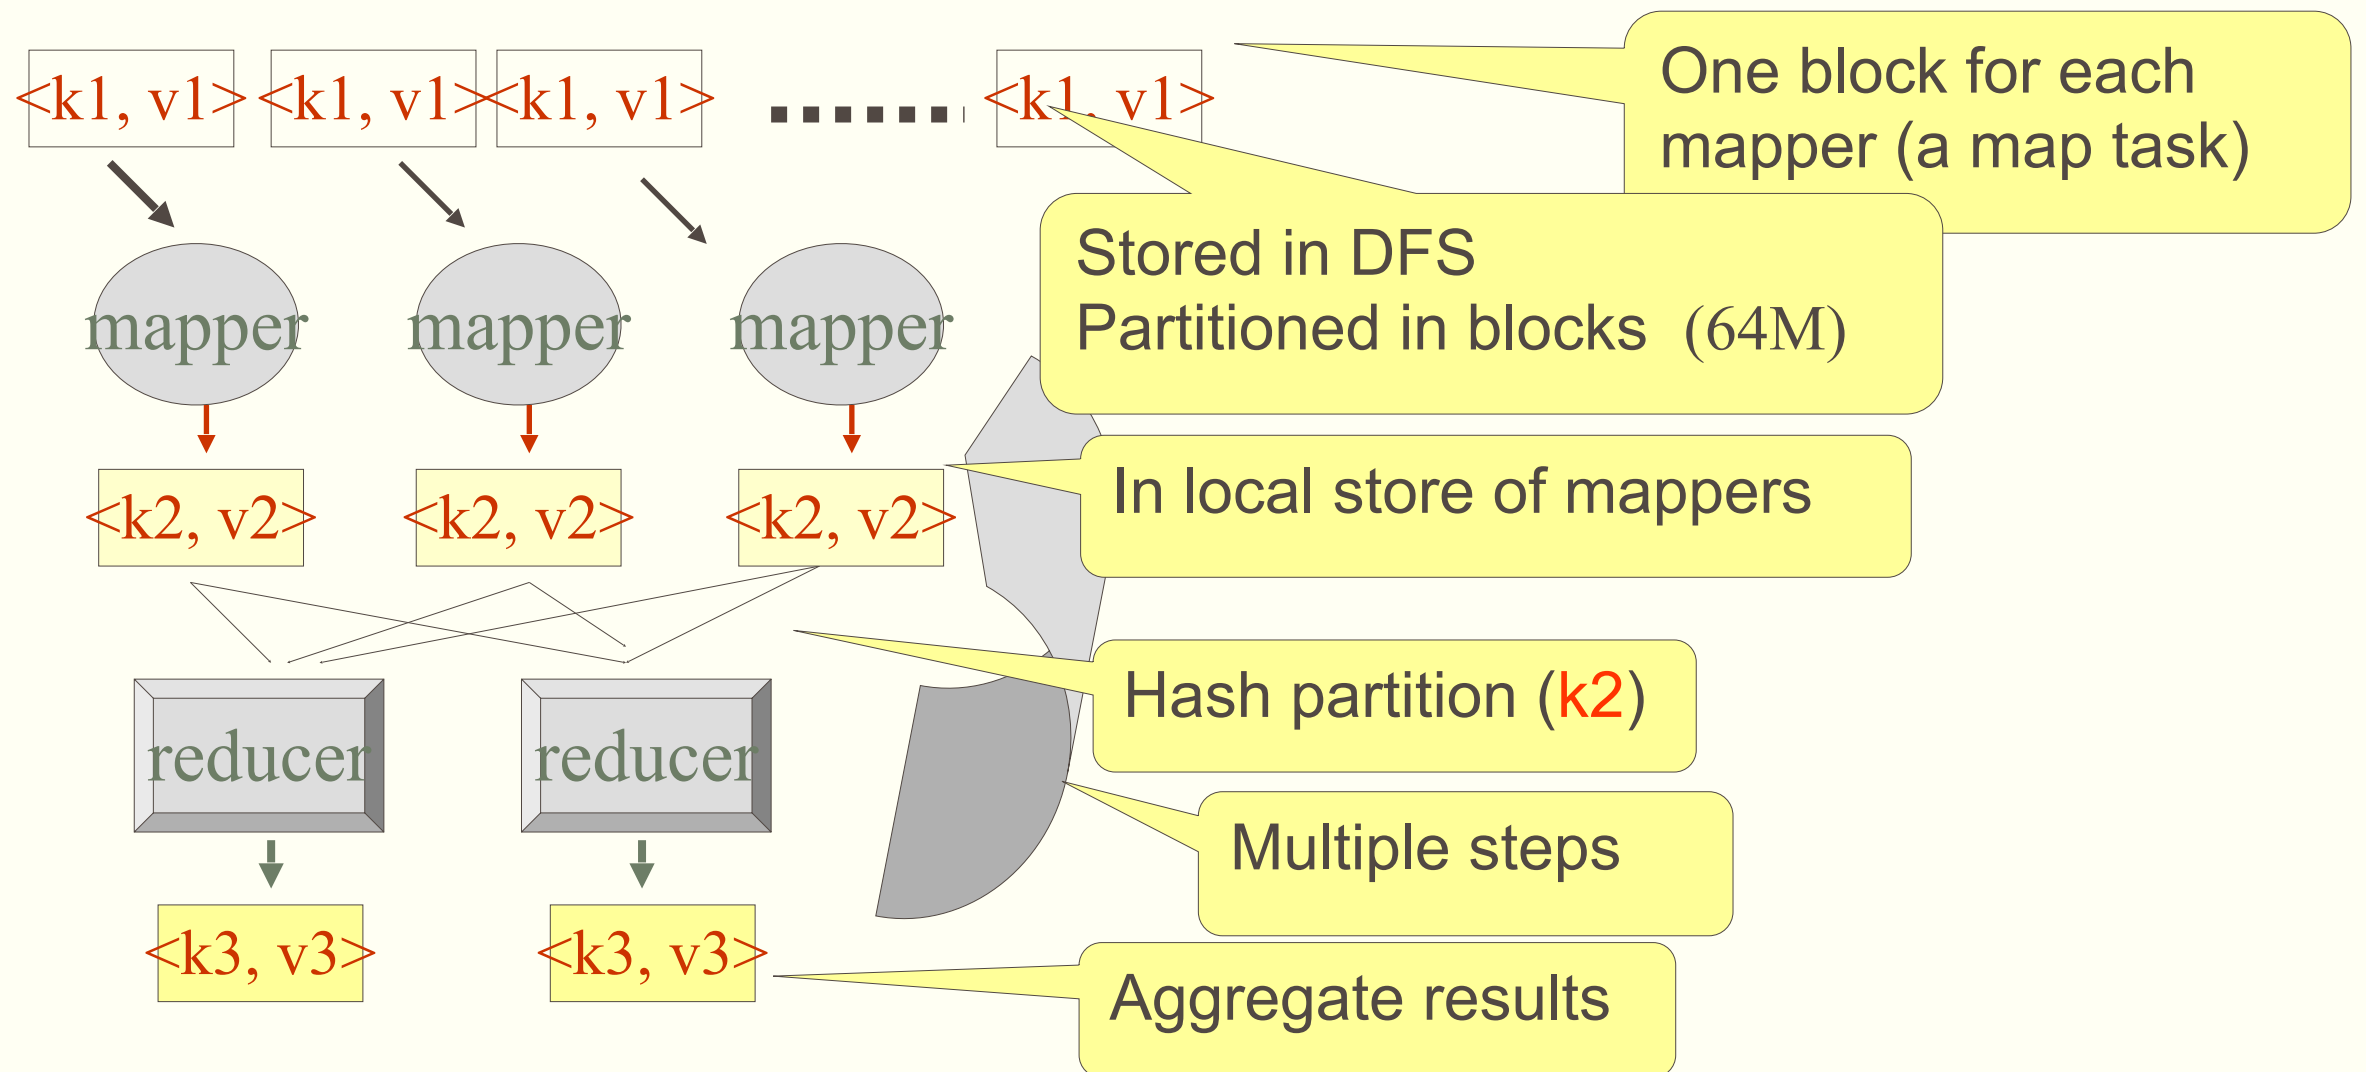
MapReduce

- MapReduce model
- MapReduce for relational operators
- MapReduce for graph querying

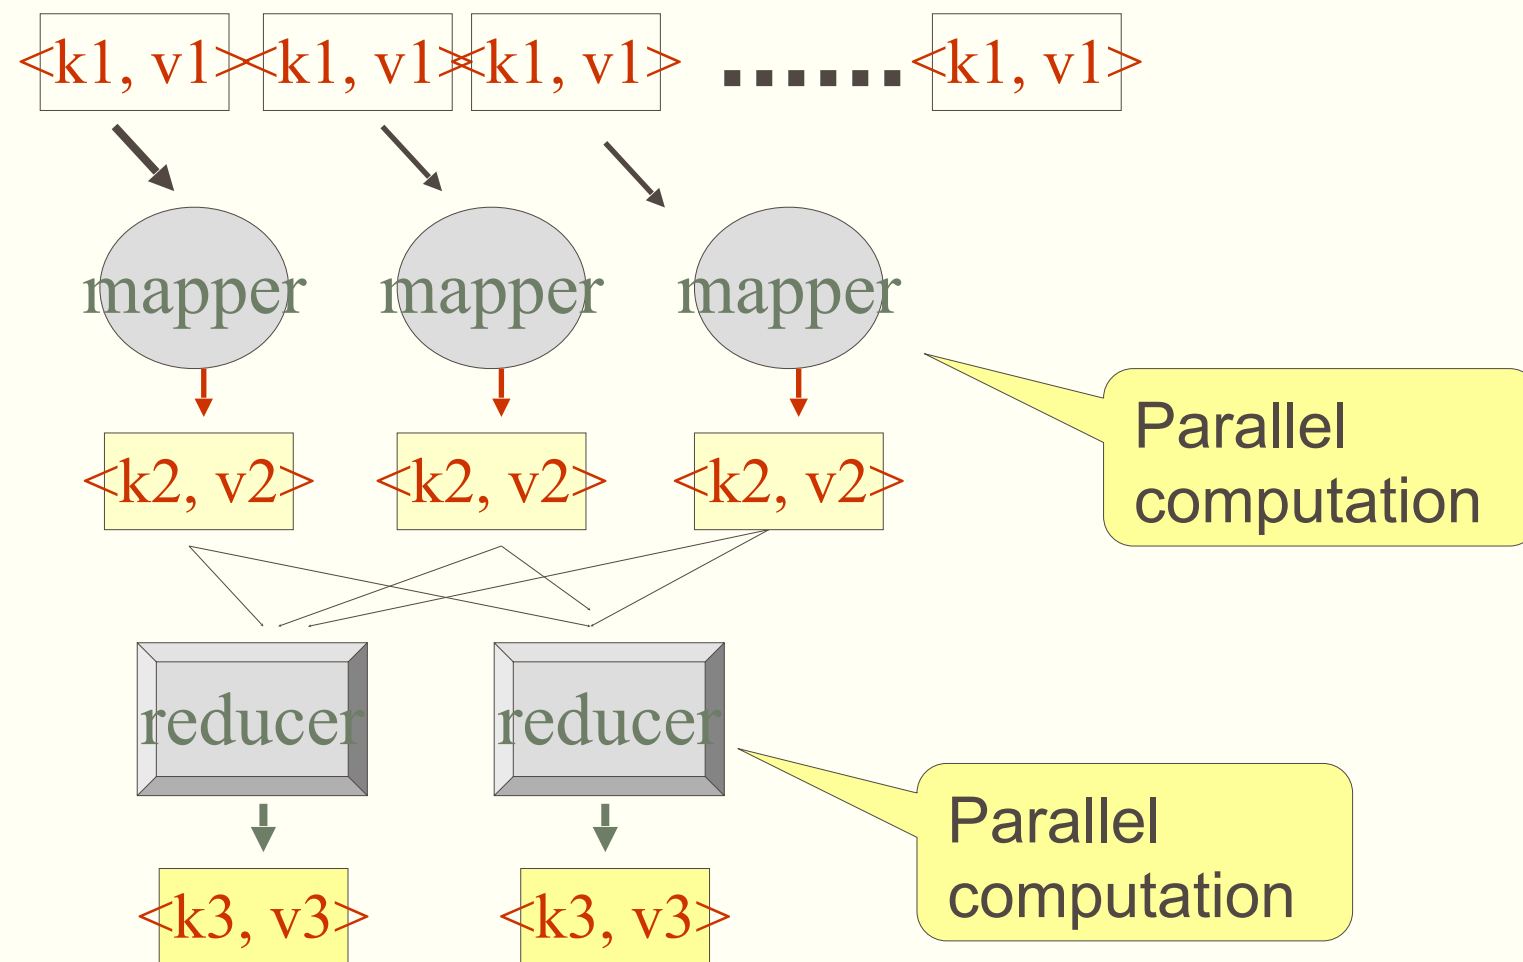
What is “MapReduce”

- A programming model with two primitive functions
 - Map: $\langle k_1, v_1 \rangle \rightarrow \text{list}(k_2, v_2)$
 - Reduce: $\langle k_2, \text{list}(v_2) \rangle \rightarrow \text{list}(k_3, v_3)$
- Input: a list $\langle k_1, v_1 \rangle$ of key-value pairs
- Map:
 - Applied to each pair, computes key-value pairs $\langle k_2, v_2 \rangle$
 - The intermediate key-value pairs are hash-partitioned based on k_2
 - Each partition $(k_2, \text{list}(v_2))$ is sent to a reducer
- Reduce:
 - Takes a partition as input and computes key-value pairs $\langle k_3, v_3 \rangle$
- This process any iterate – multiple map/reduce steps

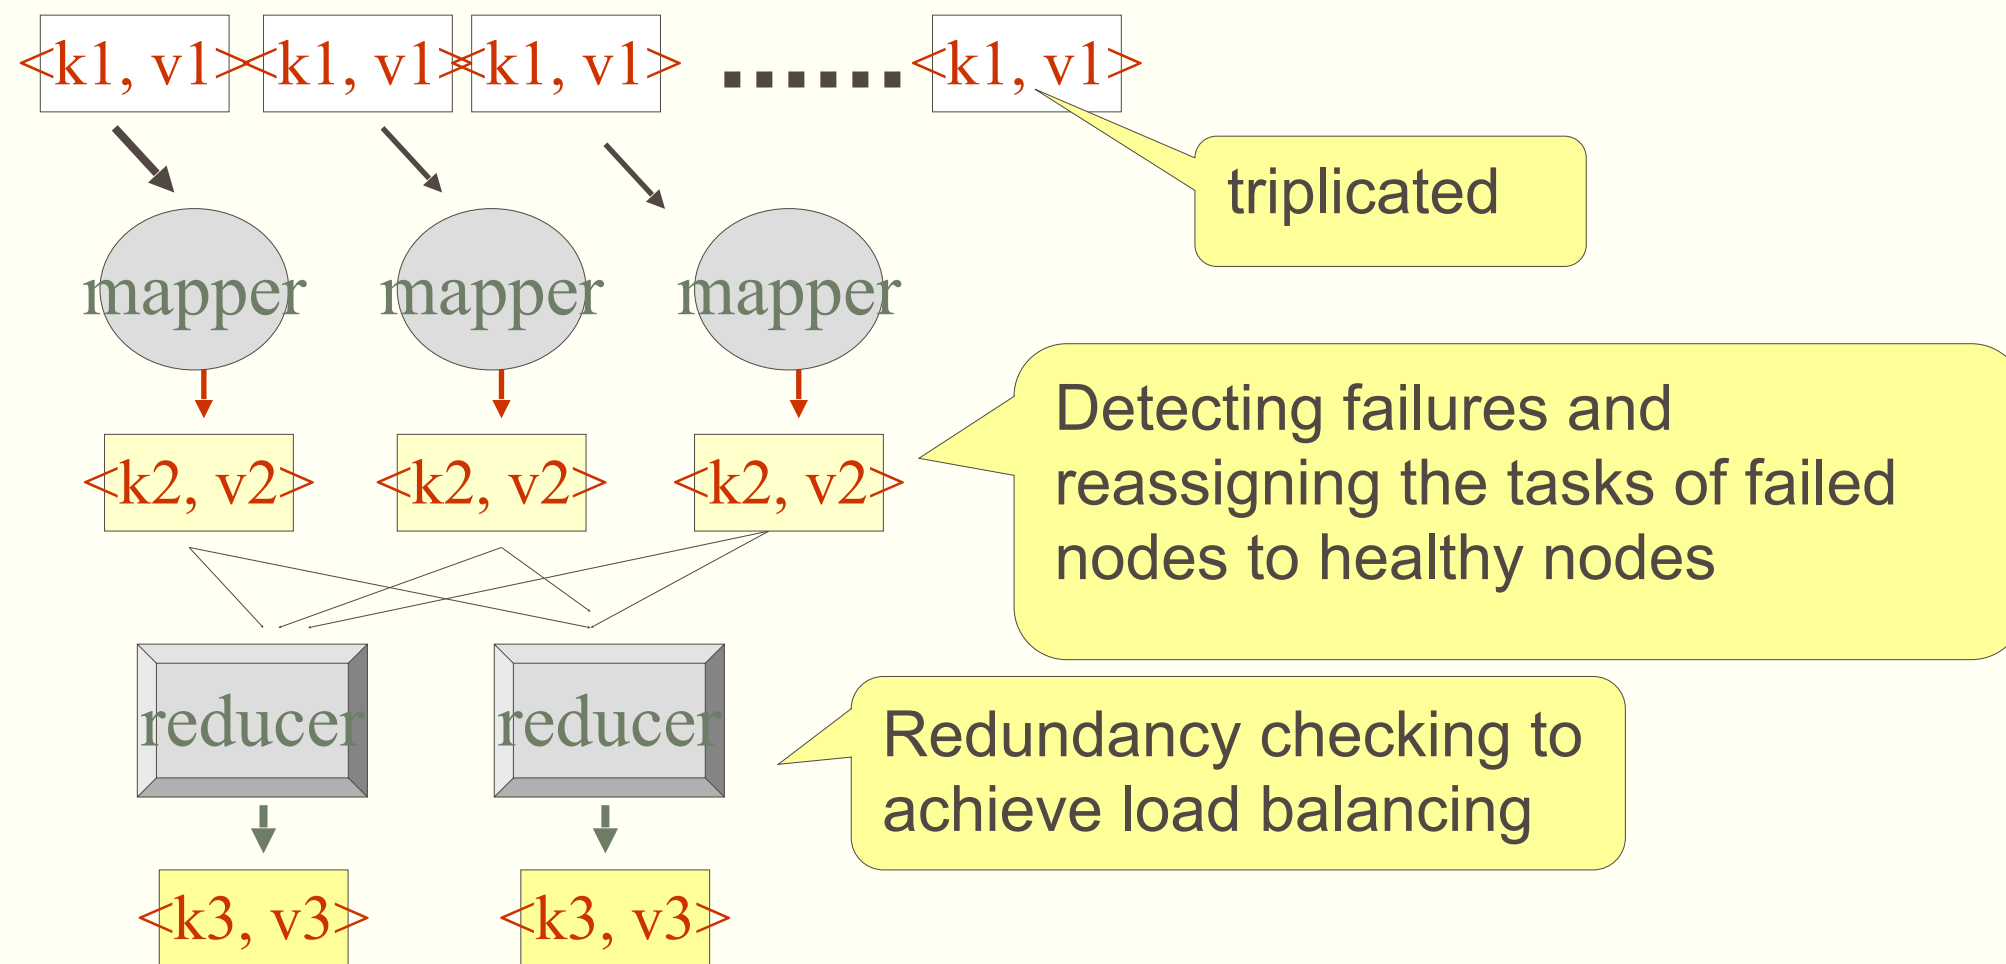
Architecture



Parallelism



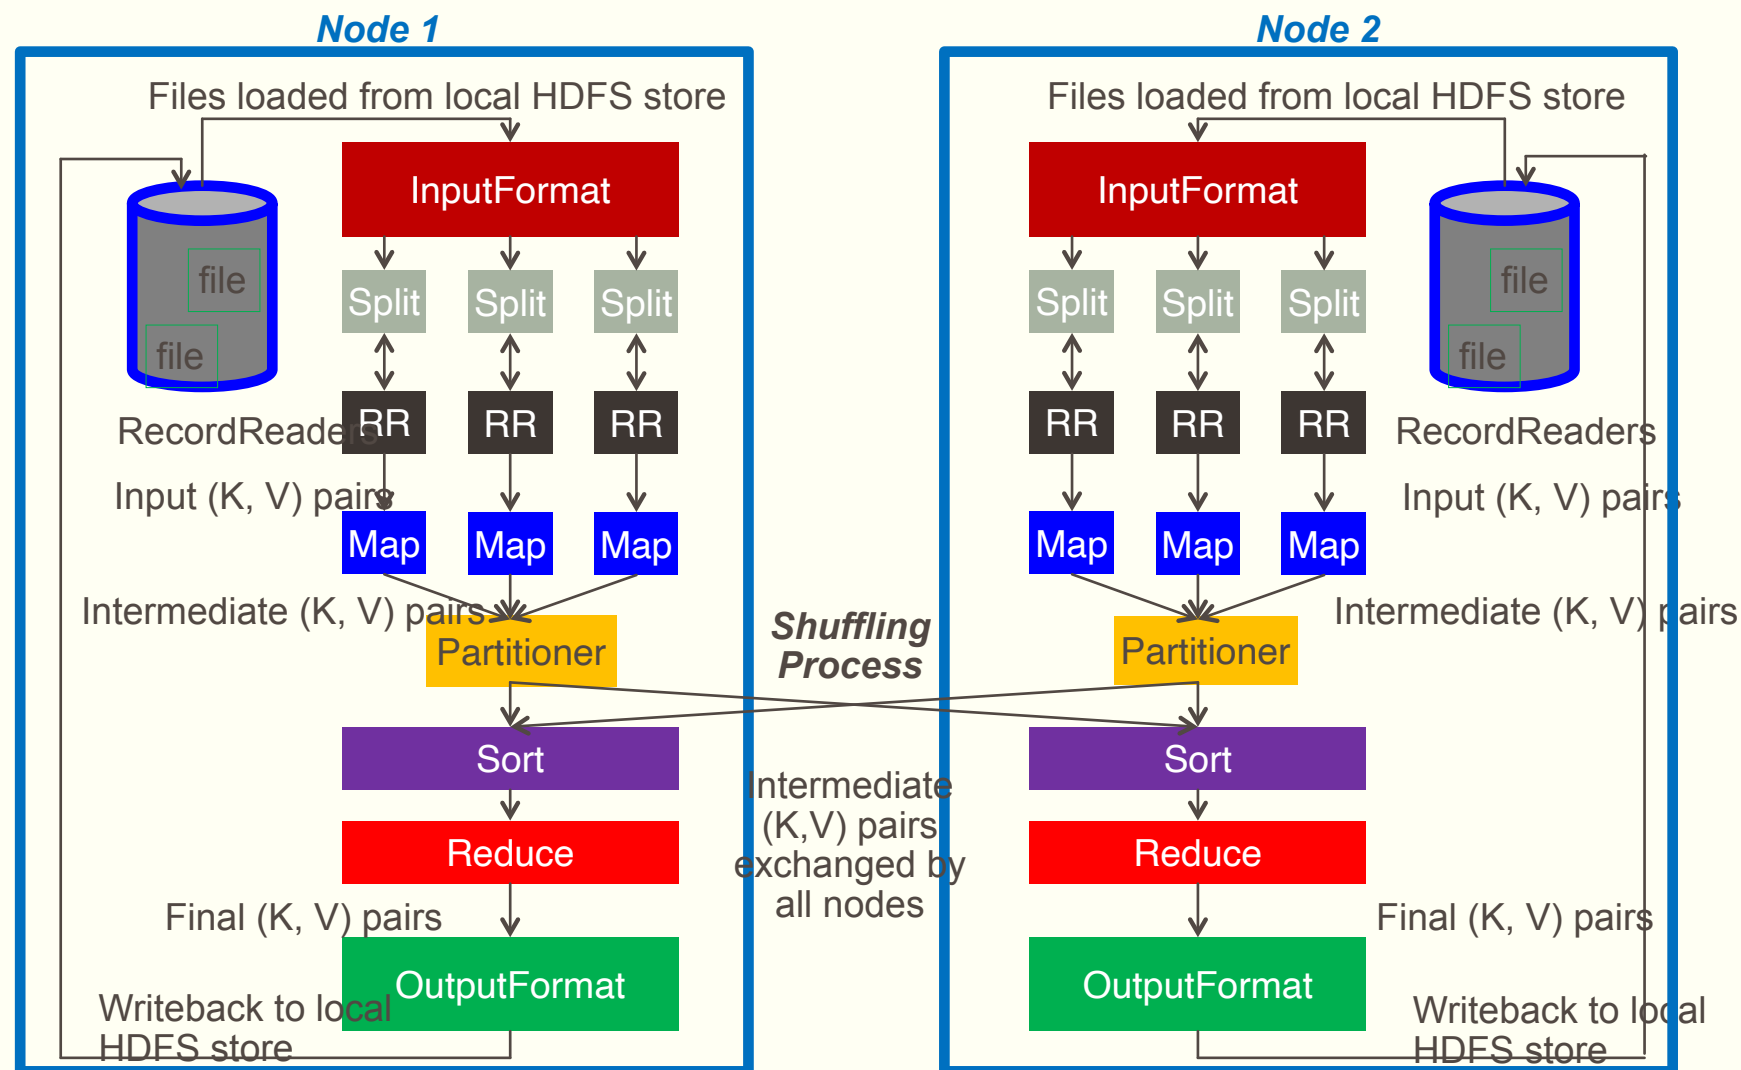
Fault Tolerance



Advantages of MapReduce


- Simple:
 - One only needs to define two functions
 - No need to worry about how the data is stored, distributed and how the operations are scheduled
- Scalability:
 - A large number of low-end machines
 - Scale Out (horizontally): Adding a new computer to a distributed software application; lost-cost “commodity”
 - Scale Up (vertically): Upgrade, add (costly) resources to a single node
- Independence
 - It can work with various storage layers
- Flexibility
 - Independent of data models and schema
- Fault tolerance

Hadoop MapReduce: A Closer Look



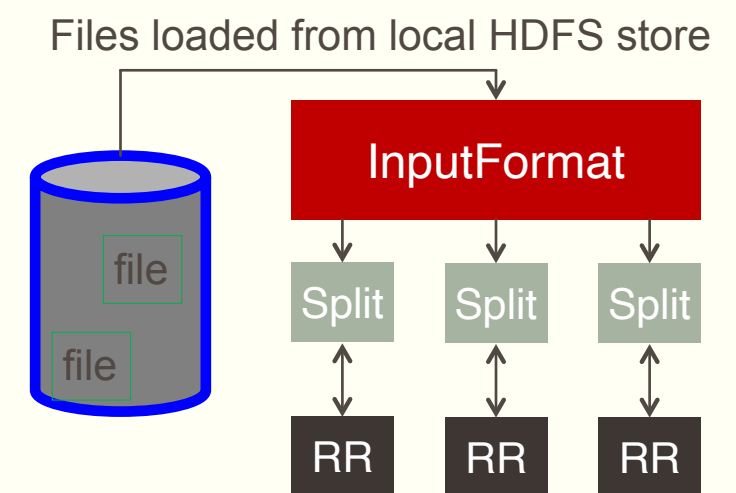
Input Splits

- An input split describes a unit of work that comprises a single map task in a MapReduce Program
- By default, the InputFormat breaks a file up into 64MB splits
- By dividing the file into splits, we allow several map task to operate on a single file in parallel
- If the file is very large, this can improve performance significantly through parallelism
- Each map task corresponds to a single input split



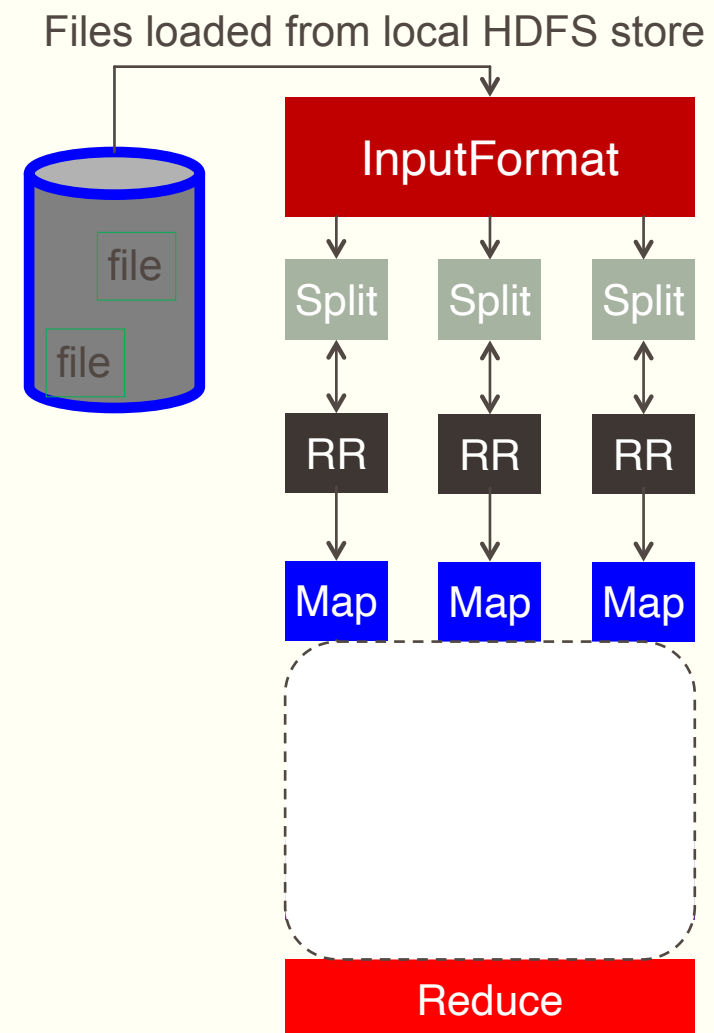
Record Reader

- The input split defines a slice of work but does not describe how to access it
- The **RecordReader** class actually loads data from its source and converts it into (K, V) pairs suitable for reading by Mappers
- The RecordReader is invoked repeatedly on the input until the entire split is consumed
- Each invocation of the RecordReader leads to another call of the map function defined by the programmer



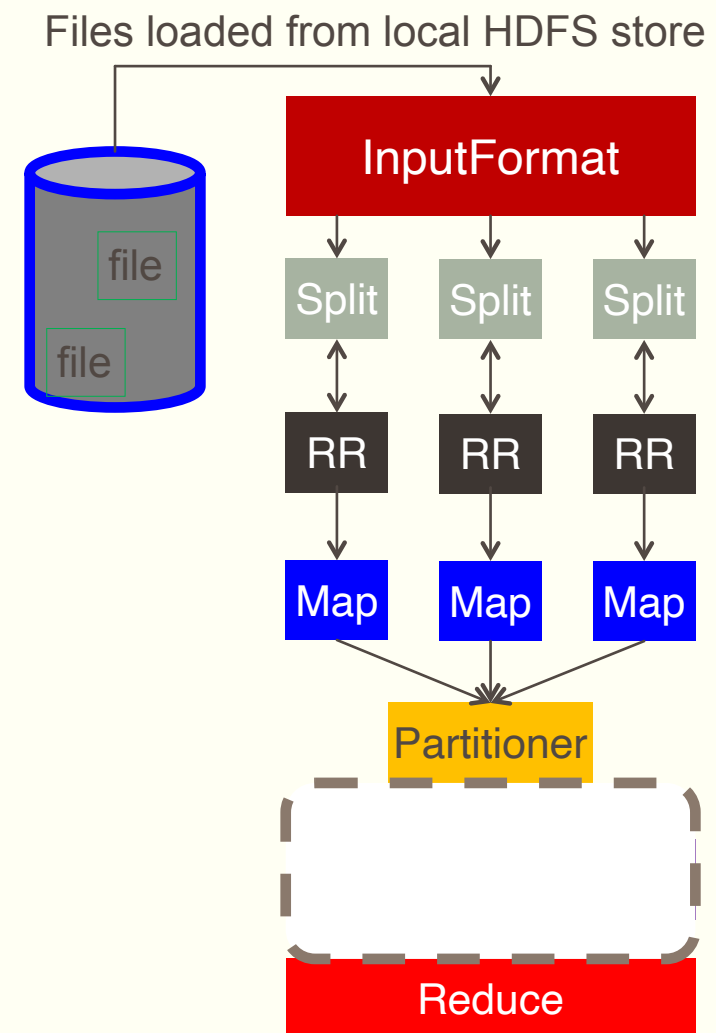
Mapper and Reducer

- The Mapper performs the user-defined work of the first phase of the MapReduce program
- A new instance of Mapper is created for each split
- The Reducer performs the user-defined work of the second phase of the MapReduce program
- A new instance of Reducer is created for each partition
- For each key in the partition assigned to a Reducer, the Reducer is called once



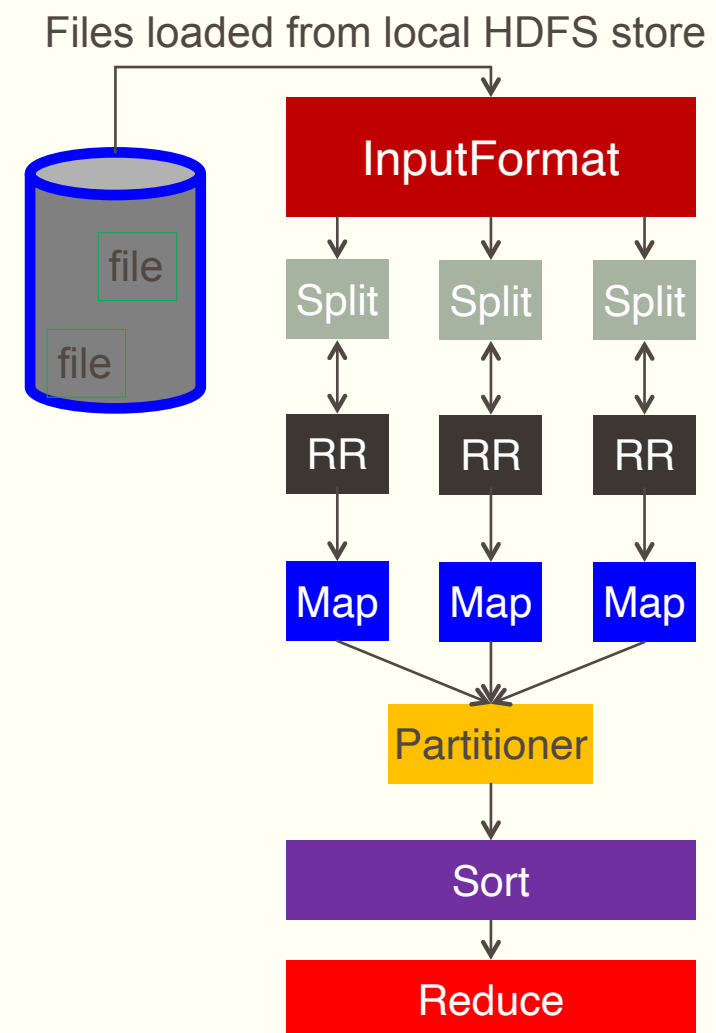
Partitioner

- Each mapper may emit (K, V) pairs to any partition
- Therefore, the map nodes must all agree on where to send different pieces of intermediate data
- The partitioner class determines which partition a given (K,V) pair will go to
- The default partitioner computes a hash value for a given key and assigns it to a partition based on this result



Sort

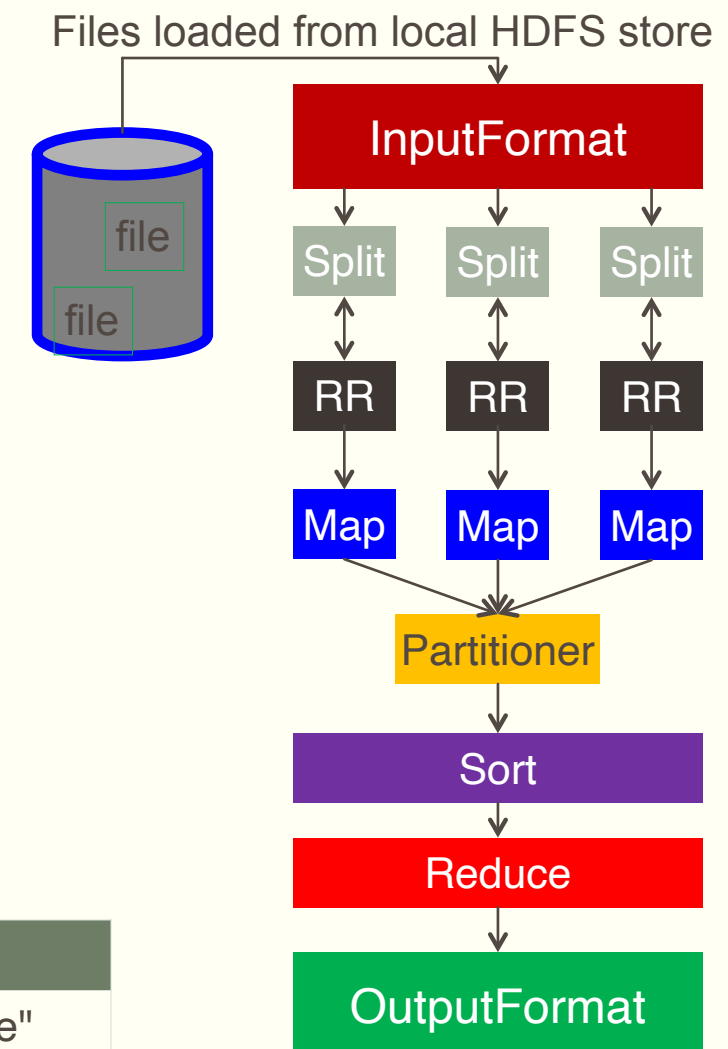
- Each Reducer is responsible for reducing the values associated with (several) intermediate keys
- The set of intermediate keys on a single node is automatically sorted by MapReduce before they are presented to the Reducer



Output Format

- The `OutputFormat` class defines the way (K,V) pairs produced by Reducers are written to output files
- The instances of `OutputFormat` provided by Hadoop write to files on the local disk or in HDFS
- Several `OutputFormats` are provided by Hadoop:

OutputFormat	Description
<code>TextOutputFormat</code>	Default; writes lines in "key \t value" format
<code>SequenceFileOutputFormat</code>	Writes binary files suitable for reading into subsequent MapReduce jobs
<code>NullOutputFormat</code>	Generates no output files

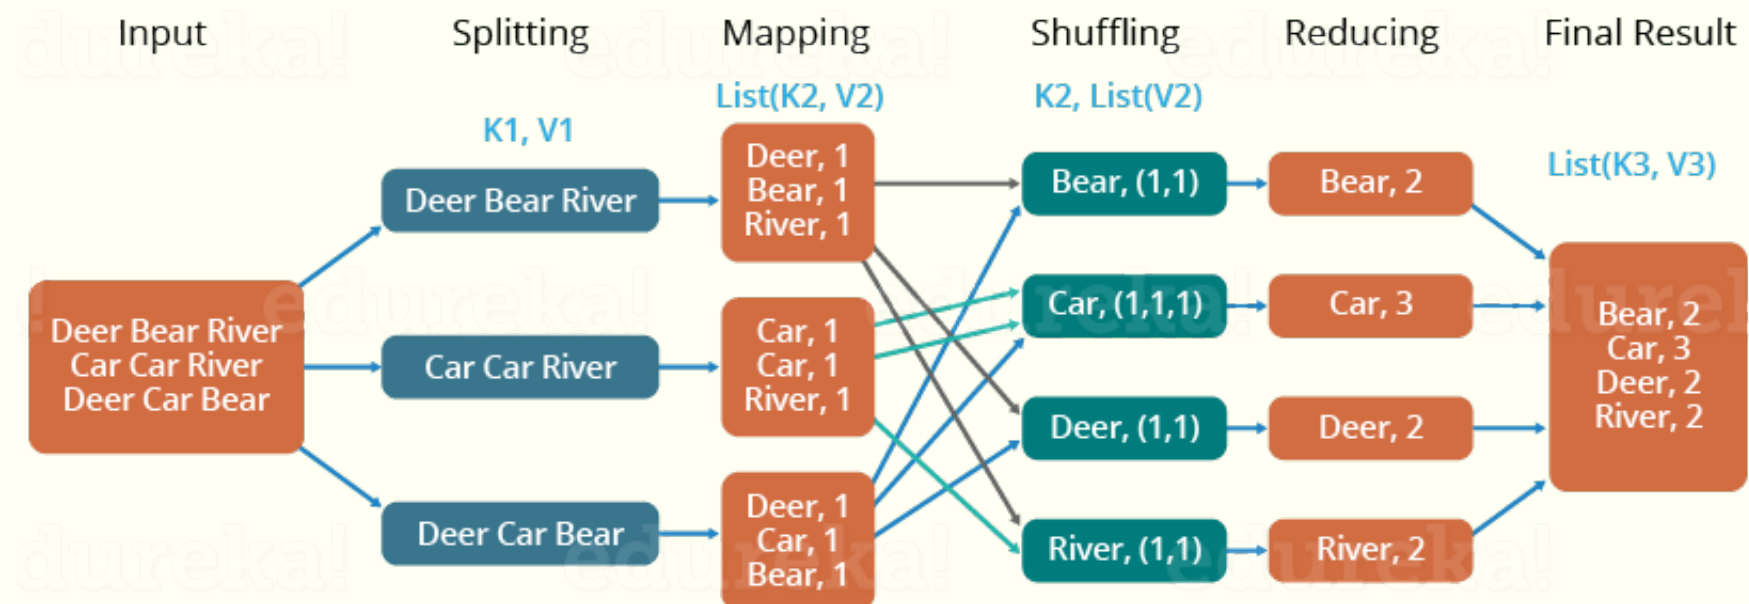


Example: Word Count

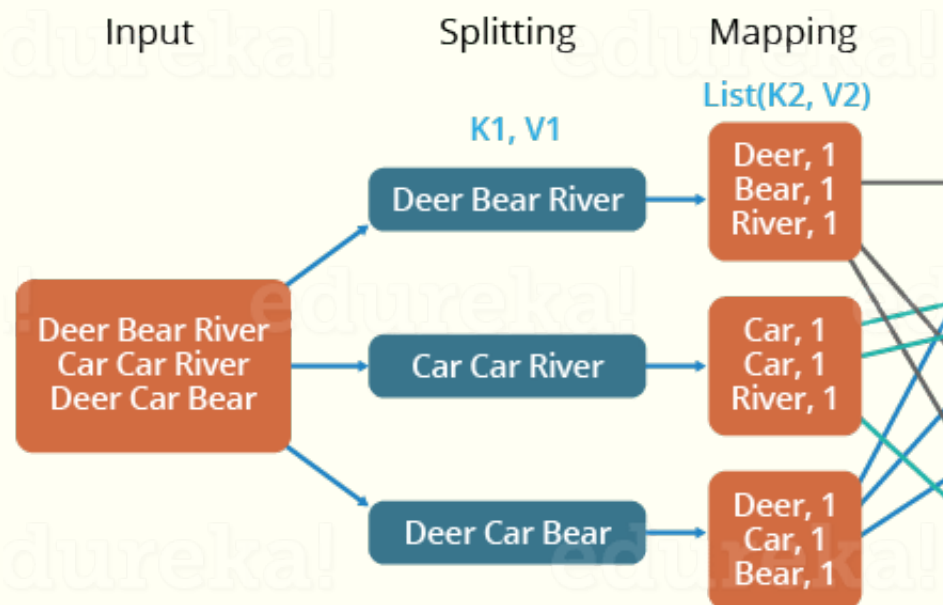
```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:     for all term  $t \in \text{doc } d$  do
4:       EMIT(term  $t$ , count 1)
5:
6: class REDUCER
7:   method REDUCE(term  $t$ , counts  $[c_1, c_2, \dots]$ )
8:      $sum \leftarrow 0$ 
9:     for all count  $c \in \text{counts } [c_1, c_2, \dots]$  do
10:       $sum \leftarrow sum + c$ 
11:     EMIT(term  $t$ , count  $s$ )
```

The Overall MapReduce Word Count Process

edureka!



Java Implementation: Word Count



```
import org.apache.hadoop.mapreduce.Mapper;
public class WordCountMapper extends Mapper{

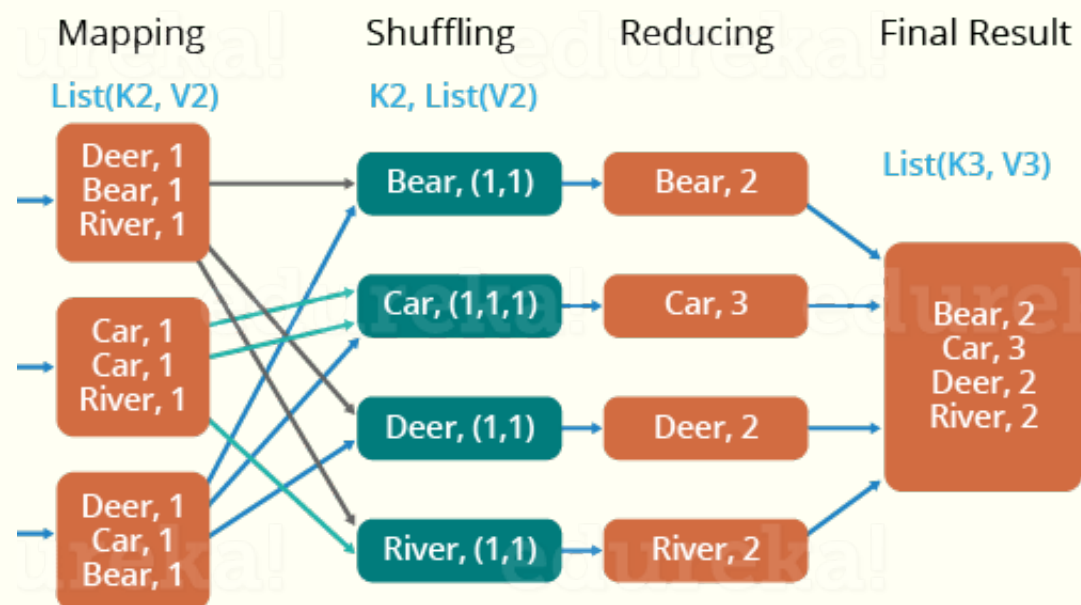
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    @Override
    protected void map(LongWritable key, Text value,
        Context context)
        throws IOException, InterruptedException {

        //Get the text and tokenize the word using space as separator.
        String line = value.toString();
        StringTokenizer st = new StringTokenizer(line, " ");

        //For each token aka word, write a key value pair with
        //word and 1 as value to context
        while(st.hasMoreTokens()){
            word.set(st.nextToken());
            context.write(word, one);
        }
    }
}
```


WordCount: Reducer



```
import org.apache.hadoop.mapreduce.Reducer;  
public class WordCountReducer extends Reducer{
```

```
@Override
```

```
protected void reduce(Text key, Iterable values,  
                        Context context)
```

```
throws IOException, InterruptedException {
```

```
    int sum = 0;
```

```
    Iterator valuesIt = values.iterator();
```

```
    //
```

```
    // For each key value pair, get the value and adds to the sum
```

```
    // to get the total occurrences of a word
```

```
    while(valuesIt.hasNext()){
```

```
        sum = sum + valuesIt.next().get();
```

```
    }
```

```
    // Writes the word and total occurrences as
```

```
    // key-value pair to the context
```

```
    context.write(key, new IntWritable(sum));
```

```
}
```

```
}
```

WordCount: Driver

```
import org.apache.hadoop.conf.Configured;
public class WordCount extends Configured implements Tool{

    public static void main(String[] args) throws Exception{
        int exitCode = ToolRunner.run(new WordCount(), args);
        System.exit(exitCode);
    }

    public int run(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.printf("Usage: %s needs two arguments, input
and output
files\n", getClass().getSimpleName());
            return -1;
        }

        // Create a new Jar and set the driver class(this class) as the
        // main class of jar
        Job job = new Job();
        job.setJarByClass(WordCount.class);
        job.setJobName("WordCounter");

        // Set the input and the output path from the arguments
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

```
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        job.setOutputFormatClass(TextOutputFormat.class);

        //Set the map and reduce classes in the job
        job.setMapperClass(WordCountMapper.class);
        job.setReducerClass(WordCountReducer.class);

        //Run the job and wait for its completion
        int returnValue = job.waitForCompletion(true) ? 0:1;

        if(job.isSuccessful()) {
            System.out.println("Job was successful");
        } else if(!job.isSuccessful()) {
            System.out.println("Job was not successful");
        }

        return returnValue;
    }
}
```

```
$ bin/hadoop jar wordcount.jar WordCount input output
```

MapReduce Implementation of Relational Operators

- Projection
- Selection
- Union
- Set Difference
- Join
 - Reduce-side Join
 - Map-side Join
 - In-memory Join
- Aggregation

MapReduce: Projection

- Projection $\pi_A R$
- Input: for each tuple t in R , a pair (key, value), where value = t

not necessarily a key of R

- Map(key, t)

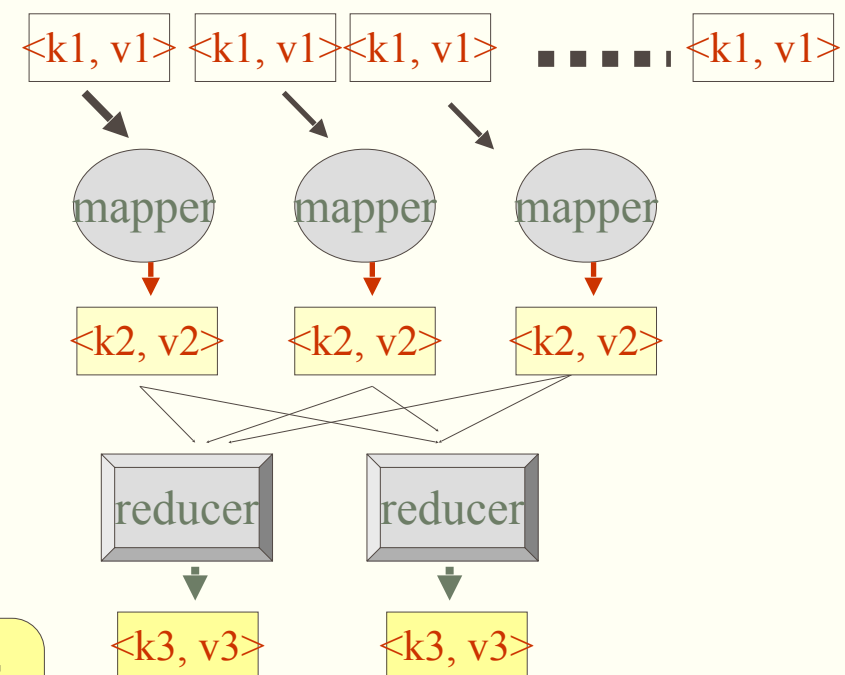
- Emit ($t.A$, $t.A$)

Apply to each input tuple, in parallel; emit new tuples with projected attributes

- Reduce (hkey, hvalue[])

- Emit(hkey, hkey)

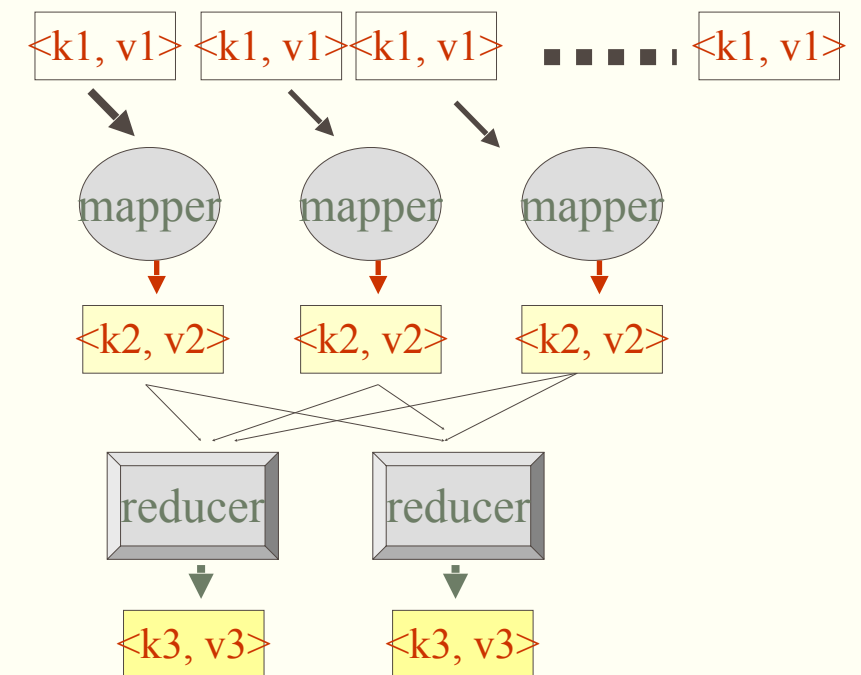
the reducer is not necessary; but it eliminates duplicates.



MapReduce: Selection

- Selection $\sigma_C R$
- Input: for each tuple t in R , a pair (key, value) where value = t
- Map (key, t)
 - If $C(t)$, Then emit (t , “1”)
- Reduce (hkey, hvalue[])
 - emit(hkey, hkey)

Apply to each input tuple, in parallel;
select tuples that satisfy condition C



MapReduce: Union

- Union $R_1 \cup R_2$

A mapper is assigned chunks from either R1 or R2

- Input: for each tuple t in R_1 and s in R_2 , a pair (key, value)

- Map (key, t)

- Emit (t , "1")

A mapper just passes an input tuple to a reducer

- Reduce (hkey, hvalue[])

- emit(hkey, hkey)

Reducers simply eliminate duplicates

MapReduce: Set Difference

- Set difference $R_1 - R_2$
- Input, for each tuple t in R_1 and s in R_2 , a pair (key, value)
- Map (key, t)
 - If t is in R_1 , then emit (t , "1"), else emit (t , "2")
- Reduce (hkey, hvalue[])
 - If only "1" appears in the list hvalue, then emit (hkey, hkey)

distinguishable

tag each tuple with its source

Reducers do the checking

Join Algorithms in MapReduce

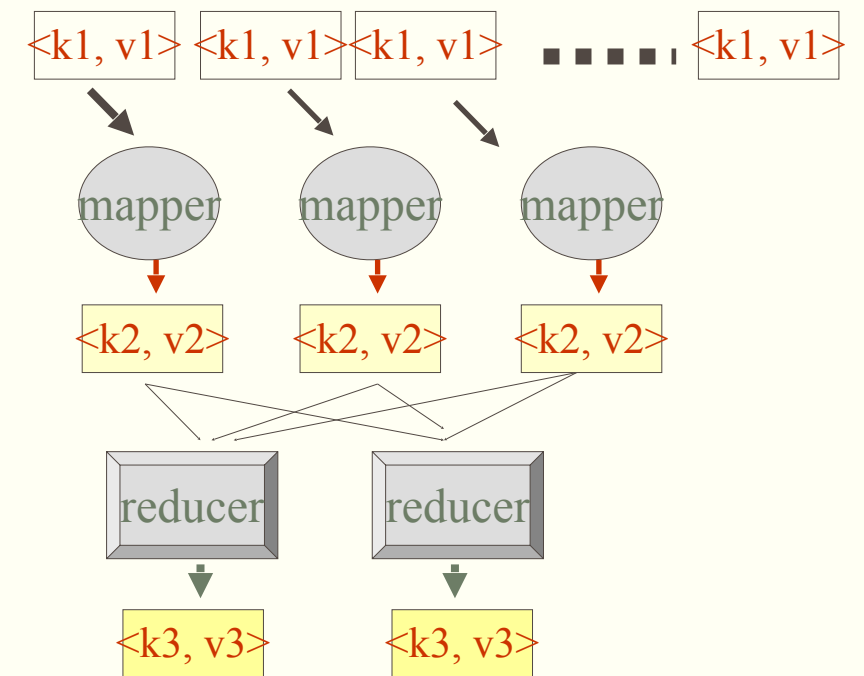
- Reduce-side join
- Map-side join
- In-memory join
 - Striped variant
 - Memcached variant

Reduce-side Join

- Natural Join: $R_1 \bowtie_{R_1.A=R_2.B} R_2$, where $R_1[A, C]$, $R_2[B, D]$
- Input: for each tuple t in R_1 and s in R_2 , a pair (key, value)
- Map (key, t)
 - If t is in R_1
 - then emit ($t.[A]$, (“1”, $t.[C]$))
 - Else emit ($t.[B]$, (“2”, $t.[D]$))
- Reduce (hkey, hvalue[])
 - For each (“1”, $t.[C]$) and each (“2”, $s.[D]$) in the list hvalue[]
 - Emit ((hkey, $t.[C]$, $s.[D]$), hkey)

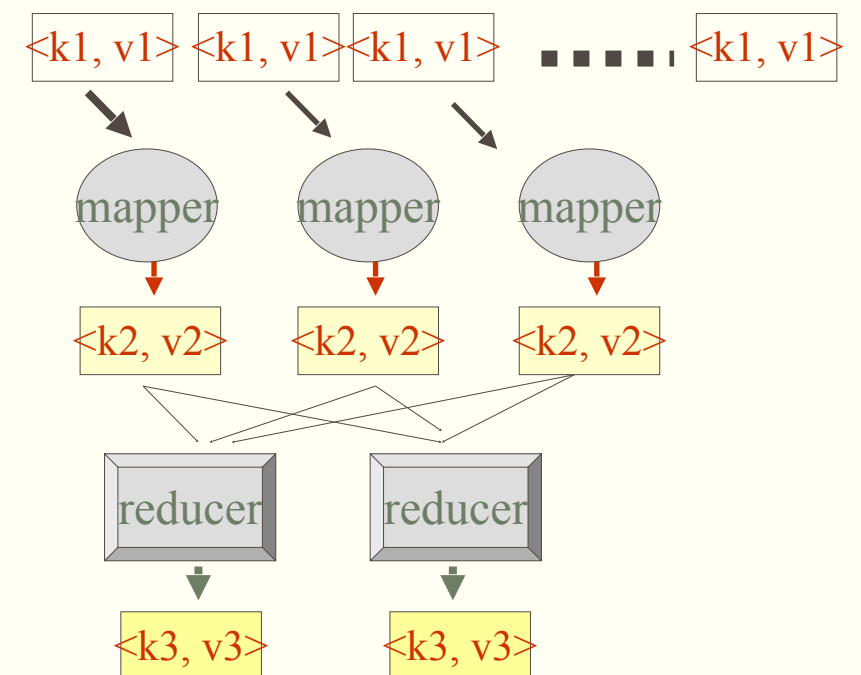
Hashing on join attributes

Nested loop




Map-Side Join

- Recall $R_1 \bowtie_{R_1.A=R_2.B} R_2$
 - Partition R_1 and R_2 into n partitions, by the same partitioning function (range/hash)
 - Compute $R_1^i \bowtie_{R_1.A=R_2.B} R_2^i$ locally
 - Merge the local results
- Map-side Join
 - Input relations are partitioned and sorted based on join keys
 - Map over R_1 and read from the corresponding partition of R_2
- Map(key, t)
 - Read R_2^i
 - For each tuple s in relation R_2^i
 - If $t[A] = s[B]$ then emit $((t[A], t[C], s[D]), t[A])$
- Reduce(hkey, hvalue[])
 - Emit (hkey, hkey)



In-Memory Join (Broadcast Join)

- Recall $R_1 \bowtie_{R_1.A < R_2.B} R_2$
 - Partition R1 into n partitions, by the same partitioning function (range/hash)
 - Replicate the other relation R2
 - Compute $R_1^i \bowtie_{R_1.A < R_2.B} R_2$ locally at each processor i
 - Merge the local results
- Broadcast Join
 - A smaller relation is broadcast to each node and stored in its local memory
 - The other relation is partitioned and distributed across mappers
- Map(key, t)
 - Read R_2
 - For each tuple s in relation R_2
 - If $t[A] < s[B]$ then emit $((t[A], t[C], s[D]), t[A])$
- Reduce(hkey, hvalue[])
 - Emit (hkey, hkey)



MapReduce: Aggregation

- $R(A, B, C)$, compute $\text{sum}(B)$ group by A

- Map (key, t)
 - Emit ($t[A]$, $t[B]$)

Grouping: done by MapReduce framework

- Reduce (hkey, hvalue[])
 - Sum = 0
 - For each value s in the list hvalue[]
 - Sum = Sum + 1
 - Emit (hkey, Sum)