



MapReduce II

MapReduce for Graph Queries

- Input: Query Q and graph G
- Output: answers $Q(G)$ to Q in G
- Map (key: $node$, value: ($adjacency-list$, $others$))
 - Computation
 - Emit ($mkey$, $mvalue$)
- Reduce(key: $mkey$, value: list[$mvalue$])
 - Computation
 - Emit ($rkey$, $rvalue$)

Match $rkey$, $rvalue$ when multiple iterations of MapReduce are needed

Match $mkey$, $mvalue$

Dijkstra's Algorithm for Distance Query

- Distance: single-source shortest-path problem
 - Input: A directed weighted graph G and a node s in G
 - Output: The length of shortest paths from s to all nodes in G

- Dijkstra (G, s, w)

- For all nodes $v \in V$ do

- $d[v] \leftarrow \infty$

- $d[s] \leftarrow 0; Q \leftarrow V$

Use a priority queue Q ;
 $w(u, v)$: weight of edge (u, v) ;
 $d(u)$: the distance from s to u

- While Q is non-empty, do

- $u \leftarrow \text{ExtractMin}(Q)$

Extract one with the minimum $d(u)$

- For all nodes $v \in \text{adj}(u)$ do


- $d[v] > d[u] + w(u, v)$ then $d[v] \leftarrow d[u] + w(u, v)$

- Complexity:

- $O(|V| \log |V| + |E|)$

Finding the Shortest Path

- Consider simple case of equal edge weights: solution to the problem can be defined inductively
- Intuition:
 - Define: b is reachable from a if b is on adjacency list of a
 - $d[s] = 0$
 - For all nodes p reachable from s , $d[p] = 1$
 - For all nodes n reachable from some other set of nodes M ,
$$d[n] = 1 + \min_{m \in M} d[m]$$



Shortest Path: From Intuition to Algorithm

- Input: Graph G , represented by adjacency lists
- Key: node ID n
- Value: node value N
 - $N.distance$: from start node s to n
 - $N.adjList$: $\left[(m, w(n, m)) \right]$, node id and weight of edge (n, m)
- Initialization:
 - For all n , $N.distance = \infty$

Shortest Path: From Intuition to Algorithm

- Mapper

- $\forall m \in N. AdjList: \text{Emit } (m, d + w(n, m))$

- Sort/Shuffle

- Groups distances by reachable nodes

- Reducer:

- Selects minimum distance path for each reachable node
 - Additional book-keeping needed to keep track of actual path

Shortest Path: Mapper

- Map (nid n , nvalue N)
 - $d \leftarrow N.distance$
 - $emit(n, N)$
 - For each $(m, w) \in N.AdjList$
 - $emit(m, d + w(n, m))$
- Parallel Processing
 - All nodes are processed in parallel, each by a mapper
 - $emit(n, N)$ preserve graph structure for iterative processing
 - For each node m adjacent to n , emit a revised distance via n


Shortest Path: Reducer

- Reduce (nid m , list[d_1, d_2, \dots])
 - $d_{min} \leftarrow \infty$
 - $\forall d \in [d_1, d_2, \dots]$
 - If d is Node
 - $M \leftarrow d$
 - Else
 - If $d < d_{min}$ then $d_{min} \leftarrow d$
 - $M.distance \leftarrow d_{min}$
 - $emit(m, M)$
- list for m :
 - distances from all predecessors so far
 - Node value M : must exist (from Mapper)

Iteration and Termination

- Each MapReduce iteration advances the “known frontier” by one hop
 - Subsequent iteration include more and more reachable nodes as frontier expands
 - Multiple iterations are needed to explore entire graph
- Termination: when the intermediate result no longer changes
 - Controlled by a **non-MapReduce Driver**
 - Use a flag – inspected by a **non-MapReduce Driver**

Visualizing Parallel BFS



Efficiency

- MapReduce explores all path in parallel
- Each MapReduce iteration advances the “unknown frontier” by one hop
 - Redundant work, since useful is only done at the “frontier”
- Dijkstra’s algorithm can be more efficient
 - At any step, it only pursues edges from the minimum-cost path inside the frontier

MapReduce: A Closer Look

- Data partitioned parallelism
 - Local computation at each node in mapper, in parallel:
 - Attributes of the node, adjacent edges and local link structure
 - Propagating Computations
 - Transversing the graph, this may involve iterative MapReduce
- Tips
 - Adjacency lists
 - Local computation in mapper
 - Pass along partial results via outlinks, keyed by destination node
 - Perform aggregation in reducer on inlinks to a node
 - Iterate until convergence: controlled by external “driver”
 - Pass graph structures between iterations
- Need a way to test for convergence!

MapReduce: PageRank

- The likelihood that a page v is visited by a random walk:

$$\alpha \left(\frac{1}{|V|} \right) + (1 - \alpha) \sum_{u \in L(v)} \frac{P(u)}{C(u)}$$

- Recursive computation:
 - For each page $v \in G$
 - Compute $P(v)$ using $P(u)$ for all $u \in L(v)$
 - Until
 - Converge: no changes to any $P(v)$
 - After a fixed number of iteration

PageRank: MapReduce Algorithm

- Input: Graph G , represented by adjacency lists
- Key: Node ID n
- Value: node value N :
 - Rank: the page rank of a node
 - AdjList: Adjacency list
- Simplified Version:

$$(1 - \alpha) \sum_{u \in L(v)} \frac{P(u)}{C(u)}$$



PageRank: Mapper

- Map (nid n , nvalue N)
 - $p \leftarrow N.rank / |N.AdjList|$
 - $emit(n, N)$
 - For each $(m, w) \in N.AdjList$
 - $emit(m, p)$
- Parallel Processing
 - All nodes are processed in parallel, each by a mapper
 - For each node m adjacent to n , emit PageRank contribution from n
 - $emit(n, N)$ preserve graph structure for iterative processing



PageRank: Reducer

- Reduce (nid m , list[p_1, p_2, \dots])
 - $s \leftarrow 0$
 - $\forall p \in [p_1, p_2, \dots]$
 - If p is Node
 - $M \leftarrow p$
 - Else
 - $s \leftarrow s + p$
 - $M.rank \leftarrow p$
 - $emit(m, M)$

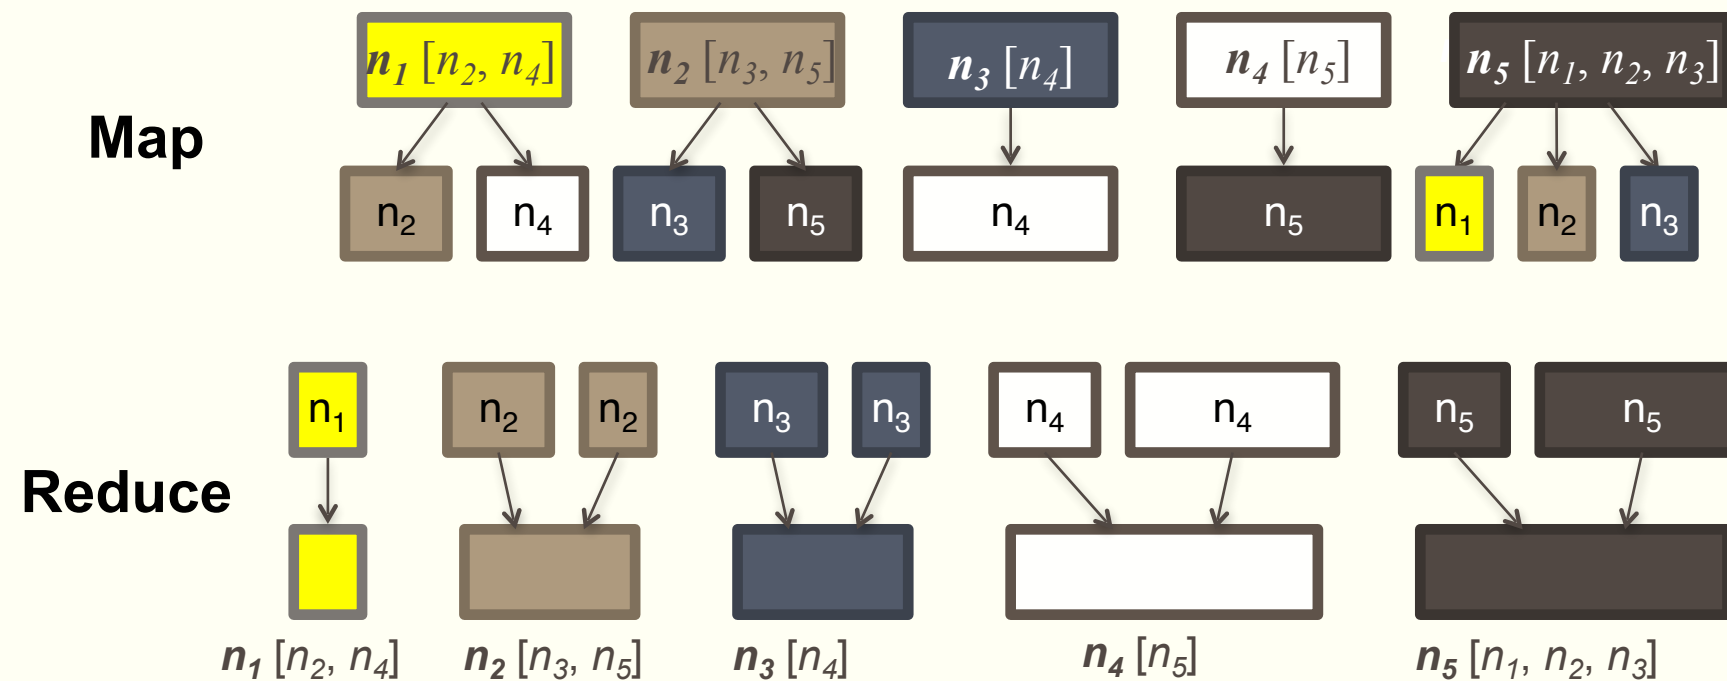
Sample PageRank: Iteration 1



Sample PageRank: Iteration 2



PageRank in MapReduce



The need for parallel models beyond MapReduce

- Inefficiency:
 - Blocking
 - Intermediate result shipping (all to all)
 - Disk I/O in each step, even for invariant data in a loop
- Does not support iterative graph computation
 - Need for external driver
 - No mechanism to support global data structures that can be accessed and updated by all mappers and reducers
- Support for incremental Computation?
- Have to recast algorithm in MapReduce
- General model, not limited to graphs