




# In Memory DBMS



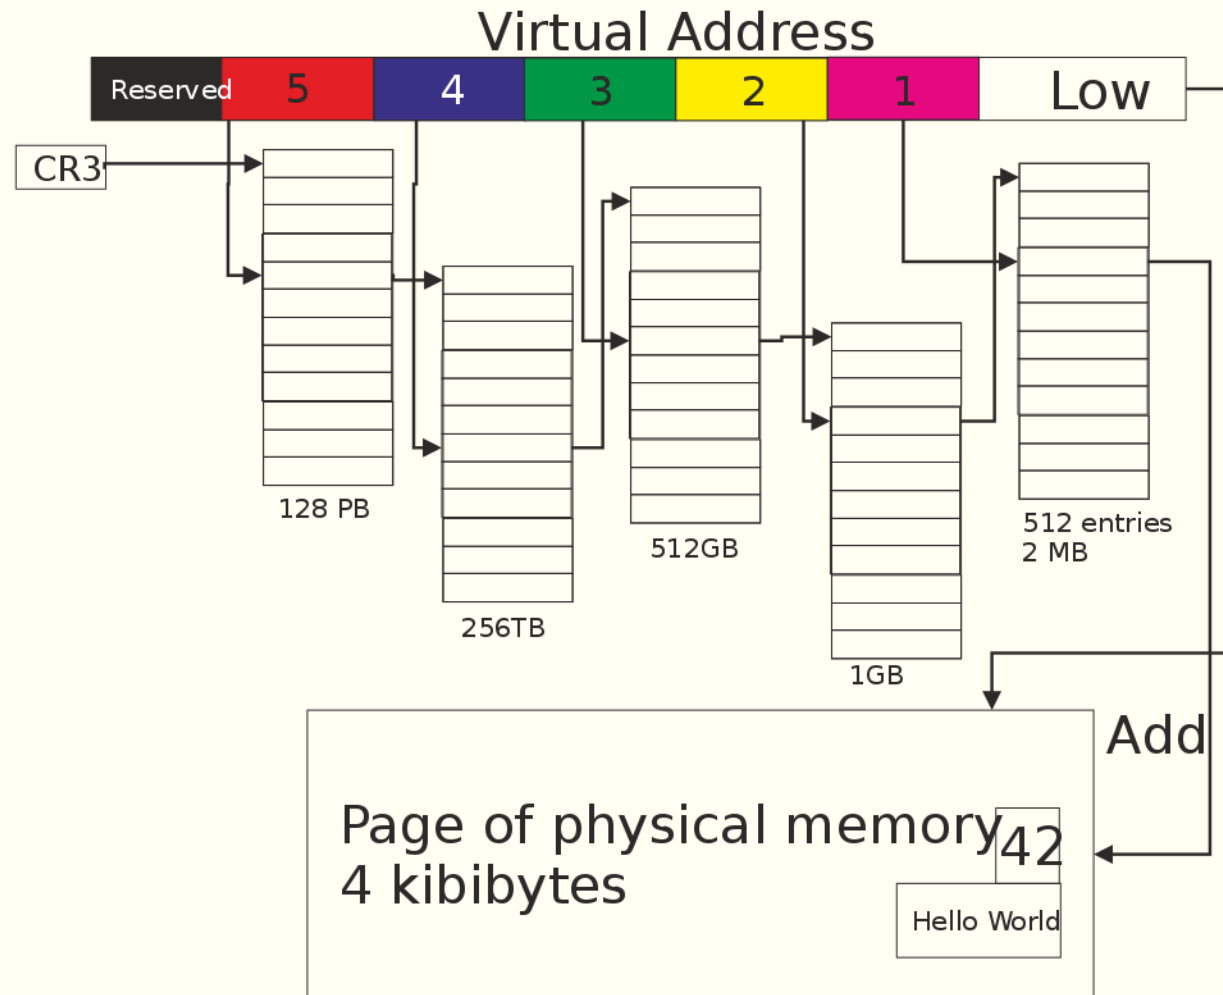
# Computer Architecture

---



# How Memory is Managed?

---



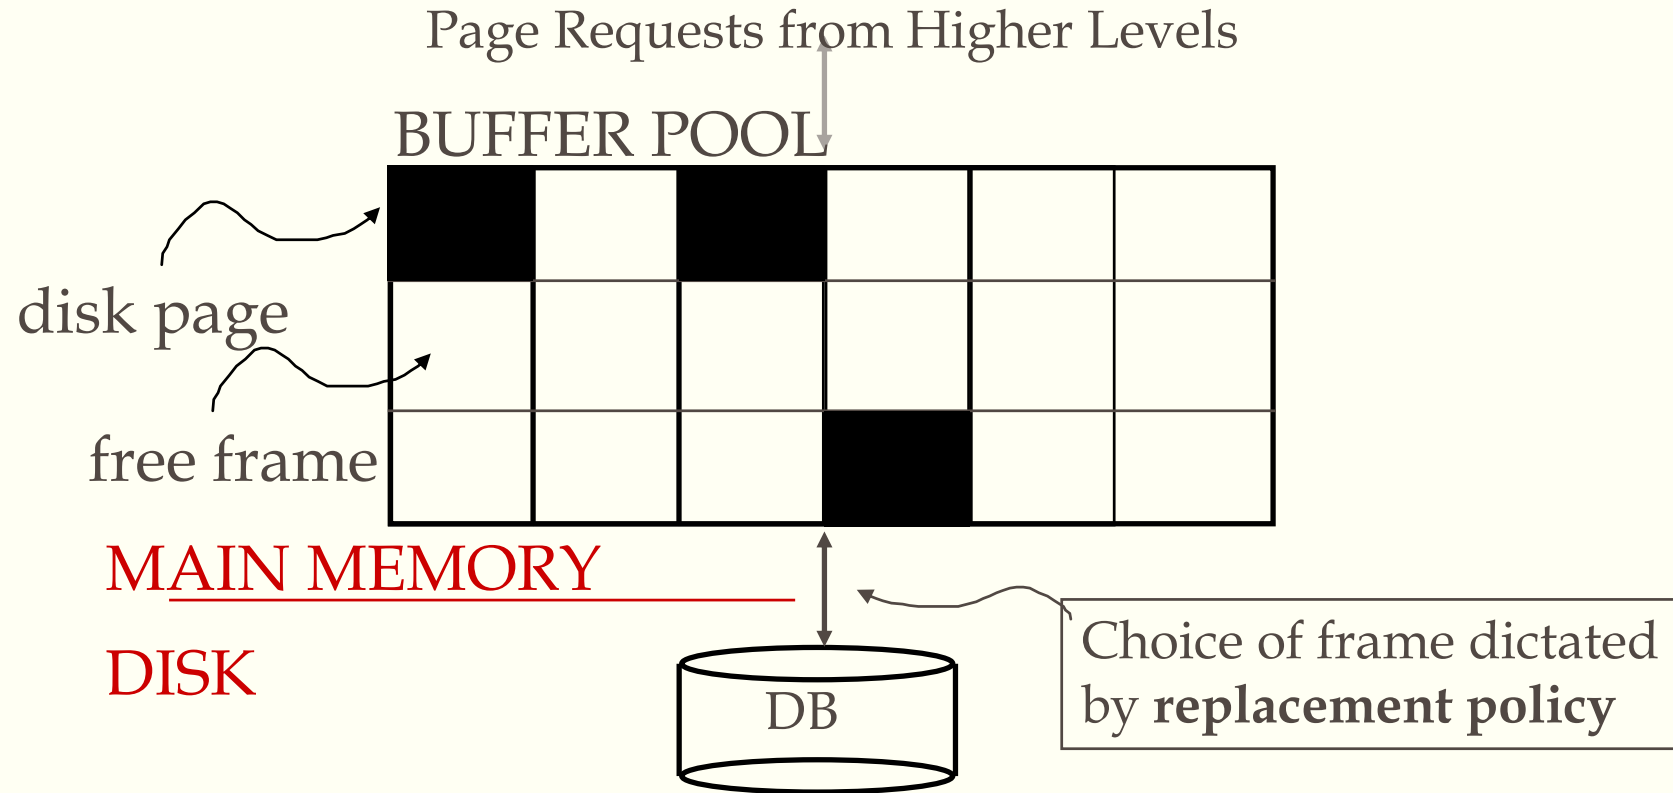
# Disk Oriented DBMS

---

- The primary storage location of the database is on non-volatile storage (e.g., HDD, SSD).
  - → The database is organized as a set of fixed-length blocks called slotted pages.
- The system uses an in-memory buffer pool to cache blocks fetched from disk.
  - → Its job is to manage the movement of those blocks back and forth between disk and memory.

# Page request

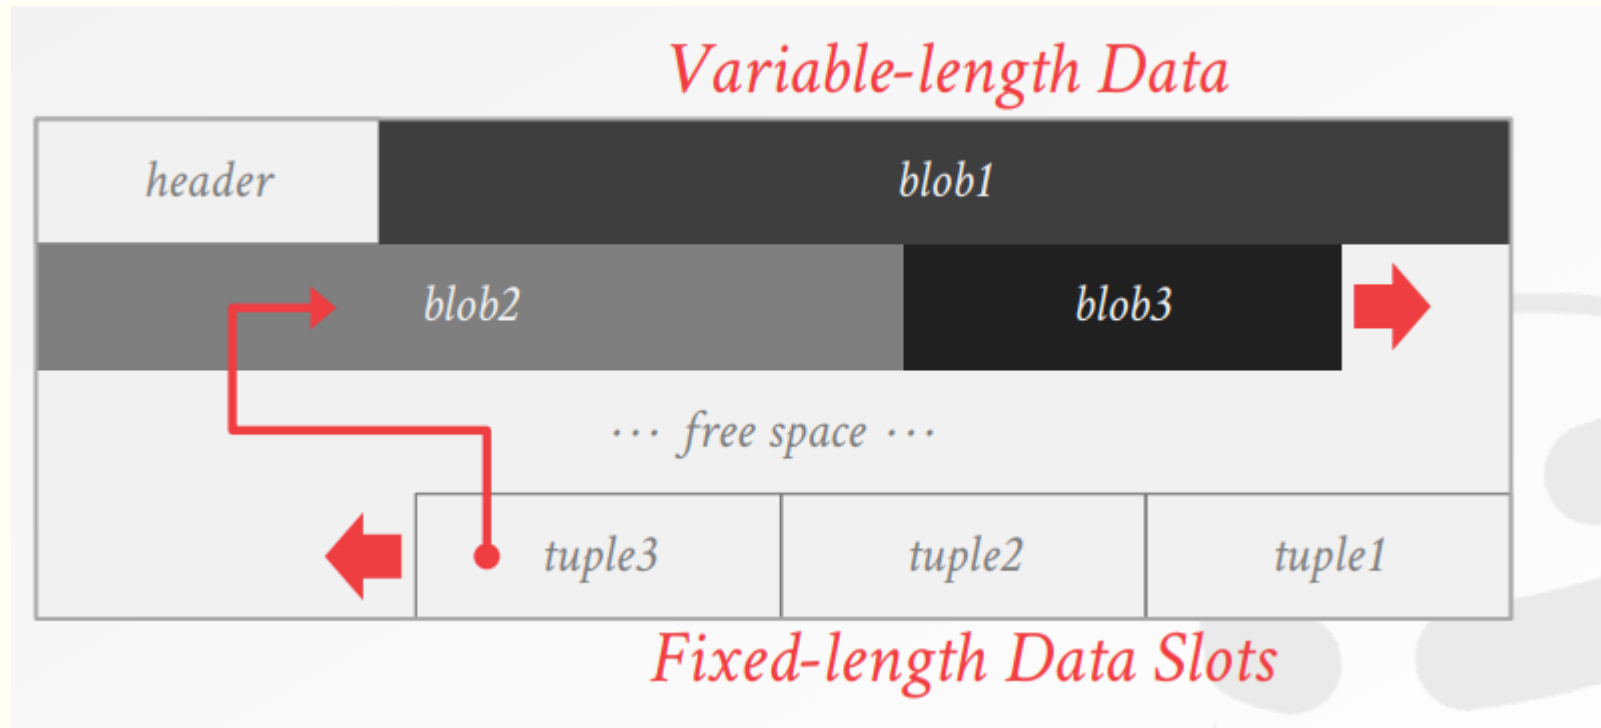
---



- Data must be in RAM for DBMS to operate on it!
- Table of <frame#, pageId> pairs is maintained.

# A Page

---



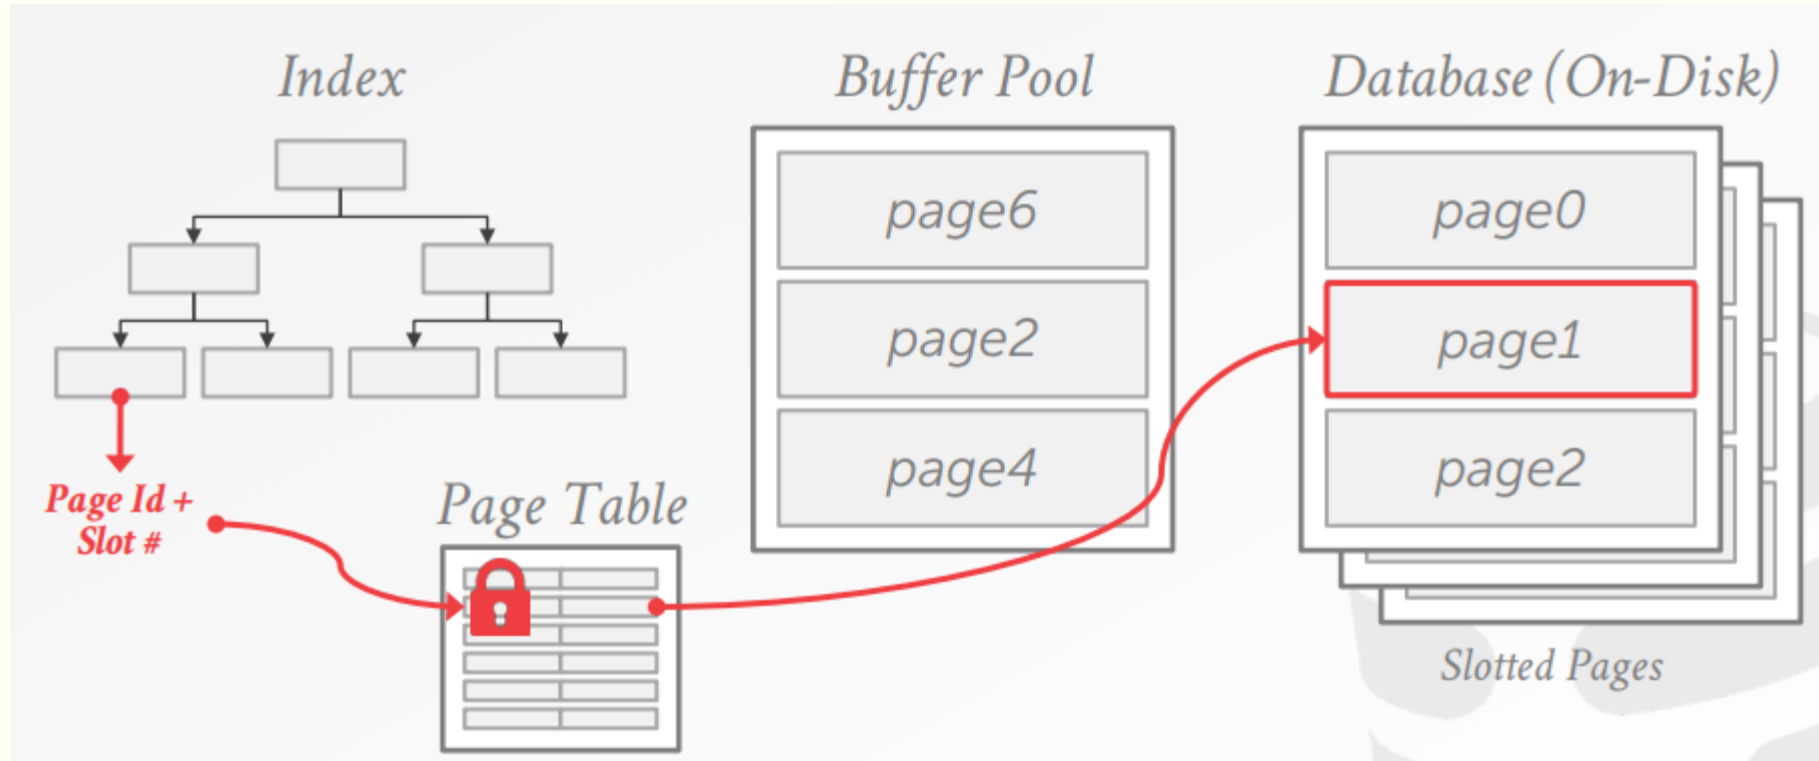
# Buffer Pool Management

---

- When Query access a page:
  - DBMS checks if page is in memory
  - No – retrieve from disk and copy into frame of buffer pool
    - No free frames: find a page to evict
    - Dirty page evicted – write back to disk
  - Yes – translate on-disk addresses to in-memory addresses

## Example: page request

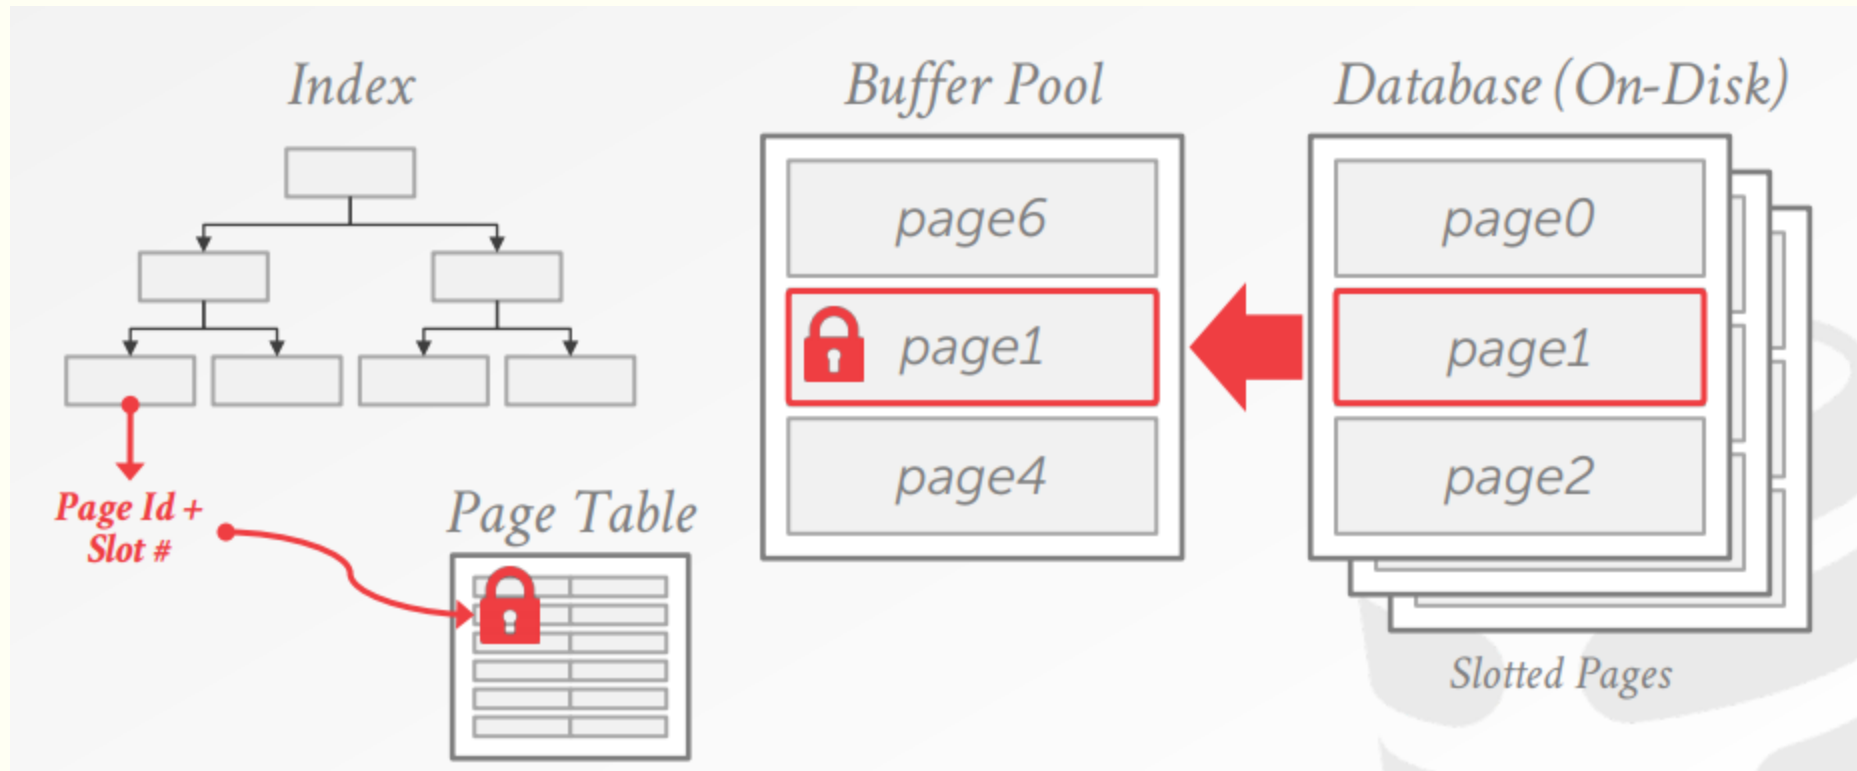
---





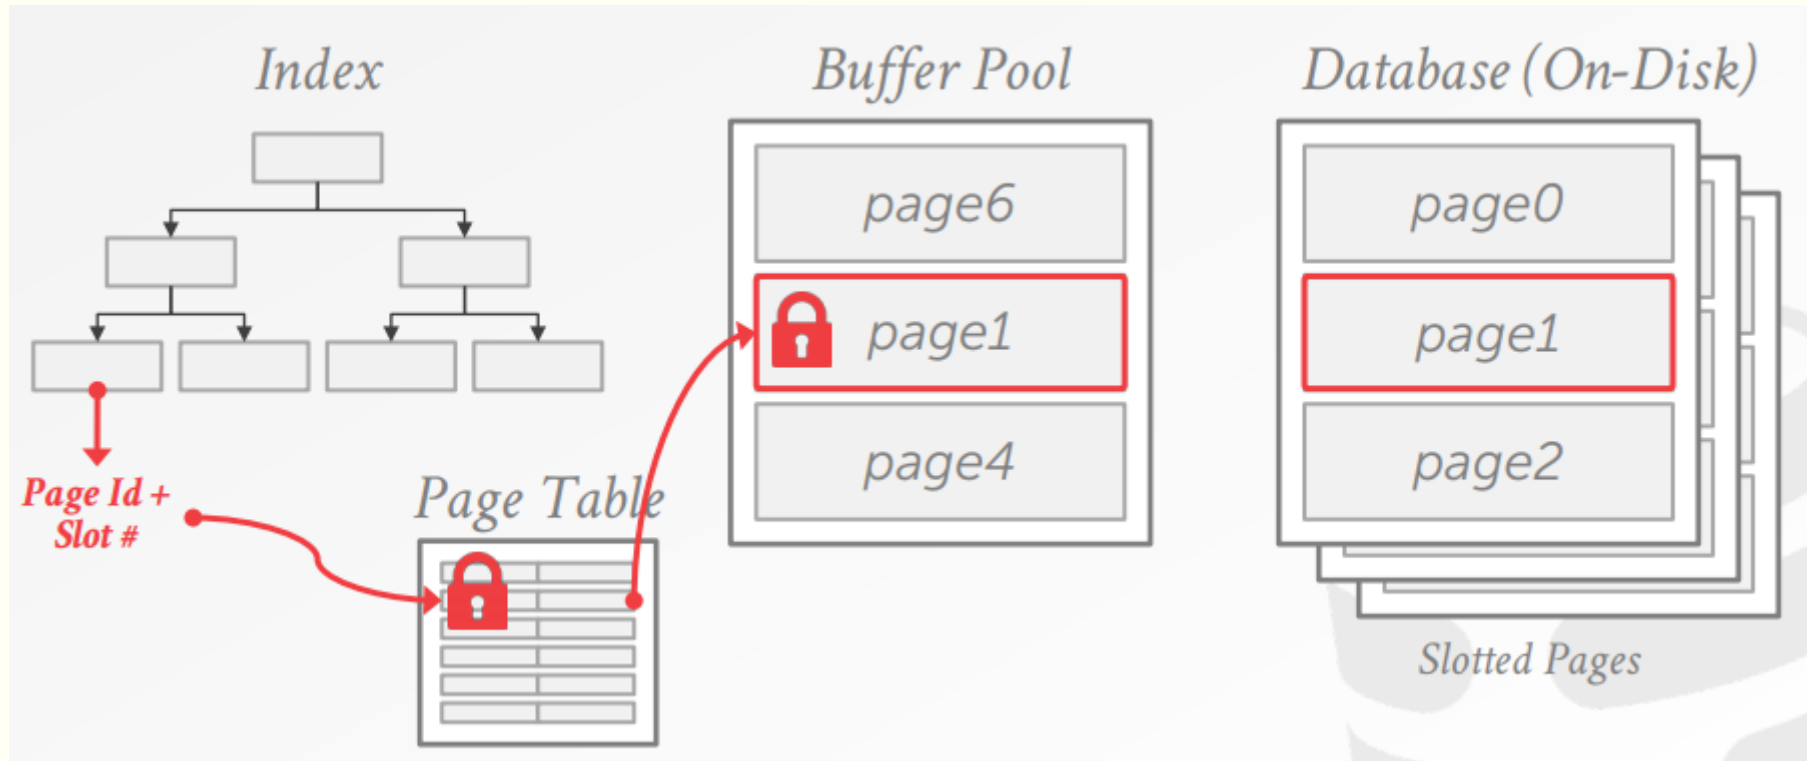
## Example: page request

---



## Example: page request

---




# Overhead

---

- Buffer pool: Each tuple access needs to go through buffer pool management regardless of whether the data will always be in memory
  - Translate tuple record id to memory location
  - Pin pages to make sure they are not swapped to disk
- Concurrency control
  - ACID guarantee: set locks and latches
- Logging & recovery
  - “steal” + “no-force” buffer pool policy
  - Log contains before and after images of modified record.

# Overhead

---



“Removing those overheads and running the database in main memory would yield orders of magnitude improvements in database performance”

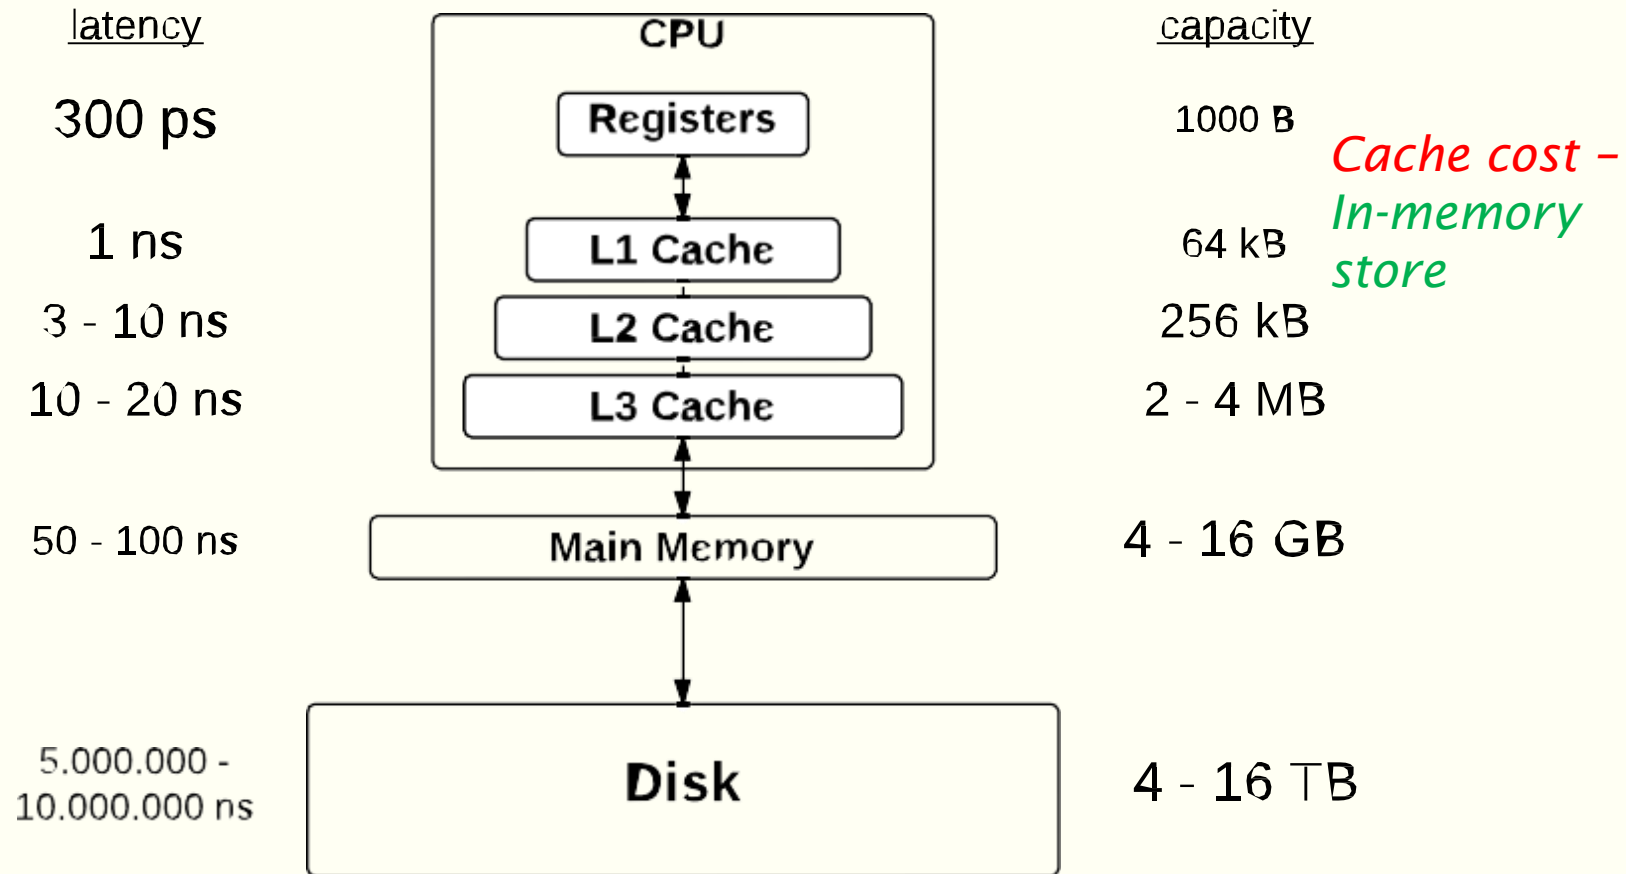
# NewSQL Design Principles

---

- SQL + ACID + performance and scalability through modern innovative software architecture
- *Principle 1: minimizing or stay away from locking*
- *Principle 2: rely on main memory*
- *Principle 3: try to avoid latching*
- *Principle 4: cheaper solutions for HA*

# In-Memory Databases

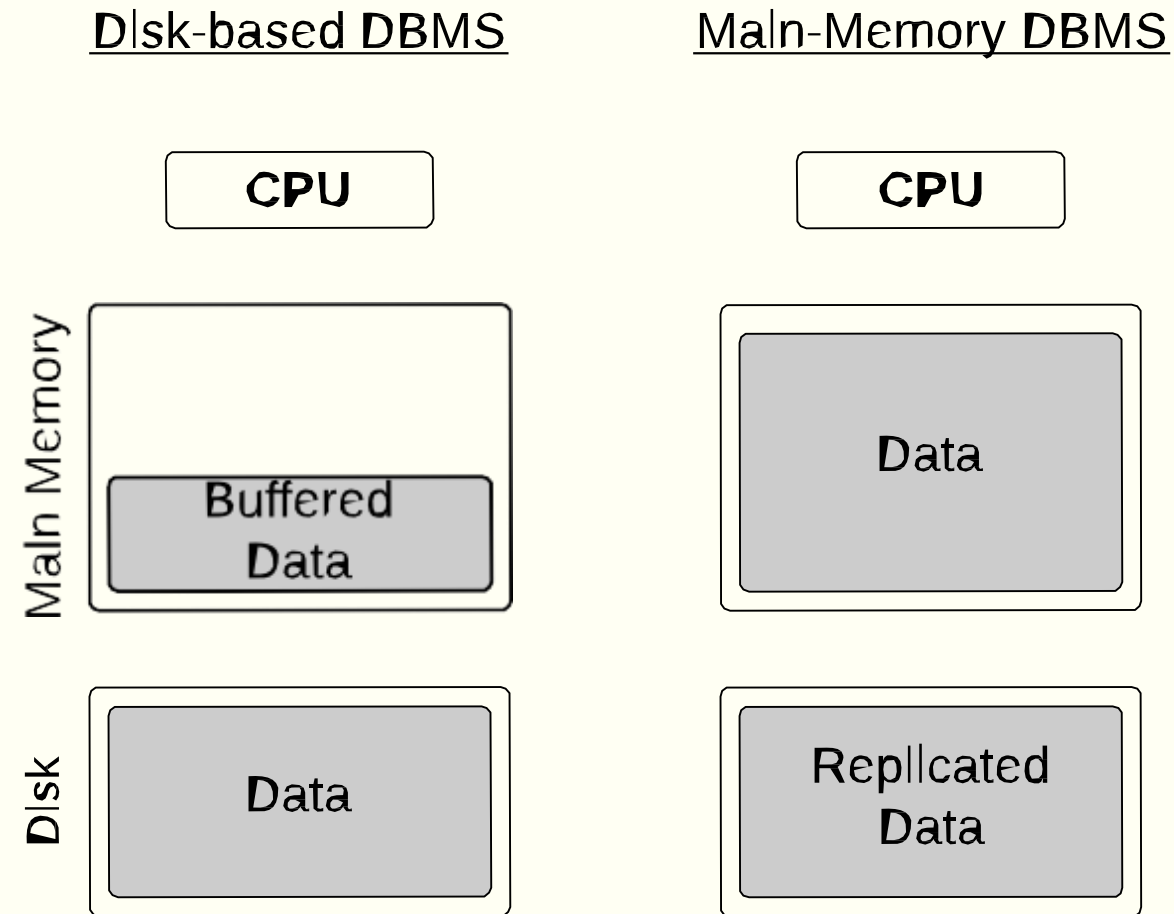
---



Data taken from [Hennessy and Patterson, 2012]

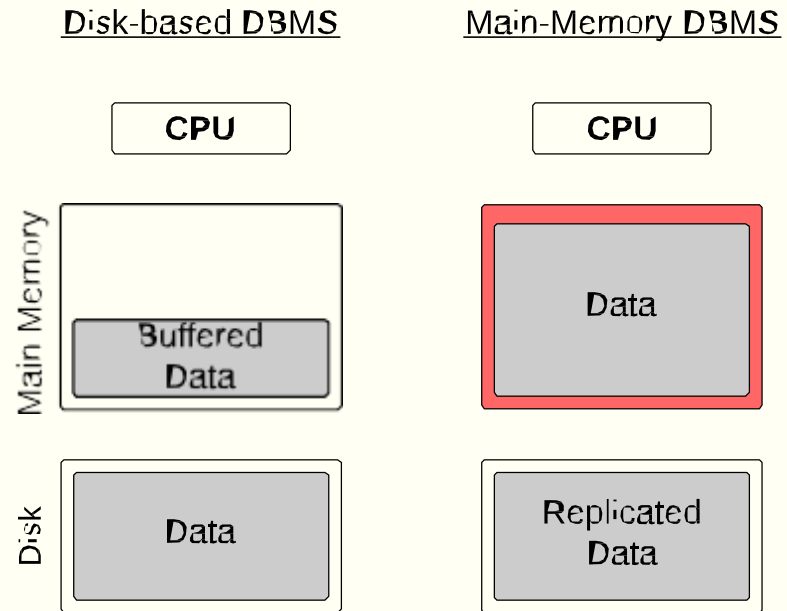
# Disk-based vs. Main-Memory DBMS

---



# Disk-based vs. Main-Memory DBMS (2)

---



**ATTENTION:** Main-memory storage != No Durability

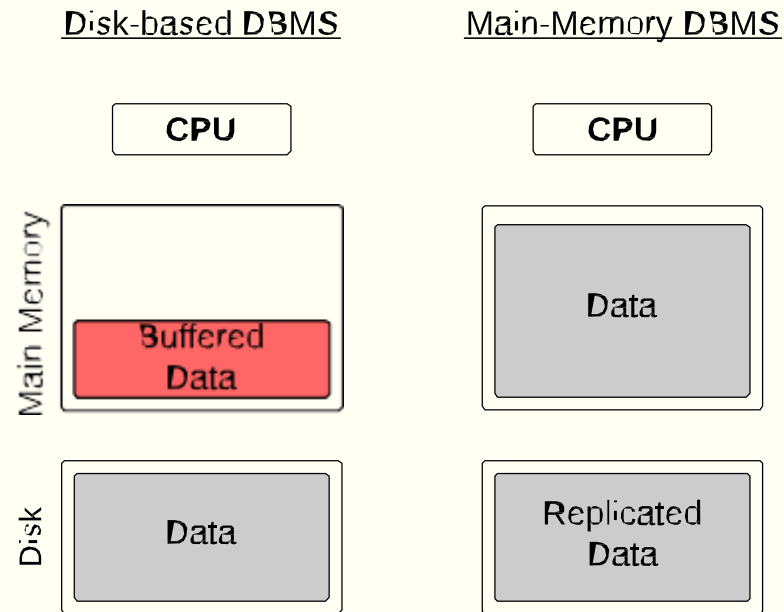
→ ACID properties have to be guaranteed

→ However, there are new ways of guaranteeing it, such as a second machine in hot standby



# Disk-based vs. Main-Memory DBMS (3)

---




Having the database in main memory allows us to remove buffer manager and paging

→ Remove level of indirection

→ Results in better performance

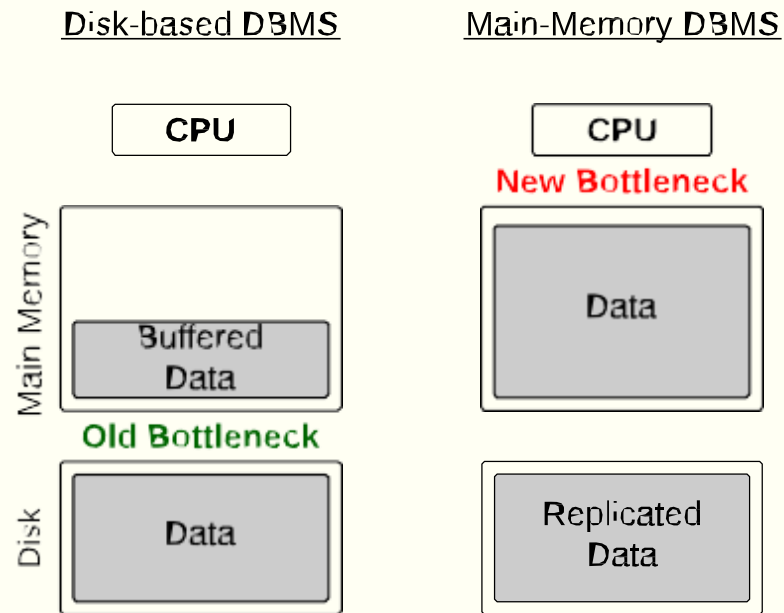
# In-memory Data Organization

---



# Disk-based vs. Main-Memory DBMS (4)

---




Disk bottleneck is removed as database is kept in main memory

→ Access to main memory becomes new bottleneck

# The New Bottleneck: Memory Access

---



Picture taken from [Manegold et al., 2000]

Accessing main-memory is much more expensive than accessing CPU registers.

→ Is main-memory the new disk?

# New Bottleneck

---

- When disk I/O is no longer the bottleneck...
  - Locking/latching
  - Cache-line misses
  - Data movement

# Rethink the Architecture of DBMSs

---

Even if the complete database fits in main memory, there are significant overheads of traditional DBMSs:

- Many function calls → stack manipulation overhead + instruction-cache misses
- Adverse memory access → data-cache misses

→ Be aware of the caches!

# Cache Awareness: Principle of Locality


---

- Caches take advantage of the principle of locality.
  - The hot set of data often fits into caches.
  - 90 % execution time spent in 10 % of the code.
- Spatial Locality:
  - Related data is often spatially close.
  - Code often contains loops.
- Temporal Locality:
  - Programs tend to re-use data frequently.
  - Code may call a function repeatedly, even if it is not spatially close.

# CPU Cache Internals

---

- To guarantee speed, the overhead of caching must be kept reasonable.
- Organize cache in cache lines.
- Only load/evict full cache lines.
- Typical cache line size: 64 bytes.



cache line

The organization in cache lines is consistent with the principle of (spatial) locality.



# Memory Access


---

- On every memory access, the CPU checks if the respective cache line is already cached.
- Cache Hit:
  - Read data directly from the cache.
  - No need to access lower-level memory.
- Cache Miss:
  - Read full cache line from lower-level memory.
  - Evict some cached block and replace it by the newly read cache line.
  - CPU stalls until data becomes available.
- Modern CPUs support out-of-order execution and several in-flight cache misses.

# Block Placement: Direct-Mapped Cache

---


- In a direct-mapped cache, a block has only one place it can appear in the cache.
  - Much simpler to implement.
  - Easier to make fast.
  - Increases the chance of conflicts.



# Block Placement: Fully Associative Cache

---

- In a fully associative cache, a block can be loaded into any cache line
- Provide freedom to block replacement strategy.
- Does not scale to large caches
  - → 4 MB cache,
  - line size: 64 B: 65,536 cache lines.




# Block Placement: Set-Associative Cache

---

A compromise are set-associative caches.

- Group cache lines into sets.
- Each memory block maps to one set.
- Block can be placed anywhere within a set.
- Most processor caches today are set-associative.



# Block Replacement

---

- When bringing in new cache lines, an existing entry has to be evicted:
- Least Recently Used (LRU)
  - Evict cache line whose last access is longest ago.
  - → Least likely to be needed any time soon.
- First In First Out (FIFO)
  - Behaves often similar like LRU.
  - But easier to implement.
- Random
  - Pick a random cache line to evict.
  - Very simple to implement in hardware.
- Replacement has to be decided in hardware and fast.

# What Happens on a Write?

---

- To implement memory writes, CPU makers have two options:
- Write Through
  - Data is directly written to lower-level memory (and to the cache).
  - → Writes will stall the CPU.
  - → Greatly simplifies data coherency.
- Write Back
  - Data is only written into the cache.
  - A dirty flag marks modified cache lines (Remember the status field.)
  - → May reduce traffic to lower-level memory.
  - → Need to write on eviction of dirty cache lines.
- Modern processors usually implement write back.

## Putting it all Together

---

- To compensate for slow memory, systems use caches.
  - Typically multiple levels of caching (memory hierarchy).
  - Caches are organized into cache lines.
  - Set associativity: A memory block can only go into a small number of cache lines (most caches are set-associative).
- In-memory DBMS will benefit from locality of data and code.

# Processing Models

---


- There are basically two alternative processing models that are used in modern DBMSs:
  - Tuple-at-a-time volcano model [Graefe, 1990]
  - Operator requests next tuple, processes it, and passes it to the next operator
  - Operator-at-a-time bulk processing [Manegold et al., 2009]
  - Operator consumes its input and materializes its output



# Tuple-At-A-Time Processing

---

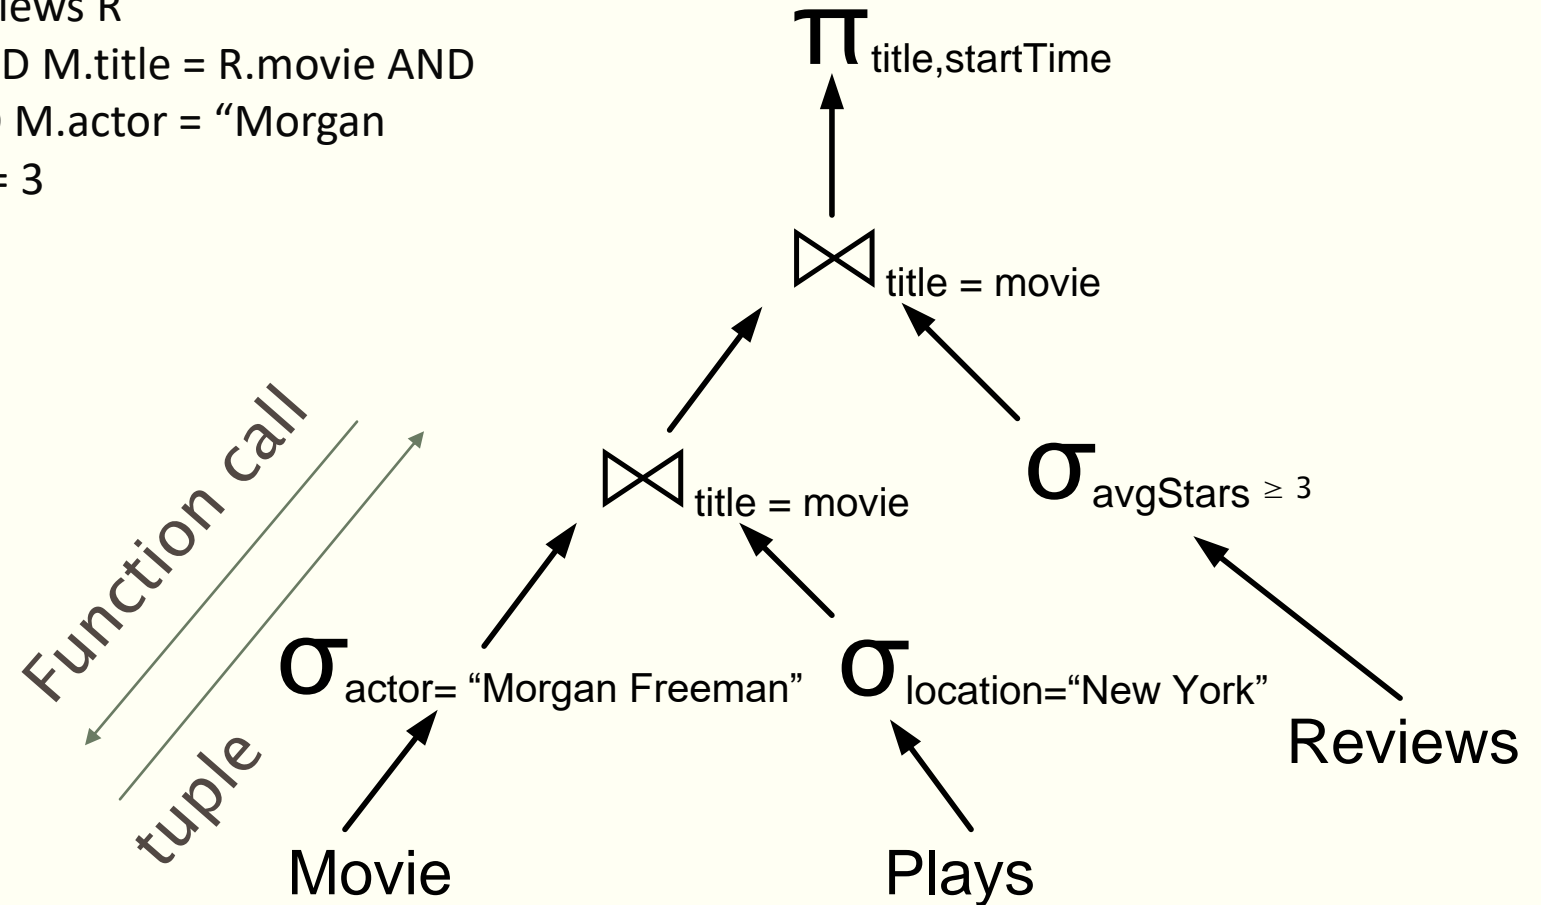
- Most systems implement the Volcano iterator model:
  - Operators request tuples from their input using `next ()`.
  - Data is processed tuple at a time.
  - Each operator keeps its own state.



## Example Logical Query Plan (revisit Lecture 7)

---

SELECT title, startTime  
FROM Movie M, Plays P, Reviews R  
WHERE M.title = P.movie AND M.title = R.movie AND  
P.location = "New York " AND M.actor = "Morgan  
Freeman" AND R.avgStars >= 3



# Tuple-At-A-Time Processing - Consequences

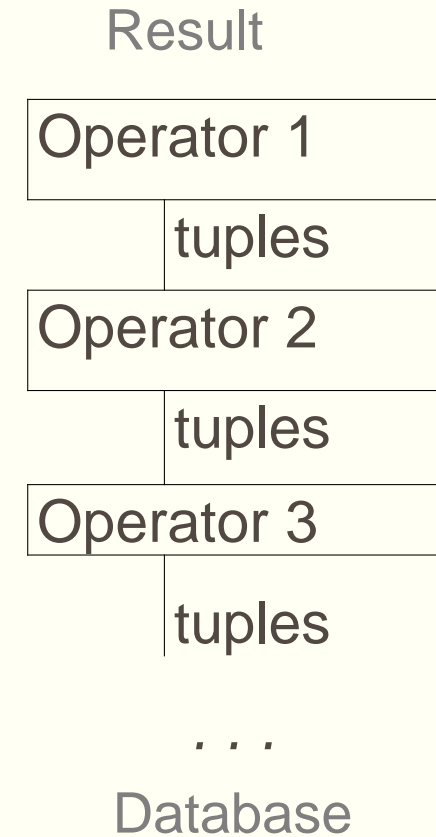
---

- Pipeline-parallelism
  - → Data processing can start although data does not fully reside in main memory
  - → Small intermediate results
- All operators in a plan run tightly interleaved.
  - → Their combined instruction footprint may be large.
  - → Instruction cache misses.
- Operators constantly call each other's functionality.
  - → Large function call overhead.
- The combined state may be too large to fit into caches.
  - E.g., hash tables, cursors, partial aggregates.
  - → Data cache misses.
- Not a good option for in-memory DBMS (especially OLAP DBMS)

# Operator-At-A-Time Processing

---

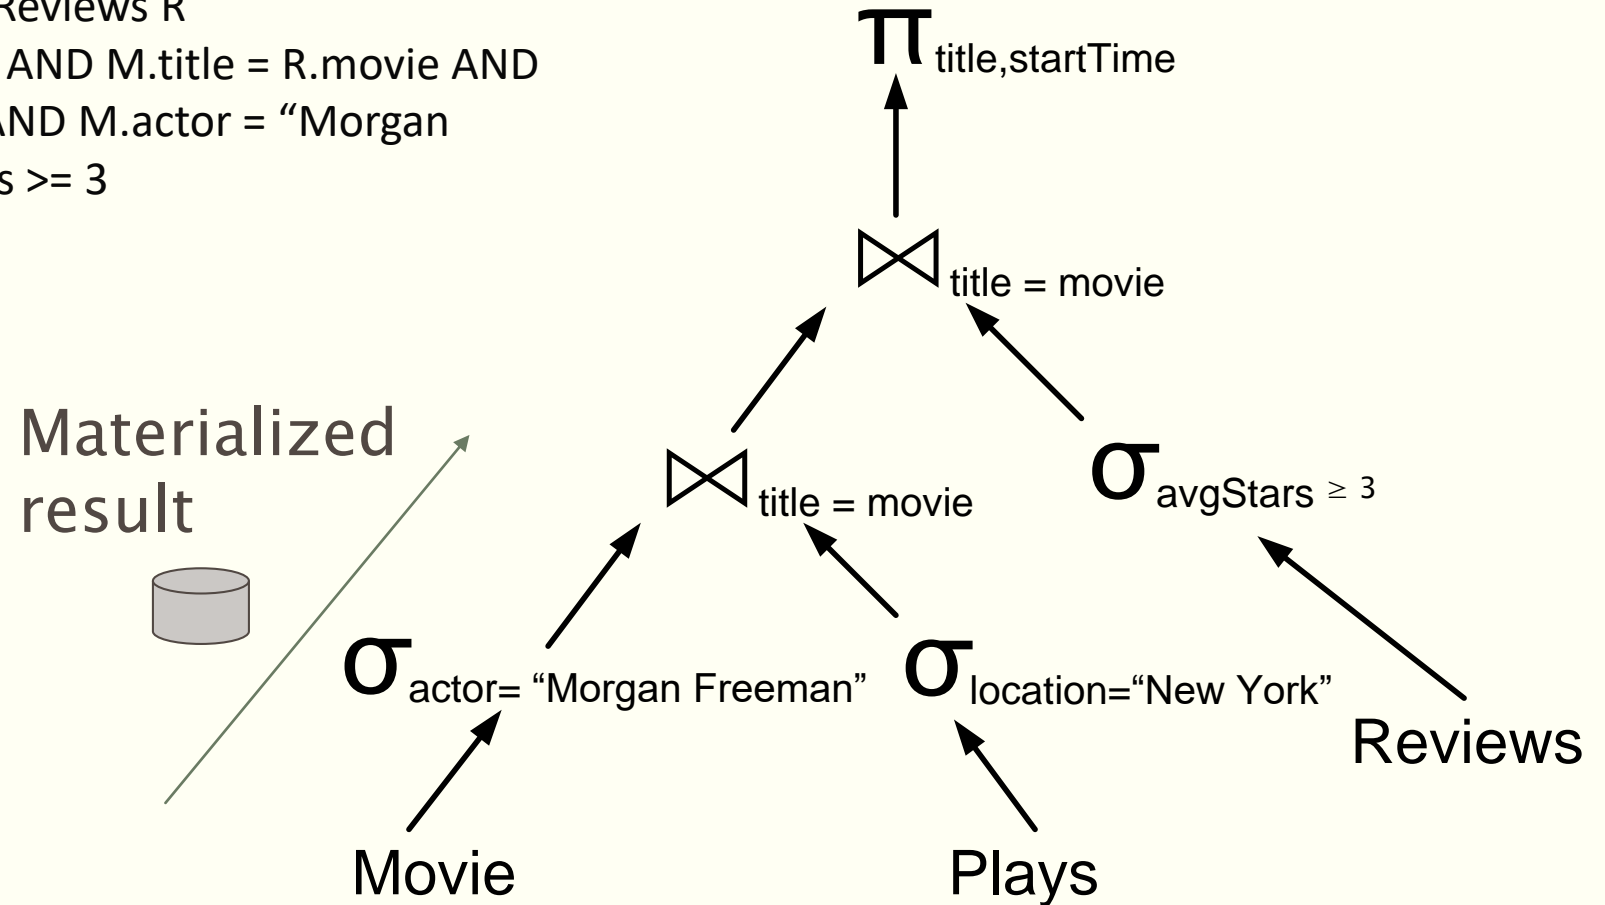
- Operators consume and produce full tables.
- Each (sub-)result is fully materialized (in memory).
- No pipelining (rather a sequence of statements).
- Each operator runs exactly once.



# Example Logical Query Plan (revisit Lecture 7)

---

SELECT title, startTime  
FROM Movie M, Plays P, Reviews R  
WHERE M.title = P.movie AND M.title = R.movie AND  
P.location = "New York " AND M.actor = "Morgan  
Freeman" AND R.avgStars >= 3



# Operator-At-A-Time Consequences


---

- Parallelism: Inter-operator and intra-operator
- Function call overhead is now replaced by extremely tight loops that
  - conveniently fit into instruction caches,
  - can be optimized effectively by modern compilers
- Function calls are now out of the critical code path.
- No per-tuple field extraction or type resolution.
  - Operator specialization, e.g., for every possible type.
  - Implemented using macro expansion.
  - Possible due to column-based storage.
- Implemented in H-store and VoltDB
  - Fine for OLTP.
  - What about OLAP?

# Vectorized Execution Model

---

- Idea:
  - Use Volcano-style iteration,
- but:
  - for each `next ()` call return a large number of tuples
  - → a so called “vector”
- Choose vector size
  - large enough to compensate for iteration overhead (function calls, instruction cache misses, . . . ), but
  - small enough to not thrash data caches.



# What if larger-than-memory?

---

- Hybrid workload:
  - OLAP + OLTP
  - Small, frequently updated: “Hot Data” -- OLTP
    - News, social activities, posts, fresh data, fast data
    - Main Memory
  - Large, infrequent updated but support analytical queries: “Cold Data” – OLAP
    - SSD, Hard disk
- A comparison with Disk-based systems. Hot vs. Cold.



# A Vision

---

- Non-volatile memory – storage level memory
  - Same read/write speed as DRAM
  - Persistent guarantee of SSD
- High-speed DRAM networks & Systems-on-a-Chip
  - Game changer for parallel/distributed algorithm design
- In-memory Data Analytics Systems
  - Big Data in your laptop!

# Conclusion

---

- Overhead of Disk-based DBMS
  - Buffer pool
  - Concurrency control
  - Locking/latching
- In-memory DBMS
  - Data organization
  - Cache awareness
  - Query processing models
- What we haven't talked: Indexing?
  - (T-trees: read: <http://www.vldb.org/conf/1986/P294.PDF>)