




Distributed Query Processing

How to Make Big Data Small?

- **Input:** A class \mathcal{Q} of queries
- **Question:** Can we effectively find, given queries $Q \in \mathcal{Q}$ and any (possibly big) data D , a small D_Q such that $Q(D) = Q(D_Q)$?



- Data synopsis
- Boundedly evaluable queries
- Query answering using views
- Incremental evaluation
- Distributed Query Processing


Parallel DBMS

- Why parallel DBMS?
- Architectures
- Parallelism
 - Intra-query Parallelism
 - Inter-query Parallelism
 - Intra-operation Parallelism
 - Inter-operation Parallelism


Performance of a DBMS

- Throughput:
 - The number of tasks finished in a given time interval
- Response Time (Latency):
 - The amount of time to finish a single task from the time it is submitted
- Can we do better given more resources (CPU, disk, memory, ...)?
- Parallel DBMS: exploring parallelism
 - Divide a big problem into many smaller ones to be solved in parallel

Traditional DBMS




Parallel DBMS



Degree of Parallelism: Speedup


- Speedup: TS/TL , for a given task
 - TS : time taken by a traditional DBMS
 - TL : time taken by a parallel DBMS with more resources
 - TS/TL : Ideally, more resources mean proportionally less time for a task
- Linear speedup:
 - The speedup is N while the parallel system has N times resources of the traditional system
 - Can we do better?

Speed: throughput response time



Degree of Parallelism: Scaleup


- Scaleup: TS/TL
 - Factor that expresses how much more work can be done in the same time period by a larger system
 - A task Q and a task Q_N , which is N times bigger than Q
 - A DBMS M_S and a parallel DBMS M_L , N times more resources
 - TS : time taken by a traditional DBMS
 - TL : time taken by a parallel DBMS with more resources
 - Linear scaleup if $\frac{TS}{TL} = 1$
 - The time is constant if the resource increases in proportion to increase in problem size



Why can't it be better than linear scaleup/speedup?

- Startup costs: initializing each process
- Interference: competing for shared resources (network, disk, memory or locks)
- Skew:
 - It is difficult to divide a task into exactly equal-sized parts
 - The response time (latency) is determined by the largest part
- Amdahl's law:
$$\text{Speedup} = \frac{1}{f + \frac{1-f}{s}}$$

f : “sequential fraction” of the program
 s : Amount of parallel resources

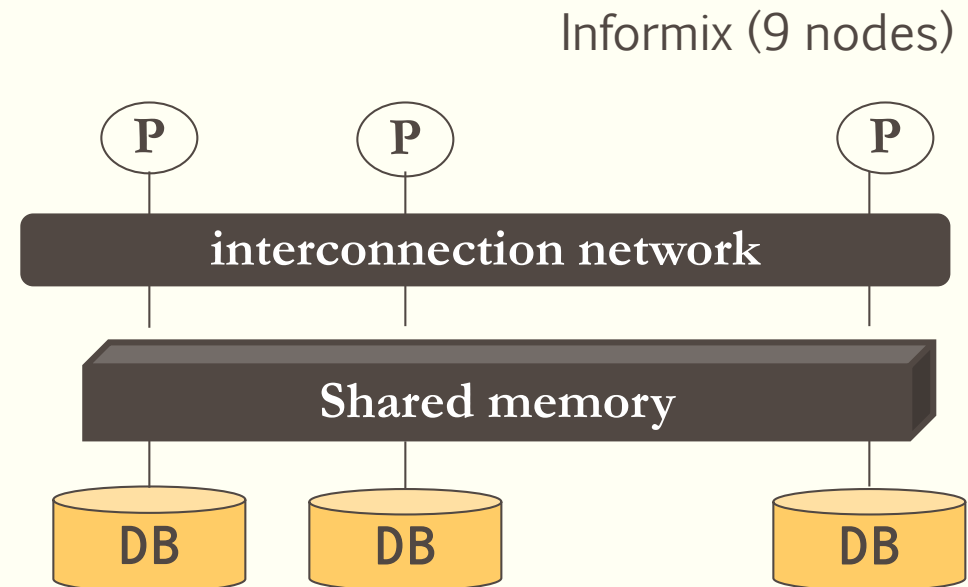


Why Parallel DBMS

- Improve Performance
- Almost died 25 years ago, but with renewed interests
 - Big Data: Data collected from the web
 - Decision Support Queries: Costly on large data
 - Hardware has become much cheaper
- Improve **reliability** and **availability**
- **MapReduce**


Architecture: Shared Memory

- Efficient Communication
 - Via data in memory, accessible by all
- Not Scalable:
 - Shared memory and network become bottleneck – interference
 - Not scalable beyond 32/64 processors
 - Adding memory cache to each processor?
 - Cache coherence problem when data is updated



Shared Disk

- Fault tolerance:
 - If a processor fails, the other can take over, since the database is resident on disk
- Scalability:
 - Better than shared memory: memory is no longer a bottleneck
 - But disk subsystem is a bottleneck
- Interference:
 - All I/O to go through a single network
 - Not scalable beyond a couple of hundred processors




Shared Nothing

- Scalable:
 - Only queries and result relations pass through the network
- Communication costs and access to non-local disks
 - Sending data involves software interaction at both ends

Teradata: 400 nodes


IBM SP2/DB2: 128 nodes

Informix SP2: 48 nodes



Architecture Summary


Shared Memory (SMP)



Easy to program
Expensive to build
Difficult to scaleup

Informix, RedBrick
Sequent, SGI, Sun
scale: 9 nodes


Shared Disk



Better scalability
Fault tolerance

VMSccluster,
Oracle (170 nodes)
DEC Rdb (24 nodes)

Shared Nothing (network)



Hard to program
Cheap to build
Easy to scaleup

| | |
|--------------|-----------|
| Teradata: | 400 nodes |
| Tandem: | 110 nodes |
| IBM/SP2/DB2: | 128 nodes |
| Informix/SP2 | 48 nodes |

Architectures of Parallel DBMS

- Tradeoffs
 - Scalability
 - Communication Speed
 - Cache Coherence

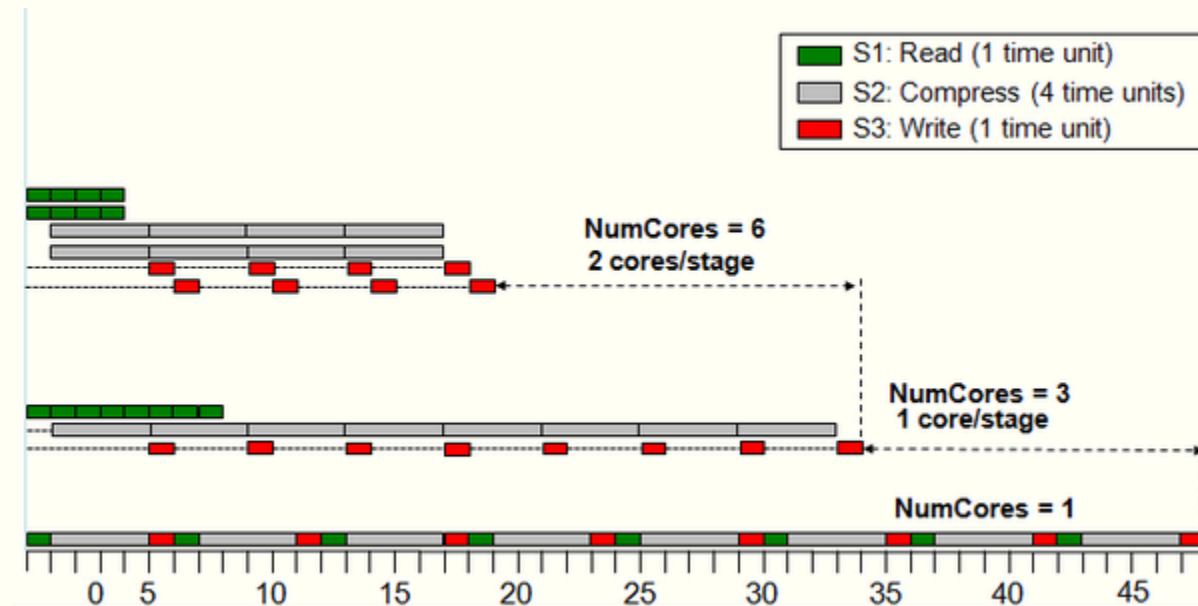
Pipelined

- The output of operation A is consumed by another operation B, before A has produced the entire output
 - Many machines, each doing one step in a multi-step process
 - May lead to increase in response time
- Does not scale up well when
 - The computation does not provide sufficiently long chain to provide a high degree of parallelism
 - Relational operators do not produce output until all inputs have been accessed (blocking)
 - A's computation cost is much higher than that of B

Pipelined Parallelism: Compress a File


```
while(!done){  
    Read block from file;  
    Compress the block;  
    Write block to file;  
}
```

```
while(!program end){  
    If work is available in the  
    in-Queue of thread  
    Compress the block;  
    Write block to file;  
}
```



Data partitioned Parallelism

- Many machines performing the **same** operation on **different** pieces of data (similar to SIMD)



Partitioning in RDMS

- Partition a relation and distribute it to different processors
 - Maximize processing at each individual processor
 - Minimize data shipping
- Query types
 - Scan a relation
 - Point access $r.A = v$
 - Range Queries $v < r.A < v'$

Partitioning Strategy

- Assume N disks, a relation R
- Round Robin
 - Send the j -th tuple of R to disk number $j \bmod N$
 - Even distribution: Good for scanning
 - Not good for equal joins (point queries) and range queries (all disk have to be involved for the search)
- Range Partitioning:
 - Partitioning attribute A , vector $[v_1, \dots, v_{n-1}]$
 - Send tuple t to disk j if $t[A]$ in $[v_{j-1}, v_j]$
 - Good for point and range queries on partitioning attributes (using only a few disks, while leaving the others free)
 - Execution skew: distribution may not be even, and all operations occur in one or few partitions (scanning)

Partitioning Strategies (Cont'd)


- Assume N disks, a relation R
- Hash Partitioning:
 - Hash function $f(t)$ in the range of integer $[0, N - 1]$
 - Send tuple t to disk $f(t)$
 - Good for point queries on partitioning attributes, and sequential scanning if the hash function is even
 - No good for point queries on non-partitioning attributes and range queries

Question: how to partition $R1(A, B): \{(i, i + 1)\}$, with 5 processors?


- Round-robin
- Range partitioning: partitioning attribute A
- Hash partitioning

Automatic Data Partitioning


- Partitioning a table



Good for group-by,
range queries, and
also equip-join



Good for equijoins



Good to spread load;
Most flexible
Not good for equi-join
and range


Shared-disk and -memory less sensitive to partitioning,
Shared nothing benefits from "good" partitioning

Inter-query Vs. Intra-query Parallelism

- Inter-query:
 - Different queries or transactions execute in parallel
 - Improve transaction throughput
 - Easy on shared-memory: Traditional DBMS tricks will do
 - Shared-nothing/Disk: Cache coherence problem
 - Ensure that each processor has the latest version of the data in its buffer pool
 - Flush updated pages to shared disk before releasing the lock
- Intra-Query
 - A single query in parallel on multiple processors
 - Speed up single complex long running queries
 - **Interoperation**: operator tree
 - **Intraoperation**: parallelize the same operation on different sets of the same relations: Parallel sorting, Parallel join; Selection; Projection; Aggregation

Parallel Query Answering

- Given data D , and n processors S_1, S_2, \dots, S_n
 - D is partitioned into fragments (D_1, D_2, \dots, D_n)
 - D is distributed to n processors: D_i is stored at S_i
- Each processor S_i processes operations for a query on its local fragment D_i , in parallel



Relational Operators

- Projection $\pi_A R$
- Selection $\sigma_C R$
- Join $R_1 \bowtie_C R_2$
- Union $R_1 \cup R_2$
- Set Difference $R_1 - R_2$
- Group by and Aggregation
 - Max, min, count, average, ...

Intra-operation Parallelism: Projection

- Projection $\pi_A R$, where R is partitioned across n processors
 - Read tuples of R at all processors involved, in parallel
 - Conduct projection on tuples
 - Merge local results – to eliminate duplicate elimination (via sorting?)

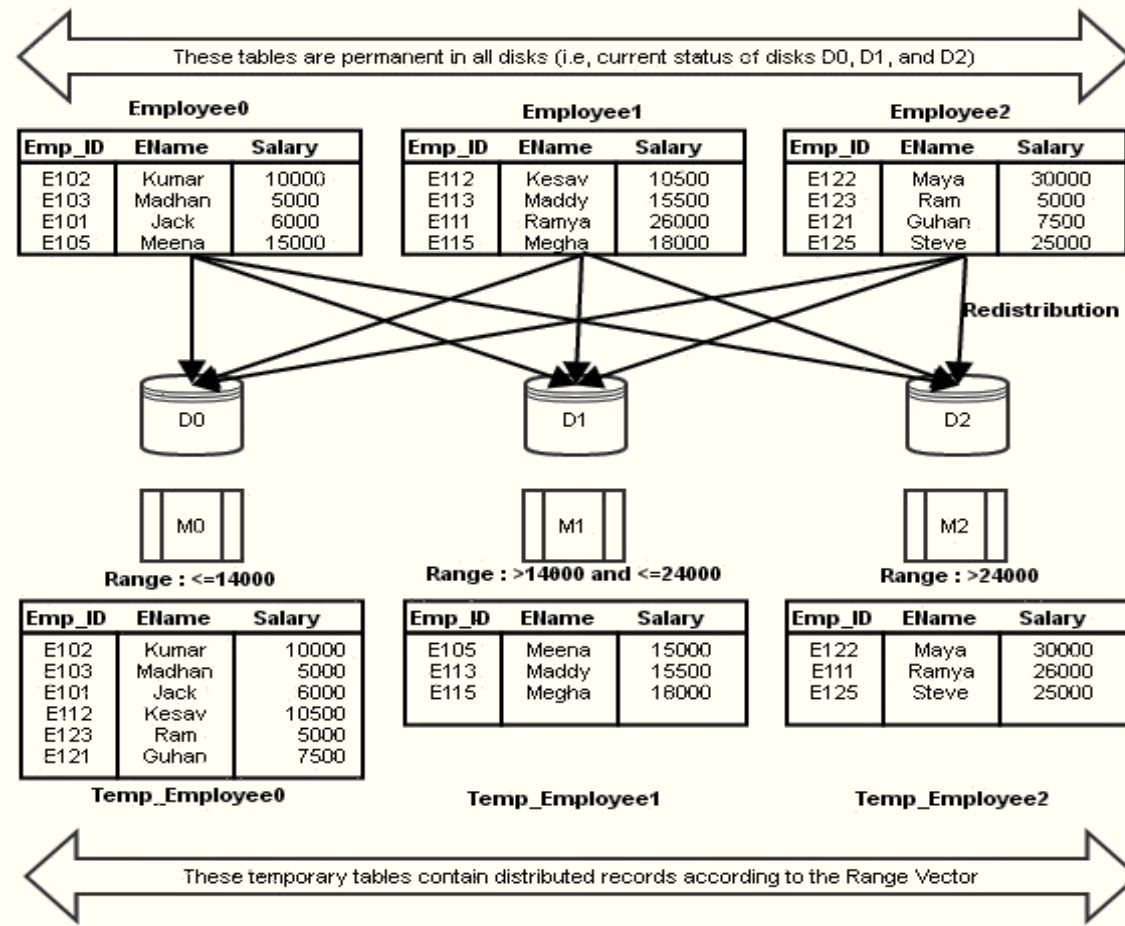
Intra-Operation Parallelism: Selection

- Selection $\sigma_C R$, where R is partitioned across n processors
- If A is the partitioning attribute
 - Point query $C: r.A = v$
A single processor that holds $r.A = v$ is involved
 - Range query $C: v \leq r.A \leq v'$
Only processors whose partition overlaps with the range are involved
- If A is not the partitioning attribute
 - Compute selection at each individual processor
 - Merge local result

Intra-Operation Parallelism: Sort

- Sort R on attribute A , where R is partitioned across n processors
- If A is the partitioning attribute: Range-partitioning
 - Sort each partition
 - Concatenate the result
- If A is not the partitioning attribute: Range-partitioning sort
 - Range partitioning R based on A , redistributed the tuples in R
 - Every processor works in parallel: read tuples and send them to corresponding processors
 - Each processor sorts its new partition locally when the tuple come in
 - Merge local results
- Issue: skew
- Solution: sample the data to determine the partitioning vector

Example: Parallel Sort




Intra-Operation Parallelism: Join

- Partitioned join: for equi-joins and natural joins
- Fragment-and replication join: inequality
- Partitioned parallel hash-join: equal or natural join
 - where R_1, R_2 are too large to fit in memory
 - Almost always the winner for equi-joins


Partitioned Join

- $R_1 \bowtie_{R_1.A=R_2.B} R_2$
 - Partition R_1 and R_2 into n partitions, by the same partitioning function in $R_1.A$ and $R_2.B$, via either
 - range partitioning, or
 - hash partitioning
 - Compute $R_1^i \bowtie_{R_1.A=R_2.B} R_2^i$ locally at processor i
 - Merge the local results
- Question: how to perform partitioned join on the following, with 2 processors?
 - $R_1(A, B): \{(1, 2), (3, 4), (5, 6)\}$
 - $R_2(B, C): \{(2, 3), \{3, 4)\}$



Fragment and Replicate join

- $R_1 \bowtie_{R_1.A < R_2.B} R_2$
 - Partition R_1 into n partitions, by any partitioning method, and distribute it across n processors
 - Replicate the other relation R_2 across all processors
 - Compute $R_1^j \bowtie_{R_1.A < R_2.B} R_2$ locally at processor j
 - Merge the local results
- Question: how to perform fragment and replicate join on the following, with 2 processors?
 - $R1(A, B): \{(1, 2), (3, 4), (5, 6)\}$
 - $R2(B, C): \{(2, 3), \{3, 4)\}$



Partitioned Parallel Hash Join

- $R_1 \bowtie_{R_1.A=R_2.B} R_2$
 - Hash partitioning R_1 and R_2 using hash function h on partitioning attributes $R_1.A$ and $R_2.B$, respectively
 - For $i \in [1, k]$, process the join of i-th partition $R_1^i \bowtie R_2^i$, with hash join

Intra-Operation Parallelism: Aggregation

- Aggregate on the Attribute B of R , grouping on A
- Decomposition:
 - $count(S) = \sum count(S_i)$; similar for sum
 - $avg(S) = sum(S)/count(S)$
- Strategy 1:
 - Range partitioning R based on A : redistribute the tuples in R
 - Each processor computes sub-aggregate (data parallelism)
 - Merge local results as above
- Strategy 2:
 - Each processor computes sub-aggregate (data parallelism)
 - Range partitioning local results based on A , redistribute partial results
 - Compose the local results

Example: Aggregation

- Describe a good processing strategy to parallelize the following query
- ```
SELECT branch-name, avg(balance)
FROM account
GROUP BY branch-name
```

- The schema for the account is

account(account-id, branch-name, balance)

- Strategy:
  - Range or hash partition account by using branch-name as the partitioning attribute. This creates table  $account_j$  at each site  $j$ .
  - At each site  $j$ , compute  $\frac{sum(account_j)}{count(account_j)}$ ;
  - output  $\frac{sum(account_j)}{count(account_j)}$ : the union of these partial results is the final query answer

# Inter-Operation Parallelism

---

- Consider  $R_1 \bowtie R_2 \bowtie R_3 \bowtie R_4$
- Pipelined:
  - $Temp1 \leftarrow R_1 \bowtie R_2$
  - $Temp2 \leftarrow R_3 \bowtie Temp1$
  - $Result \leftarrow R_4 \bowtie Temp2$
- Independent
  - $Temp1 \leftarrow R_1 \bowtie R_2$
  - $Temp2 \leftarrow R_3 \bowtie R_4$
  - $Result \leftarrow Temp1 \bowtie Temp2$  (Pipelined Stage)

# Cost Model

---

- Cost model: partitioning, skew, resource contention, scheduling
  - Partitioning:  $T_{part}$
  - Cost of assembling local answers:  $T_{asm}$
  - Skew:  $\max(T_0, \dots, T_n)$
  - Estimation:  $T_{part} + T_{asm} + \max(T_0, \dots, T_n)$

May also include startup costs and contention for resources (in each  $T_j$ )
- Query optimization: find the “best” parallel query plan
  - Heuristic 1: parallelize all operations across all processors -- partitioning, cost estimation (Teradata)
  - Heuristic 2: best sequential plan, and parallelize operations -- partition, skew, ... (Volcano parallel machine)

# Practice: Validation of Functional Dependencies

---

- Develop a parallel algorithm that given a relation  $D$  and an Functional dependency  $FD$ , computes all the violations
  - Partitioned Join
  - Partitioned and replication Join
- Question: what can we do if we are given a set of FDs to validate?

## Practice: Implement Set Difference

---

- Develop a parallel algorithm that  $R1$  and  $R2$ , compute  $R1-R2$ , by using:
  - partitioned join
  - partitioned and replicated
- Questions: what can we do if the relations are too large to fit in memory?